# Oracle® Big Data SQL User's Guide





Oracle Big Data SQL User's Guide, Release 3 (3.2)

E99050-01

Copyright © 2012, 2018, Oracle and/or its affiliates. All rights reserved.

This software and related documentation are provided under a license agreement containing restrictions on use and disclosure and are protected by intellectual property laws. Except as expressly permitted in your license agreement or allowed by law, you may not use, copy, reproduce, translate, broadcast, modify, license, transmit, distribute, exhibit, perform, publish, or display any part, in any form, or by any means. Reverse engineering, disassembly, or decompilation of this software, unless required by law for interoperability, is prohibited.

The information contained herein is subject to change without notice and is not warranted to be error-free. If you find any errors, please report them to us in writing.

If this is software or related documentation that is delivered to the U.S. Government or anyone licensing it on behalf of the U.S. Government, then the following notice is applicable:

U.S. GOVERNMENT END USERS: Oracle programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, delivered to U.S. Government end users are "commercial computer software" pursuant to the applicable Federal Acquisition Regulation and agency-specific supplemental regulations. As such, use, duplication, disclosure, modification, and adaptation of the programs, including any operating system, integrated software, any programs installed on the hardware, and/or documentation, shall be subject to license terms and license restrictions applicable to the programs. No other rights are granted to the U.S. Government.

This software or hardware is developed for general use in a variety of information management applications. It is not developed or intended for use in any inherently dangerous applications, including applications that may create a risk of personal injury. If you use this software or hardware in dangerous applications, then you shall be responsible to take all appropriate fail-safe, backup, redundancy, and other measures to ensure its safe use. Oracle Corporation and its affiliates disclaim any liability for any damages caused by use of this software or hardware in dangerous applications.

Oracle and Java are registered trademarks of Oracle and/or its affiliates. Other names may be trademarks of their respective owners.

Intel and Intel Xeon are trademarks or registered trademarks of Intel Corporation. All SPARC trademarks are used under license and are trademarks or registered trademarks of SPARC International, Inc. AMD, Opteron, the AMD logo, and the AMD Opteron logo are trademarks or registered trademarks of Advanced Micro Devices. UNIX is a registered trademark of The Open Group.

This software or hardware and documentation may provide access to or information about content, products, and services from third parties. Oracle Corporation and its affiliates are not responsible for and expressly disclaim all warranties of any kind with respect to third-party content, products, and services unless otherwise set forth in an applicable agreement between you and Oracle. Oracle Corporation and its affiliates will not be responsible for any loss, costs, or damages incurred due to your access to or use of third-party content, products, or services, except as set forth in an applicable agreement between you and Oracle.

## Contents

## Preface

Audience		ix
Related Do	ocuments	ix
Conventio	ns	ix
Backus-Na	aur Form Syntax	Х
Changes i	n Oracle Big Data SQL 3.2	X
Introdu	cing Oracle Big Data SQL	
1.1 Wha	at Is Oracle Big Data SQL?	1-1
1.1.1	About Oracle External Tables	1-2
1.1.2	About the Access Drivers for Oracle Big Data SQL	1-2
1.1.3	About Smart Scan for HDFS	1-2
1.1.4	About Storage Indexes	1-3
1.1.5	About Predicate Push Down	1-5
1.1.6	Predicate Pushdown Issue with Parquet Files Created From Hive Source	1-6
1.:	1.6.1 Other Issues With Parquet Statistics	1-8
1.1.7	About Pushdown of Character Large Object (CLOB) Processing	1-8
1.1.8	About Oracle Big Data SQL Statistics	1-10
1.2 Insta	allation	1-11
Using C	Dracle Big Data SQL for Data Access	
2.1 Crea	ating an Oracle External Table for Hive Data	2-1
2.1.1	Obtaining Information About a Hive Table	2-2
2.1.2	Using the CREATE_EXTDDL_FOR_HIVE Function	2-3
2.1.3	Using Oracle SQL Developer to Connect to Hive	2-4
2.1.4	Developing a CREATE TABLE Statement for ORACLE_HIVE	2-6
2.:	1.4.1 Using the Default ORACLE_HIVE Settings	2-6
2.:	1.4.2 Overriding the Default ORACLE_HIVE Settings	2-7
2.2 Crea	ating an Oracle External Table for Oracle NoSQL Database  Creating a Hive External Table for Oracle NoSQL Database	2-8 2-8



	2.2.2	Crea	ting the Oracle Database Table for Oracle NoSQL Data	2-9
	2.2.3	Abou	ut Column Data Type Mappings	2-9
	2.2.4	Exar	nple of Accessing Data in Oracle NoSQL Database	2-10
	2.2	2.4.1	Creating the Oracle NoSQL Database Example Table	2-11
	2.2	2.4.2	Creating the Example Hive Table for vehicleTable	2-12
	2.2	2.4.3	Creating the Oracle Table for VEHICLES	2-13
2.3	Crea	ting a	n Oracle External Table for Apache HBase	2-14
	2.3.1	Crea	ting a Hive External Table for HBase	2-14
	2.3.2	Crea	ting the Oracle Database Table for HBase	2-14
2.4	Crea	ting a	n Oracle External Table for HDFS Files	2-15
	2.4.1	Usin	g the Default Access Parameters with ORACLE_HDFS	2-15
	2.4.2	Over	riding the Default ORACLE_HDFS Settings	2-15
	2.4	.2.1	Accessing a Delimited Text File	2-15
	2.4	.2.2	Accessing Avro Container Files	2-16
2.5	Crea	ting a	n Oracle External Table for Kafka Topics	2-17
	2.5.1		g Oracle's Hive Storage Handler for Kafka to Create a Hive	
			rnal Table for Kafka Topics	2-17
	2.5.2		ting an Oracle Big Data SQL Table for Kafka Topics	2-20
2.6	Usin	g the (	Custom Parquet Reader for Oracle Big Data SQL	2-21
2.7	Abou	it the	SQL CREATE TABLE Statement	2-22
	2.7.1	Basi	c Syntax	2-22
	2.7.2	Abou	ut the External Table Clause	2-22
	2.7	'.2.1	TYPE Clause	2-22
	2.7	'.2.2	DEFAULT DIRECTORY Clause	2-23
	2.7	'.2.3	LOCATION Clause	2-23
	2.7	'.2.4	ORACLE_HDFS LOCATION Clause	2-23
	2.7	'.2.5	ORACLE_HIVE LOCATION Clause	2-24
	2.7	7.2.6	REJECT LIMIT Clause	2-24
	2.7	'.2.7	ACCESS PARAMETERS Clause	2-24
2.8	Abou	ıt Data	a Type Conversions	2-24
2.9	Quer	ying E	External Tables	2-25
	2.9.1	Gran	iting User Access	2-25
	2.9.2	Abou	ut Error Handling	2-26
	2.9.3	Abou	ut the Log Files	2-26
2.1			acle Big Data SQL on the Database Server (Oracle Exadata or Other)	2-26
	2.10.1	Abo	out the bigdata_config Directory	2-26
	2.10.2	Cor	mmon Configuration Properties	2-26
	2.1	0.2.1	bigdata.properties	2-27
	2.1	0.2.2	bigdata-log4j.properties	2-29
	2.10.3	Abo	out the Cluster Directory	2-29



## 3 Storing Oracle Data in Hadoop

	3.1 Using Copy to Hadoop	3-1
	3.1.1 What Is Copy to Hadoop?	3-1
	3.1.2 Getting Started Using Copy to Hadoop	3-2
	3.1.2.1 Table Access Requirements for Copy to Hadoop	3-3
	3.1.3 Using Oracle Shell for Hadoop Loaders With Copy to Hadoop	3-3
	3.1.3.1 Introducing Oracle Shell for Hadoop Loaders	3-3
	3.1.4 Using Copy to Hadoop to do Direct Copies	3-4
	3.1.5 Using Copy to Hadoop to do Staged Copies	3-5
	3.1.6 Querying the Data in Hive	3-5
	3.1.7 About Column Mappings and Data Type Conversions	3-5
	3.1.7.1 About Column Mappings	3-6
	3.1.7.2 About Data Type Conversions	3-6
	3.1.8 Working With Spark	3-6
	3.1.9 Using Oracle SQL Developer with Copy to Hadoop	3-7
	3.2 Storing Oracle Tablespaces in HDFS	3-8
	3.2.1 Advantages and Limitations of Tablespaces in HDFS	3-8
	3.2.2 About Tablespaces in HDFS and Data Encryption	3-8
	3.2.3 Moving Tablespaces to HDFS	3-10
	3.2.3.1 Using bds-copy-tbs-to-hdfs	3-10
	3.2.3.2 Manually Moving Tablespaces to HDFS	3-14
	3.2.4 Smart Scan for TableSpaces in HDFS	3-16
4	Oracle Big Data SQL Reference	
	4.1.1 DBMS_HADOOP PL/SQL Package	4-1
	4.1.1.1 CREATE_EXTDDL_FOR_HIVE	4-1
	4.1.1.1.1 Example	4-2
	4.1.2 DBMS_BDSQL PL/SQL Package	4-3
	4.1.2.1 ADD_USER_MAP	4-3
	4.1.2.2 REMOVE_USER_MAP	4-5
	4.1.2.3 Multi-User Authorization Security Table	4-5
	4.1.3 CREATE TABLE ACCESS PARAMETERS Clause	4-7
	4.1.3.1 Syntax Rules for Specifying Properties	4-7
	4.1.3.2 ORACLE_HDFS Access Parameters	4-8
	4.1.3.2.1 Default Parameter Settings for ORACLE_HDFS	4-8
	4.1.3.2.2 Optional Parameter Settings for ORACLE_HDFS	4-9
	4.1.3.3 ORACLE_HIVE Access Parameters	4-9



	4.1.3.3.2 Optional	Parameter Values for ORACLE_HIVE	4-10
	4.1.3.4 com.oracle.bigo	data.buffersize	4-10
	4.1.3.5 com.oracle.bigo	data.colmap	4-11
	4.1.3.6 com.oracle.bigo	data.datamode	4-12
	4.1.3.7 com.oracle.bigo	data.erroropt	4-12
	4.1.3.8 com.oracle.bigo	data.fields	4-14
	4.1.3.9 com.oracle.bigo	data.fileformat	4-16
	4.1.3.10 com.oracle.big	gdata.log.exec	4-17
	4.1.3.11 com.oracle.big	gdata.log.qc	4-17
	4.1.3.12 com.oracle.big	gdata.overflow	4-18
	4.1.3.13 com.oracle.big	gdata.rowformat	4-19
	4.1.3.14 com.oracle.big	gdata.tablename	4-21
	4.1.4 Static Data Dictionary	y Views for Hive	4-21
	4.1.4.1 ALL_HIVE_DA	TABASES	4-22
	4.1.4.2 ALL_HIVE_TAI	BLES	4-22
	4.1.4.3 ALL_HIVE_CO	DLUMNS	4-23
	4.1.4.4 DBA_HIVE_DA	ATABASES	4-24
	4.1.4.5 DBA_HIVE_TA	ABLES	4-24
	4.1.4.6 DBA_HIVE_CC	OLUMNS	4-25
	4.1.4.7 USER_HIVE_D	DATABASES	4-25
	4.1.4.8 USER_HIVE_T	ΓABLES	4-25
	4.1.4.9 USER_HIVE_C	COLUMNS	4-25
Part	Appendices		
C	Using mtactl to Mana	age the MTA extproc	
F	Change History for P	Previous Releases	
	F.1 Changes in Oracle Big	Data SQL 3.1	F-1
	F.2 Changes in Oracle Big	Data SQL 3.0.1	F-3
А	Manual Steps for Usi	ing Copy to Hadoop for Staged Copies	
	A.1 Generating the Data P	rump Files	A-1
	A.1.1 About Data Pum	p Format Files	A-1
	A.1.2 Identifying the Ta	arget Directory	A-2
	A.1.3 About the CREA	TE TABLE Syntax	A-2
	A.2 Copying the Files to HI	DFS	A-3

4.1.3.3.1 Default Parameter Settings for ORACLE\_HIVE



4-10

A.5	Creating a five rable	A-3
Α.3	3.1 About Hive External Tables	A-3
A.4	Example Using the Sample Schemas	A-4
A.4	4.1 About the Sample Data	A-4
Α.4	1.2 Creating the EXPDIR Database Directory	A-4
A.4	1.3 Creating Data Pump Format Files for Customer Data	A-5
	A.4.3.1 CREATE TABLE Example With a Simple SELECT Statement	A-5
	A.4.3.2 CREATE TABLE Example With a More Complex SQL SELECT Statement	A-5
Α.4	1.4 Verifying the Contents of the Data Files	A-6
Α.4	1.5 Copying the Files into Hadoop	A-6
Α.4	1.6 Creating a Hive External Table	A-6
Usin	g Copy to Hadoop With Direct Copy	
B.1	Manual Steps for Using Copy to Hadoop for Direct Copies	B-1
B.2	Copy to Hadoop Property Reference	B-4
Diag	nostic Tips and Details	
D.1	Running Diagnostics with bdschecksw	D-1
D.2	How to do a Quick Test	D-4
D.3	Oracle Big Data SQL Database Objects	D-5
D.4	Other Database-Side Artifacts	D-6
D.5	Hadoop Datanode Artifacts	D-11
D.6	Step-by-Step Process for Querying an External Table	D-12
D.7	Step-by-Step for a Hive Data Dictionary Query	D-15
D.8	Key Adminstration Tasks for Oracle Big Data SQL	D-16
D.9	Additional Java Diagnostics	D-18
D.10	Checking for Correct Oracle Big Data SQL Patches	D-18
Licer	nsing Information	
E.1	Oracle Big Data SQL Licensing	E-1
E.:	L.1 ANTLR 4.7	E-1
E.:	L.2 JLine 2.13	E-2
E.:	L.3 Apache Commons Exec 1.3	E-2
E.:	L.4 Apache Licensed Code	E-3
E.:	L.5 Apache License	E-3





## **Preface**

The *Oracle Big Data SQL User's Guide* describes how to use and manage the Oracle Big Data SQL product.

## **Audience**

This guide is intended for administrators and users of Oracle Big Data SQL, including:

- Application developers
- Data analysts
- Data scientists
- Database administrators
- System administrators

The guide assumes that the reader has sufficient background knowledge about the database server and the particular Hadoop platform which will host the software in order to follow the instructions successfully.

## **Related Documents**

See the Oracle Big Data Installation Guide for instructions on installing the product.

The following publications are recommended for information about using the product with the Oracle Big Data Appliance and Oracle Big Data Connectors:

- Oracle Big Data Appliance Owner's Guide
- Oracle Big Data Appliance Software User's Guide
- Oracle Big Data Connectors User's Guide

You can find more information about Oracle's Big Data solutions at the Oracle Help Center

Users installing Oracle Big Data SQL on Hortonworks HDP should refer to the Hortonworks documentation site at <a href="http://docs.hortonworks.com/index.html">http://docs.hortonworks.com/index.html</a> for supplementary information

## Conventions

The following text conventions are used in this document:



Convention	Meaning
boldface	Boldface type indicates graphical user interface elements associated with an action, or terms defined in text or the glossary.
italic	Italic type indicates book titles, emphasis, or placeholder variables for which you supply particular values.
monospace	Monospace type indicates commands within a paragraph, URLs, code in examples, text that appears on the screen, or text that you enter.
# prompt	The pound (#) prompt indicates a command that is run as the Linux root user.

## **Backus-Naur Form Syntax**

The syntax in this reference is presented in a simple variation of Backus-Naur Form (BNF) that uses the following symbols and conventions:

Symbol or Convention	Description
[]	Brackets enclose optional items.
{}	Braces enclose a choice of items, only one of which is required.
	A vertical bar separates alternatives within brackets or braces.
	Ellipses indicate that the preceding syntactic element can be repeated.
delimiters	Delimiters other than brackets, braces, and vertical bars must be entered as shown.
boldface	Words appearing in boldface are keywords. They must be typed as shown. (Keywords are case-sensitive in some, but not all, operating systems.) Words that are not in boldface are placeholders for which you must substitute a name or value.

## Changes in Oracle Big Data SQL 3.2

Oracle Big Data SQL Release 3.2 includes major improvements in performance, secure network connectivity, authentication, and user administration, as well as installation and configuration.

#### **JSON CLOB Predicate Pushdown**

Much improved filtering and parsing of JSON CLOB data in Hadoop enables Oracle Big Data SQL to push more processing for these large objects down to the Hadoop cluster. JSON Data can now be filtered on the Oracle Big Data SQL cells in Hadoop for CLOB columns up to 1 MB, depending on character set of the input document. The eligible JSON filter expressions for storage layer evaluation include simplified syntax, JSON\_VALUE, and JSON\_QUERY. In addition, Oracle Big Data SQL can project up to 32 KB of CLOB data from select list expression evaluation in Hadoop to Oracle Database. Processing falls back to Oracle Database only when column sizes exceed these two values.

Customers can disable or re-enable this functionality to suit their own needs.



In Release 3.2, this enhancement currently applies only to JSON expressions returning CLOB data. The same support will be provided for other CLOB types (such as substr and instr) as well as for BLOB data in a future release.

#### Note:

The new JSON CLOB predicate pushdown functionality requires Oracle Database version 12.1.0.2.180417 or greater, as well as the following patches:

- The April 2018 Proactive DBBP (Database Bundle Patch). This is patch 27486326.
- The one-off patch 27767148.

Install the one-off patch on all database compute nodes.

The one-off patch 26170659, which is required on top of earlier DBBPs, is not required on top of the April DBBP.

This functionality is not available through the January 2018 and August 2017 Proactive DBBPs

See the Oracle Big Data SQL Master Compatibility Matrix (Doc ID 2119369.1 in My Oracle Support) for the most up-to-date information on software version and patch requirements.

#### **Support for Querying Kafka Topics**

Release 3.2 provides Hive and Oracle Big Data SQL the ability to query Kafka topics via a new Hive storage handler. You can use this storage handler to create external Hive tables backed by data residing in Kafka. Oracle Big Data SQL or Hive can then query the Kafka data through the external tables. The Kafka key, value, offset, topic name, and partition id are mapped to Hive columns. You can explicitly designate the offset for each topic/partition pair, otherwise the offset will start from the earliest offset in the topic and end with the latest offset in the topic for each partition.

#### **Improved Processing of Parquet Files**

Oracle has introduced its own Parquet reader for processing data in Parquet format. This new reader provides significant performance and resource utilization improvements over the existing Hive Parquet driver. These include:

- More intelligent column processing retrieval. The reader uses "lazy materialization" to process only columns with rows that satisfy the filter, thereby improving I/O.
- Leveraging of dictionaries during filter predicate processing to improve CPU usage.
- Streamlined data conversion, which also contributes to more efficient CPU usage.

The Big Data SQL installation enables the Oracle's Parquet reader by default. You have the option to disable it and revert to the generic Parquet reader.

#### **Multi-User Authorization**

In previous releases of Oracle Big Data SQL, all queries against Hadoop and Hive data are executed as the oracle user and there is no option to change users. Although



oracle is still the underlying user in all cases, Oracle Big Data SQL 3.2 now uses Hadoop Secure Impersonation to direct the oracle account to execute tasks on behalf of other designated users. This enables HDFS data access based on the user that is currently executing the guery, rather than the singular oracle user.

Administrators set up the rules for identifying the query user. They can provide rules for identifying the currently connected user and mapping the connected user to the user that is impersonated. Because there are numerous ways in which users can connect to Oracle Database, this user may be a database user, a user sourced from LDAP, from Kerberos, or a user from other sources. Authorization rules on the files apply for that user and HDFS auditing identifies the actual user running the query.



Administration for Multi-User Authorization is done through the DBMS\_BDSQL PL/SQL Package, which is documented in the *Oracle Big Data SQL User's Guide*.

#### Authentication Between Oracle Database and Oracle Big Data SQL Cells

This authentication is between Oracle Database and the Big Data SQL cells on the Hadoop cluster, facilitating secure communication. The Database Authentication enhancement provides a safeguard against impersonation attacks, in which a rogue service attempts to connect to the Oracle Big Data offload server process running on a cluster node.

#### **Kerberos Ticket Renewal Automation**

On a Kerberos-secured network you can configure the installation to set up automated Kerberos ticket renewal for the oracle account used by Oracle Big Data SQL. This is done for both the Hadoop cluster and Oracle Database sides of the installation. You must provide the principal name and the path to the keytab file.in the bds-config.json configuration file. A template is provided in the configuration file:

```
"kerberos" : {
"principal" : "oracle/mycluster@MY.DOMAIN.COM",
"keytab" : "/home/oracle/security/oracle.keytab"
}
```

If you provide the Kerberos parameters in the configuration file, then Oracle Big Data SQL installation sets up cron jobs on both the Hadoop cluster and Oracle Database servers. These jobs renew the Kerboeros tickets for the principal once per day.

The principal and keytab file must already exist.

#### **Automatic Upgrade**

The current release can now be installed over an earlier release with no need to remove the older software on either the Hadoop or Oracle Database side. The previous installation is upgraded to the current release level.

#### **Common Installation Bundle for all Platforms**

In previous releases, customers needed to unpack the Oracle Big Data SQL installation bundle and choose the correct package for their Hadoop system (CDH or



HDP). Now the bundle contains a single installation package that works for all supported Hadoop systems.

#### Simpler and Faster Installation with the new "Jaguar" Installer

The Jaguar installer replaces  $\mathtt{setup-bds.sh}$ , the installer in previous releases. Jaguar includes these changes:

#### Automatic Check for Installation Prerequisites on Hadoop Nodes

Jaguar checks for installation readiness on each Hadoop DataNode and reports any missing prerequisites.

#### No Need to Manually Generate the Database-Side Installation Bundle

The database-side installation bundle that previously was manually generated by the customer can now be generated automatically. You still need to copy the bundle to the Oracle Database nodes and install it.

#### Faster Overall Installation Time on the Hadoop Side

Installation time will vary, but on the Hadoop Side the installation may take approximately eight minutes if all resources are local, possibly 20 minutes if Hadoop clients must be downloaded from the Internet, depending on download speed.

#### Prerequisite Apache Services on CDH can now be Installed as Either Packages or Parcels

Previously on CDH systems, the Oracle Big Data SQL installation required that the HDFS, YARN, and HIVE components had been installed as parcels. These components can now be installed on CDH as either packages or parcels. There is no change for HDP, where they must be installed as stacks.



On CDH systems, if the Hadooop services required by Oracle Big Data SQL are installed as packages, be sure that they are installed from within Cloudera Manager. Otherwise, Cloudera Manager will not be able to manage these services. This is not an issue with parcels.

#### In the CLI, the Jaguar utility Replaces ./setup-bds

The Jaguar utility is now the primary tool for Hadoop-side installation, deinstallation, and configuration changes, as in these examples:

```
# ./jaguar install bds-config.json
# ./jaguar reconfigure bds-config.json
# ./jaguar uninstall bds-config.json
```

#### The Default Configuration File Name is bds-config.json, but Alternate File Names are Also Accepted

You can now drop the explicit bds-config.json argument and allow the installer default to bds-config.json, as in the first example below. You can also specify an alternate configuration file of any name, though it must adhere to the same internal format as bds-config.json and should be given the .json file type.

```
# ./jaguar install
# ./jaguar install cluster2-config.json
```



You can create configurations files with settings that are tailored to the requirements of each cluster. For example, you may want to apply different security parameters to Oracle Big Data SQL installations on test and production clusters.

#### Configuration Parameters Have Changed Significantly

Users of previous releases will see that the Jaguar configuration file includes a number of new parameters. Most of them are "optional" in the sense that they are not uniformly required, although your particular installation may require some of them. See the Related Links section below for links to the table of installer parameters as well as an example of a configuration file that uses all available parameters.

#### New updatenodes Command for Easier Maintenance

Oracle Big Data SQL must be installed on each Hadoop cluster node that is provisioned with the DataNode role. It has no function on nodes where DataNode is not present. The new Jaguar utility includes the <code>updatenodes</code> command which scans the cluster for instances of the DataNode within the cluster. If the DataNode role has been removed or relocated, or if nodes provisioned with the DataNode have been added or removed, then the script installs or uninstalls Oracle Big Data SQL components from nodes as needed.

#### An Extra Installation Step is Required to Enable Some Security Features

If you choose to enable Database Authentication between Oracle Database and Oracle Big Data SQL cells in the Hadoop cluster, or, Hadoop Secure Impersonation, then an additional "Database Acknowledge" step is required. In this process, the installation on the database server generates a ZIP file of configuration information that you must copy back to the Hadoop cluster management server for processing.

## • On the Database Side, Connections to Clusters are no Longer Classified as Primary and Secondary.

An Oracle Database system can have Oracle Big Data SQL connections to multiple Hadoop clusters. In previous releases, the first these connections was considered the primary (and had to be uninstalled last) and the others were secondary. In the current release, management of multiple installation is simpler and --uninstall-as-primary and --uninstall-as-secondary parameters of the database-side installer are obsolete. However there is now a default cluster. The *Important Terms and Concepts* section of this guide explains the significance of the default cluster.

## Support for Oracle Tablespaces in HDFS Extended to Include All Non-System Permanent Tablespaces

Previous releases supported the move of permanent online tablespaces only to HDFS. This functionality now supports online, read-only, as well as offline permanent tablespaces.

#### Important Change in Behavior of the "mtactl start" Command

Oracle Big Data SQL 3.1 introduced the option to install Oracle Big Data SQL on servers where Oracle Grid Infrastructure is not present. In these environments, you can use the start subcommand of the mtactl utility (mtactl start) to start the MTA (Multi-Threaded Agent) extproc.



Note that in the current release, the mtactl start command works differently from the original Release 3.1 implementation.

- Current behavior: mtactl start starts an MTA extproc using the init parameter values that are stored in the repository. It uses the default values only if the repository does not exist.
- Previous behavior (Oracle Big Data SQL 3.1): mtactl start always uses the
  default init parameters regardless of whether or not init parameter values are
  stored in the repository.

See Also:

#### **Resource Requirements**

8 CPU cores and 12 GB of RAM are now recommended for each node of the Hadoop cluster.

There are some sysctl settings related to kernel, swap, core memory, and socket buffer size that are strongly recommended for optimal performance. These are part of the installation prerequisites explained in Chapter 2 of the installation guide.



1

## Introducing Oracle Big Data SQL

## 1.1 What Is Oracle Big Data SQL?

Oracle Big Data SQL supports queries against non-relational data stored in multiple big data sources, including Apache Hive, HDFS, Oracle NoSQL Database, Apache HBase, and other NoSQL databases. It enables unified query for distributed data and therefore the ability to view and analyze data from disparate data stores seamlessly, as if it were all stored in an Oracle database.

Oracle Big Data SQL enables you to execute highly complex SQL SELECT statements against data in the Hadoop ecosystem, either manually or through your existing applications. For example, if you are a user of Oracle Advanced Analytics, Oracle Big Data SQL enables you to extend your Oracle Database data mining models to big data in Hadoop.

#### Components of an Oracle Big Data SQL Installation

The Oracle Big Data SQL architecture consists of an installation on an Oracle Database system (single node or RAC) that works in conjunction with a parallel installation on a Hadoop (or NoSQL) cluster. The two systems may be networked via either Ethernet or InfiniBand. Hadoop and Hive clients on the compute nodes of the Oracle Database system enable communication between the database and the Oracle Big Data SQL process (known a Oracle Big Data SQL "cell") that runs on each of the DataNodes of the Hadoop cluster. Through this mechanism, Oracle Database can query data on the Hadoop cluster.

Since data in the Hadoop HDFS file system is stored in an undetermined format, SQL queries require some constructs to parse and interpret data for it to be processed in rows and columns. Oracle Big Data SQL leverages available Hadoop constructs to accomplish this, notably InputFormat and SerDe Java classes, optionally through Hive metadata definitions. The Oracle Big Data SQL processing cells on the DataNodes are a layer on top of this generic Hadoop infrastructure. Two key features provided by the cells are *Smart Scan* and *Storage Indexes*, which are described in this chapter.

#### See Also:

The Oracle Big Data SQL Installation Guide describes how to install and configure the software on the two sides of an Oracle Big Data SQL configuration.

The following sections of this guide provide details on using Oracle Big Data SQL:

- About Oracle External Tables
- About the Access Drivers for Oracle Big Data SQL



- About Smart Scan for HDFS
- About Storage Indexes
- About Predicate Push Down
- About Oracle Big Data SQL Statistics

#### 1.1.1 About Oracle External Tables

Oracle Big Data SQL provides external tables with next generation performance gains. An **external table** is an Oracle Database object that identifies and describes the location of data outside of a database. You can query an external table using the same SQL SELECT syntax that you use for any other database tables.

External tables use **access drivers** to parse the data outside the database. Each type of external data requires a unique access driver. Oracle Big Data SQL includes two access drivers for big data: one for data that has metadata defined in Apache Hive, and the other for accessing data stored in the Hadoop Distributed File System, with metadata specified only by an Oracle administrator.

## 1.1.2 About the Access Drivers for Oracle Big Data SQL

By querying external tables, you can access data stored in HDFS and Hive tables as if that data was stored in tables in an Oracle database. Oracle Database accesses the data by using the metadata provided when the external table was created.

Oracle Database supports these access drivers for Oracle Big Data SQL:

- ORACLE\_HIVE: Enables you to create Oracle external tables over Apache Hive data sources. Use this access driver when you already have Hive tables defined for your HDFS data sources. ORACLE\_HIVE can also access data stored in other locations, such as HBase, that have Hive tables defined for them.
- ORACLE\_HDFS: Enables you to create Oracle external tables directly over files stored
  in HDFS. This access driver uses Hive syntax to describe a data source, assigning
  default column names of COL\_1, COL\_2, and so forth. You do not need to create a
  Hive table manually as a separate step.

Instead of acquiring the metadata from a Hive metadata store the way that <code>ORACLE\_HIVE</code> does, the <code>ORACLE\_HDFS</code> access driver acquires all of the necessary information from the access parameters. The <code>ORACLE\_HDFS</code> access parameters are required to specify the metadata, and are stored as part of the external table definition in Oracle Database.

Oracle Big Data SQL uses these access drivers to optimize query performance.

## 1.1.3 About Smart Scan for HDFS

Oracle external tables do not have traditional indexes. Queries against these external tables typically require a full table scan. The Oracle Big Data SQL processing agent on the DataNodes of the Hadoop cluster extends Smart Scan capabilities (such as filter-predicate off-loads) to Oracle external tables. Smart Scan has been used for some time on the Oracle Exadata Database Machine to do column and predicate filtering in the Storage Layer before query results are sent back to the Database Layer. In Oracle Big Data SQL, Smart Scan is a final filtering pass done locally on the Hadoop server to ensure that only requested elements are sent to Oracle Database. Oracle storage



servers running on the Hadoop DataNodes are capable of doing Smart Scans against various data formats in HDFS, such as CSV text, Avro, and Parquet.

This implementation of Smart Scan leverages the massively parallel processing power of the Hadoop cluster to filter data at its source. It can preemptively discard a huge portion of irrelevant data—up to 99 percent of the total. This has several benefits:

- Greatly reduces data movement and network traffic between the cluster and the database
- Returns much smaller result sets to the Oracle Database server.

Query results are returned significantly faster. This is the direct result reduced traffic on the network and reduced load on Oracle Database.



See Storing Oracle Tablespaces in HDFS for instructions on how to set up data files for smart scanning.

See *Oracle Database Concepts* for a general introduction to external tables and pointers to more detailed information in the Oracle Database documentation library

## 1.1.4 About Storage Indexes

Oracle Big Data SQL maintains Storage Indexes automatically, which is transparent to Oracle Database. Storage Indexes contain the summary of data distribution on a hard disk for the data that is stored in HDFS. Storage Indexes reduce the I/O operations cost and the CPU cost of converting data from flat files to Oracle Database blocks. You can think of a storage index as a "negative index". It tells Smart Scan that data does not fall within a block of data, which enables Smart Scan to skip reading that block. This can lead to significant I/O avoidance.

Storage Indexes can be used only for the external tables that are based on HDFS and are created using either the ORACLE\_HDFS driver or the ORACLE\_HIVE driver. Storage Indexes cannot be used for the external tables that use StorageHandlers, such as Apache HBase and Oracle NoSQL.

A Storage Index is a collection of in-memory region indexes, and each region index stores summaries for up to 32 columns. There is one region index for each split. The content stored in one region index is independent of the other region indexes. This makes them highly scalable, and avoids latch contention.

Storage Indexes maintain the minimum and maximum values of the columns of a region for each region index. The minimum and maximum values are used to eliminate unnecessary I/O, also known as I/O filtering. The cell XT granule I/O bytes saved by the Storage Indexes statistic, available in the V\$SYSSTAT view, shows the number of bytes of I/O saved using Storage Indexes.



See Also:

Oracle® Database Reference for information about v\$sysstat view

Queries using the following comparisons are improved by the Storage Indexes:

- Equality (=)
- Inequality (<, !=, or >)
- Less than or equal (<=)</li>
- Greater than or equal (>=)
- IS NULL
- IS NOT NULL

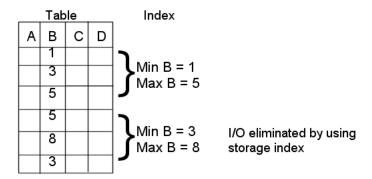
Storage Indexes are built automatically after Oracle Big Data SQL service receives a query with a comparison predicate that is greater than the maximum or less than the minimum value for the column in a region.



- The effectiveness of Storage Indexes can be improved by ordering the rows in a table based on the columns that frequently appear in the WHERE query clause.
- Storage Indexes work with any non-linguistic data type, and works with linguistic data types similar to non-linguistic index.

#### Example 1-1 Elimination of Disk I/O with Storage Indexes

The following figure shows a table and region indexes. The values in **column B** in the table range from 1 to 8. One region index stores the minimum 1, and the maximum of 5. The other region index stores the minimum of 3, and the maximum of 8.



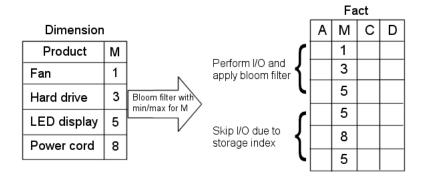
For a query such as the one below, only the first set of rows match. Disk I/O is eliminated because the minimum and maximum of the second set of rows do not match the WHERE clause of the query.

```
SELECT *
FROM TABLE
WHERE B < 2;
```

#### **Example 1-2** Improved Join Performance Using Storage Indexes

Using Storage Indexes allows table joins to skip unnecessary I/O operations. For example, the following query would perform an I/O operation and apply a Bloom filter to only the first block of the fact table. Bloom filters are the key to improved join performance. In the example, a predicate is on the dimension table - not the fact table. The Bloom Filter is created based on "dim.name=Hard drive" and this filter is then applied to the fact table. Therefore, even though the filter is on the dimension table, you are able to filter the data at its source (i.e. Hadoop) based on the results of the dimension query. This also enables optimizations like Storage Indexes to engage.

```
SELECT count(*)
FROM fact, dimension dim
WHERE fact.m=dim.m and dim.product="Hard drive";
```



The I/O for the second block of the fact table is completely eliminated by Storage Indexes as its minimum/maximum range (5,8) is not present in the Bloom filter.

### 1.1.5 About Predicate Push Down

Many Big Data systems support some level of predicate off-loading, either through the filetype itself (e.g. Apache Parquet), or through Hive's partitioning and StorageHandler APIs. Oracle Big Data SQL takes advantage of these off-load capabilities by pushing predicates from the Oracle Database into supporting systems. For example, predicate push down enables the following automatic behaviors:

- Queries against partitioned Hive tables are pruned, based on filter predicates on partition columns.
- Queries against Apache Parquet and Apache ORC files reduce I/O by testing predicates against the internal index-like structures contained within these file formats.



Predicate pushdown in queries against Parquet files is inefficient unless the files are generated through Hive using the workaround described in the next section.  Queries against Oracle NoSQL Database or Apache HBase use SARGable predicates to drive subscans of data in the remote data store.

#### Required Datatypes to Enable Predicate Push Down

Predicate push down requires that certain mappings between Hive Datatypes and Oracle Datatypes be present. These mappings are described in the following table.

Hive Datatype	Mapped To Oracle Datatype
CHAR(m)	CHAR(n), VARCHAR2(n) where n is >= m
VARCHAR(m)	CHAR(n), VARCHAR2(n) where n is $\geq$ m.
string	CHAR(n), VARCHAR2(n)
DATE	DATE
TIMESTAMP	TIMESTAMP(9) Hive TIMESTAMP has nanoseconds, 9 digit fractional seconds.
TINYINT	NUMBER(3) preferably, but NUMBER or NUMBER(n) for any value of n is valid.
SMALLINT	NUMBER(5) preferably, but NUMBER or NUMBER(n) for any value of n is valid.
INT	NUMBER(10) preferably, but NUMBER or NUMBER(n) for any value of n is valid.
BIGINT	NUMBER(19) preferably, but NUMBER or NUMBER(n) for any value of n is OK
DECIMAL(m)	NUMBER(n) where m = n preferably, but NUMBER or NUMBER(n) for any value of n is valid.
FLOAT	BINARY_FLOAT
DOUBLE	BINARY_DOUBLE
BINARY	RAW(n)
BOOLEAN	CHAR(n), VARCHAR2(n) where n is >= 5, values 'TRUE', 'FALSE'
BOOLEAN	NUMBER(1) preferably, but NUMBER or NUMBER(n) for any value of n is valid. Values 0 (false), 1 (true).

## 1.1.6 Predicate Pushdown Issue with Parquet Files Created From Hive Source

Parquet maintains  $\min/\max$  values for each column per rowgroup. However, there is a difference between the Hive/Parquet interpretation and the Oracle interpretation of empty strings that can make predicate pushdown inefficient. There is an easy solution, which is provided at the end of the problem description.

#### **Problem Description**

A Hive VARCHAR or Hive STRING type containing an empty string (a zero-length string, such as '') is stored in Parquet as a zero-length string and is considered a real value (not NULL).



Also, when Hive CHAR fields containing an empty string are stored in Parquet, these are padded with blanks to the full length of the string.

Data Type	Stored in Parquet	When Mapped From Parquet to Oracle Database VARCHAR2		
Hive VARCHAR and Hive String	<ul> <li>Real values, such as `Bob'.</li> <li>Empty, zero-length strings, which are considered a value per the ANSI VARCHAR definition. These are not considered NULL.</li> <li>NULL.</li> </ul>	are considered NULL, not real		
Hive CHAR	<ul> <li>Real values such as: 'Bob'. (Note that the CHAR length is always padded to the defined length. For example, in the case of CHAR(4) 'Bob' is stored as 'Bob'.)</li> </ul>	• Real values such as: `Bob'. (Note that the string length is not padded with trailing blanks. For example, VARCHAR2(4) `Bob' is stored as `Bob'.		
	<ul> <li>Empty, zero-length strings are blank-padded to the length of the CHAR. For example, empty CHAR(2) is store as ` `.</li> <li>NULL.</li> </ul>	NULL. (Empty, zero-length strings are also considered NULL.)		

Oracle VARCHAR2 regards an empty string (a zero-length string) as NULL and, of course, NULL as NULL. Only a non-zero-length value is considered non-NULL. In Hive and Parquet, the empty string is not NULL. Hive and Parquet store NULL as a separate value for a column. This can result in a discrepancy in Storage Index evaluation. For example, Storage Index may compute the following:

```
min='Ohio', max='Washington', number of nulls = 23
```

Parquet statistics for the same data may be:

```
min=", max='Washington', number of nulls = 0
```

In the Parquet statistics above, NULL = 0 because the data loaded into the file included 23 empty strings. These were stored as a zero-length values – actual values, not NULLs.

The problem for predicate pushdown is that when  $\min$ ", empty strings are stored as actual values, this prevents any filtering beyond the bottom end of the min/max range. Therefore, predicate pushdown is much less efficient. It also means that we cannot directly turn Parquet statistics into Storage Index statistics.

#### **Solution**

There are various ways to adjust Parquet files created from Hive source so that they work well with predicate pushdown. The easiest is to go back to the Hive source, redefine the empty string as NULL, and recreate the Parquet file. For example:

- 1. Create a Hive textfile (source) table with:
  - ROW FORMAT DELIMITED
  - FIELDS TERMINATED BY '|'
  - NULL DEFINED AS '' (Two single-quotes in a row, no space in between.)
  - STORED AS textfile



2. Use Hive to do CREATE TABLE AS SELECT from the above textfile to recreate the Parquet file.

```
CREATE TABLE parquet_str_fix
STORED AS PARQUET
LOCATION '<path...>'
AS SELECT * FROM hive str_fix txt;
```

### 1.1.6.1 Other Issues With Parquet Statistics

#### In Parquet File Creation, Statistics are Optional

Because Parquet statistics are optional, we cannot always assume they are present. Some tools provide an options to create Parquet files with or without statistics. Some (like Impala) do not create statistics at all. Therefore, Parquet files created by Impala (as well as by other tools when they do not include statistics) perform poorly when accessed by Hive and Oracle Big Data SQL. Both Hive and Oracle Big Data SQL use Parquet statistics for optimization.

#### Not All Datatypes Are Included in Parquet min/max Column Statistics

Although min/max values for critical datatypes are present (INT, FLOAT, DOUBLE, string, VARCHAR and CHAR), min/max values for the TIMESTAMP datatype are not included in column statistics.

## 1.1.7 About Pushdown of Character Large Object (CLOB) Processing

Queries against Hadoop data may involve processing large objects with potentially millions of records. It is inefficient to return these objects to Oracle Database for filtering and parsing. Oracle Big Data SQL can provide significant performance gains by pushing CLOB processing down to its own processing cells on the Hadoop cluster. Filtering in Hadoop reduces the number of rows returned to Oracle Database. Parsing reduces the amount of data returned from a column within each filtered row.

Customers can disable or re-enable CLOB processing pushdown to suit their own needs.

In Release 3.2, this functionality currently applies only to JSON expressions returning CLOB data. The eligible JSON filter expressions for storage layer evaluation include simplified syntax, JSON VALUE, and JSON QUERY.

The same support will be provided for other CLOB types (such as substr and instr) as well as for BLOB data in a future release.

Oracle Big Data SQL can push processing down to Hadoop for CLOBs within these size constraints:

- Filtering for CLOB columns up to 1 MB in size.
  - The actual amount of data that can be consumed for evaluation in the storage server may vary, depending upon the character set used.
- Parsing for columns up to 32 KB.

This limit refers to the select list projection from storage for the CLOB datatype.

Processing falls back to the Oracle Database only when column sizes exceed these two values.

#### Note:

The new JSON CLOB predicate pushdown functionality requires Oracle Database version 12.1.0.2.180417 or greater, as well as the following patches:

- The April 2018 Proactive DBBP (Database Bundle Patch). This is patch 27486326.
- The one-off patch 27767148.
  Install the one-off patch on all database compute nodes.

The one-off patch 26170659, which is required on top the earlier DBBPs, is not required on top of the April DBBP.

This functionality is not available through the earlier January 2018 and August 2017 Proactive DBBPs

See the Oracle Big Data SQL Master Compatibility Matrix (Doc ID 2119369.1 in My Oracle Support) for the most up-to-date information on software version and patch requirements.

#### **Example 1-3 JSON Document Processing**

For queries into large JSON documents, pushdown of CLOB processing to Oracle Big Data SQL processing cells in Hadoop can be highly effective. Consider the following example, where purchase orders information is stored in JSON. Assume that this record could be up to 25K in size and several millions of such records must processed.

```
{"ponumber":9764, "reference": "LSMITH-20141017", "requestor": "Lindsey Smith", "email": "Lindsey@myco.com", "company": "myco" ...}
```

You can create the external table to access this data as follows. Notice there is a single CLOB column.

```
CREATE TABLE POS_DATA

( pos_info CLOB )

ORGANIZATION EXTERNAL

( TYPE ORACLE_HDFS

DEFAULT DIRECTORY DEFAULT_DIR

LOCATION ('/data/pos/*')
)

REJECT LIMIT UNLIMITED;
```

You can then query the data with this simple syntax:

```
SELECT p.pos_info.email, p.pos_info.requestor
FROM POS_DATA p
WHERE p.pos_info.company='myco'
```

The guery example above engages two data elimination optimizations:

 The data is filtered by the Oracle Big Data SQL cells in the Hadoop cluster. Only records pertaining to the company "myco" are parsed (and after parsing only selected data from these records is returned to the database).



• The Oracle Big Data SQL cells in the cluster parse the filtered set of records and from each record only the values for the two attributes requested (p.pos\_info.email and p.pos\_info.requestor) are returned to the database.

The table below shows some other examples where CLOB processing pushdown is supported. Remember that projections (references on the select side of the CLOB column) are limited to 32 KB of CLOB data, while predicate pushdown is limited to 1 MB of CLOB data.

Query	Comment
SELECT count(*) FROM pos_data p WHERE pos_info is json;	In this case, the predicate ensures that only columns which comply with JSON format are returned.
SELECT pos_info FROM pos_data p WHERE pos_info is json;	The same predicate as in the previous case, but now the CLOB value is projected.
<pre>SELECT json_value(pos_info, '\$.reference') FROM pos_data p WHERE json_value(pos_info, '\$.ponumber') &gt; 9000</pre>	Here, the predicate is issued on a field of the JSON document, and we also execute a JSON value to retrieve field "reference" on top of the projected CLOB JSON value.
<pre>SELECT p.pos_info.reference FROM pos_data p WHERE p.pos_info.ponumber &gt; 9000;</pre>	This is functionally the same query as the previous example, but expressed in simplified syntax.
SELECT p.pos_info.email FROM po_data p WHERE json_exists(pos_info, '\$.requestor') and json_query(pos_info, '\$.requestor') is not null;	This example shows how json_exists and json_query can also be used as predicates.

## 1.1.8 About Oracle Big Data SQL Statistics

Oracle Big Data SQL provides a number of statistics that can contribute data for performance analyses.

#### Five Key Cell XT and Storage Index Statistics

If a query is off-loadable, the following XT-related statistics that can help you to determine what kind of I/O savings you can expect from the offload and from Smart Scan.

cell XT granules requested for predicate offload

Note that number of granules requested depends on a number of a factors, including the HDFS block size, Hadoop data source splittability, and the effectiveness of Hive partition elimination.

cell XT granule bytes requested for predicate offload

The number of bytes requested for the scan. This is the size of the data on Hadoop to be investigated after Hive partition elimination and before Storage Index evaluation.

cell interconnect bytes returned by XT smart scan

The number of bytes of I/O returned by an XT smart scan to Oracle Database.



#### cell XT granule predicate offload retries

The number of times that a Big Data SQL process running on a DataNode could not complete the requested action. Oracle Big Data SQL automatically retries failed requests on other DataNodes that have a replica of the data. The retries value should be zero.

#### cell XT granule IO bytes saved by storage index

The number of bytes filtered out by storage indexes at the storage cell level. This is data that was not scanned, based information provided by the storage indexes.

You can check these statistics before and after running queries as follows. This example shows the values at null, before running any queries.

```
SQL> SELECT sn.name,ms.value
FROM V$MYSTAT ms, V$STATNAME sn
WHERE ms.STATISTIC#=sn.STATISTIC# AND sn.name LIKE '%XT%';
```

NAME	VALUE
cell XT granules requested for predicate offload	0
cell XT granule bytes requested for predicate offload	0
cell interconnect bytes returned by XT smart scan	0
cell XT granule predicate offload retries	0
cell XT granule IO bytes saved by storage index	0

You can check some or all of these statistics after execution of a query to test the effectiveness of the query, as in:

```
SQL> SELECT n.name, round(s.value/1024/1024)
FROM v$mystat s, v$statname n
WHERE s.statistic# IN (462,463)
AND s.statistic# = n.statistic#;
cell XT granule bytes requested for predicate offload 32768
cell interconnect bytes returned by XT smart scan 32
```



#### Tip:

The Oracle Big Data SQL Quickstart blog, published in the Data Warehouse Insider, provides a series of code and functionality walkthroughs that show you how to use these statistics to analyze the performance of Oracle Big Data SQL. See Part 2, Part 7, and Part 10.

## 1.2 Installation

Oracle Big Data SQL requires installation of components on the Hadoop system where the data resides and also on the Oracle Database server which queries the data.

See the following resources for installation information:

- Oracle Big Data SQL Installation Guide
   This guide describes installation and configuration procedures for supported
- Oracle Big Data SQL Master Compatibility Matrix

Hadoop system/Oracle Database server combinations.



This is Document 2119369.1 in My Oracle Support. Check the matrix for up-to-date information on Big Data SQL compatibility with the following:

- Oracle Engineered Systems.
- Other systems.
- Linux OS distributions and versions.
- Hadoop distributions.
- Oracle Database releases, including required patches.



2

## Using Oracle Big Data SQL for Data Access

This chapter describes how to use Oracle Big Data SQL to create external tables and access data from Hadoop data sources as well as Oracle NoSQL Database.

It also describes some of the changes that Oracle Big Data SQL makes on the Oracle Database server.

- Creating an Oracle External Table for Hive Data
- Creating an Oracle External Table for Oracle NoSQL Database
- Creating an Oracle External Table for Apache HBase
- Creating an Oracle External Table for HDFS Files
- About the SQL CREATE TABLE Statement
- About Data Type Conversions
- Querying External Tables
- About Pushdown of Character Large Object (CLOB) Processing
- About Oracle Big Data SQL on the Database Server (Oracle Exadata Machine or Other)

## 2.1 Creating an Oracle External Table for Hive Data

Oracle Big Data SQL enables you to query Hive tables from the Oracle Database using the full power of Oracle SQL SELECT statements. It also enables you to write queries that join Oracle tables and Hive data, leverage robust Oracle Database security features, and take advantage of advanced SQL capabilities like analytic functions, JSON handling, and others.

To enable Oracle Big Data SQL to query Hive data, you must first define an Oracle external table for your Hive data. There are a number of tools available to help you create the Oracle external table definition.

#### DBMS HADOOP

DBMS\_HADOOP is a PL/SQL package that contains the CREATE\_EXTDDL\_FOR\_HIVE procedure. This procedure generates the DDL to create an Oracle external table for a given Hive table. You can optionally edit the text of the generated DDL before execution in order to customize the external table properties.

#### The Big Data SQL wizard in Oracle SQL Developer

The most recent versions of the free Oracle SQL Developer tool include a Big Data SQL wizard that guides you easily through the process of creating Oracle external table definitions.

If you have a configured Hive connection in Oracle SQL Developer, then in the **Connections** navigator, drill down from the connection entry point to a Hive table and do the following:

- Right-click on the table icon and select Use in Oracle Big Data SQL...
- When prompted, select an Oracle Database connection for the import of the Hive table.
- 3. Select an Oracle Big Data SQL-enabled target database.
- 4. In the Create Table dialog, check over the current configuration for columns, external table properties, and storage. Modify as needed. You can also preview the text of the DDL that will be generated.
- 5. Click **OK** when you are satisfied with the table definition. The wizard will create the external table at the designated location.

#### The Oracle SQL Developer Data Modeler

This is free graphical design tool that you can use to connect to a Hive metastore and generate an external table. You can select and import one or multiple Hive tables, modify table properties as needed, and then generate the DDL that you can copy into an SQL Worksheet and then run in order to create an Oracle external table. Although the Data Modeler is a more complex tool to use than the other options, its advantage is that you can use it to work on multiple Hive tables

See Oracle SQL Developer & Data Modeler Support for Oracle Big Data SQL in the Oracle Blog space for a demonstration of how to use the Data Modeler.

#### See Also:

For instructions on how to install Oracle SQL Developer and connect to Hive in order to create external tables, see Using Oracle SQL Developer to Connect to Hive

## 2.1.1 Obtaining Information About a Hive Table

The DBMS\_HADOOP PL/SQL package contains a function named CREATE\_EXTDDL\_FOR\_HIVE. It returns the data dictionary language (DDL) to create an external table for accessing a Hive table. This function requires you to provide basic information about the Hive table:

- Name of the Hadoop cluster
- Name of the Hive database
- Name of the Hive table
- Whether the Hive table is partitioned

You can obtain this information by querying the ALL\_HIVE\_TABLES data dictionary view. It displays information about all Hive tables that you can access from Oracle Database.

This example shows that the current user has access to an unpartitioned Hive table named RATINGS HIVE TABLE in the default database. A user named JDDE is the owner.

```
SQL> SELECT cluster_id, database_name, owner, table_name, partitioned FROM
all_hive_tables;
CLUSTER_ID DATABASE_NAME OWNER TABLE_NAME PARTITIONED
```



```
hadoopl default jdoe ratings_hive_table UN-PARTITIONED
```

See Also:

"Static Data Dictionary Views for Hive"

## 2.1.2 Using the CREATE\_EXTDDL\_FOR\_HIVE Function

With the information from the data dictionary, you can use the <code>CREATE\_EXTDDL\_FOR\_HIVE</code> function of <code>DBMS\_HADOOP</code>. This example specifies a database table name of <code>RATINGS\_DB\_TABLE</code> in the current schema. The function returns the text of the <code>CREATE TABLE</code> command in a local variable named <code>DDLOUT</code>, but does not execute it.

```
DECLARE
   DDLout VARCHAR2(4000);
BEGIN
   dbms_hadoop.create_extddl_for_hive(
        CLUSTER_ID=>'hadoop1',
        DB_NAME=>'default',
        HIVE_TABLE_NAME=>'ratings_hive_table',
        HIVE_PARTITION=>FALSE,
        TABLE_NAME=>'ratings_db_table',
        PERFORM_DDL=>FALSE,
        TEXT_OF_DDL=>DDLout
   );
   dbms_output.put_line(DDLout);
END;
/
```

When this procedure runs, the PUT LINE function displays the CREATE TABLE command:

```
CREATE TABLE ratings_db_table (
    c0 VARCHAR2(4000),
    c1 VARCHAR2(4000),
    c2 VARCHAR2(4000),
    c3 VARCHAR2(4000),
    c4 VARCHAR2(4000),
    c5 VARCHAR2(4000),
    c6 VARCHAR2(4000),
    c7 VARCHAR2(4000))

ORGANIZATION EXTERNAL
    (TYPE ORACLE_HIVE DEFAULT DIRECTORY DEFAULT_DIR
    ACCESS PARAMETERS
     (
        com.oracle.bigdata.cluster=hadoop1
        com.oracle.bigdata.tablename=default.ratings_hive_table
    )
    ) PARALLEL 2 REJECT LIMIT UNLIMITED
```

You can capture this information in a SQL script, and use the access parameters to change the Oracle table name, the column names, and the data types as desired before executing it. You might also use access parameters to specify a date format mask.



The ALL\_HIVE\_COLUMNS view shows how the default column names and data types are derived. This example shows that the Hive column names are C0 to C7, and that the Hive STRING data type maps to VARCHAR2(4000):

SQL> SELECT table\_name, column\_name, hive\_column\_type, oracle\_column\_type FROM all\_hive\_columns;

TABLE_NAME	COLUMN_NAME	HIVE_COLUMN_TYPE	ORACLE_COLUMN_TYPE
ratings_hive_table	c0	string	VARCHAR2(4000)
ratings_hive_table	c1	string	VARCHAR2(4000)
ratings_hive_table	c2	string	VARCHAR2(4000)
ratings_hive_table	c3	string	VARCHAR2(4000)
ratings_hive_table	c4	string	VARCHAR2(4000)
ratings_hive_table	c5	string	VARCHAR2(4000)
ratings_hive_table	c6	string	VARCHAR2(4000)
ratings_hive_table	c7	string	VARCHAR2(4000)

8 rows selected.



"DBMS HADOOP PL/SQL Package"

## 2.1.3 Using Oracle SQL Developer to Connect to Hive

Oracle SQL Developer provides methods to connect to a Hive metastore and create Oracle external tables over Hive.

Follow these steps to set up Oracle SQL Developer to work with Oracle Big Data SQL.

- 1. Install Oracle SQL Developer
- 2. Download the Hive JDBC Drivers
- 3. Add the new Hive JDBC Drivers to Oracle SQL Developer
- 4. Create a database connection to Hive.

#### **Installing Oracle SQL Developer**

Install Oracle SQL Developer 4.2 or greater. Release 4.2 is recommended because it includes support for *Copy To Hadoop*, a useful Oracle Big Data SQL tool for offloading Oracle Database tables to HDFS.

The installation is simple. Just download the package and extract it.

- Go to the Oracle SQL Developer download site on the Oracle Technology Network (OTN).
- 2. Accept the license agreement and download the version that is appropriate for your platform.

For most users, Windows 64-bit with JDK 8 included is the correct choice.

3. Extract the downloaded ZIP file to your local drive.

You can extract to any folder name.



See Installing and Getting Started with SQL Developer in the Oracle SQL Developer User's Guide for further installation and configuration details.

#### Downloading and Installing the Hive JDBC Drivers for Cloudera Enterprise

To connect Oracle SQL Developer to Hive in the Hadoop environment, you need to download and install the Hive JDBC drivers for Cloudera Enterprise. These drivers are not included in the Oracle SQL Developer download package.

#### **Note for HDP Users:**

At this time, SQL Developer 4.2 requires the Cloudera JDBC drivers for Hive. However, these drivers appear to work against Hortonworks clusters as well. HDP users should test to determine if these drivers meet their needs.

Download the latest Cloudera JDBC drivers for Hive from the Cloudera website to any local directory.

You can search for "cloudera hive jdbc drivers download" on the Cloudera website to locate the available driver packages.

You are prompted to select the driver version, OS, and OS version (32/64 bit). At this time, the latest drive version is 2.5.18. You can choose the newest version available.

2. Unzip the archive:

```
unzip hive_jdbc_<version>.zip
```

3. View the extracted content. Notice that under the top-level folder there are multiple ZIP files. Each is for a different JDBC version. For this setup, only JBDC 4.0 is usable. Select the JDBC4\_ ZIP file (JDBC4\_<version>.zip).



#### Important:

Choose only the JDBC4\_ ZIP file, which contains the drivers for JDBC 4.0. This is the only compatible version. The drivers in other packages, such as JDBC41 \*, are not compatible with SOL Developer 4.2 and will return errors upon connection attempts.

Unzip the JDBC4 archive to a target directory that is accessible to Oracle SQL Developer, for example, ./home/oracle/jdbc:

```
# unzip Cloudera_HiveJDBC4_<version>.zip -d /home/oracle/jdbc/
```

#### The extracted content should be similar to this:

```
Cloudera HiveJDBC4_2.5.18.1050\Cloudera-JDBC-Driver-for-Apache-Hive-Install-
Guide.pdf
Cloudera_HiveJDBC4_2.5.18.1050\Cloudera-JDBC-Driver-for-Apache-Hive-Release-
Notes.pdf
Cloudera_HiveJDBC4_2.5.18.1050\commons-codec-1.3.jar
Cloudera_HiveJDBC4_2.5.18.1050\commons-logging-1.1.1.jar
Cloudera_HiveJDBC4_2.5.18.1050\HiveJDBC4.jar
Cloudera_HiveJDBC4_2.5.18.1050\hive_metastore.jar
Cloudera_HiveJDBC4_2.5.18.1050\hive_service.jar
Cloudera_HiveJDBC4_2.5.18.1050\httpclient-4.1.3.jar
```



```
Cloudera_HiveJDBC4_2.5.18.1050\httpcore-4.1.3.jar
Cloudera_HiveJDBC4_2.5.18.1050\libfb303-0.9.0.jar
Cloudera_HiveJDBC4_2.5.18.1050\libthrift-0.9.0.jar
Cloudera_HiveJDBC4_2.5.18.1050\log4j-1.2.14.jar
Cloudera_HiveJDBC4_2.5.18.1050\out.txt
Cloudera_HiveJDBC4_2.5.18.1050\ql.jar
Cloudera_HiveJDBC4_2.5.18.1050\slf4j-api-1.5.11.jar
Cloudera_HiveJDBC4_2.5.18.1050\slf4j-log4j12-1.5.11.jar
Cloudera_HiveJDBC4_2.5.18.1050\rowslf4j-log4j12-1.5.11.jar
Cloudera_HiveJDBC4_2.5.18.1050\rowslf4j-log4j12-1.5.11.jar
Cloudera_HiveJDBC4_2.5.18.1050\rowslf4j-log4j12-1.5.11.jar
Cloudera_HiveJDBC4_2.5.18.1050\rowslf4j-log4j12-1.5.11.jar
```

#### Add the new Hive JDBC Drivers to Oracle SQL Developer

Next, start up SQL Developer and copy all of the extracted driver files into "Third Party JDBC Drivers" in the **Preferences** window.

- Navigate to the folder where you downloaded and extracted Oracle SQL Developer.
- Click the sqldeveloper subfolder. Then, click sqldeveloper.exe in this folder.
- 3. In the SQL Developer menu bar, select Tools>Preferences.
- In the file explorer of the Preferences window, expand Database and then click Third Party JDBC Drivers.
- 5. Click Add Entry.
- 6. Navigate to the folder where you sent the files extracted from Cloudera\_HiveJDBC4\_<version>.zip. Copy all of the JAR files from the ZIP extraction into this window and then click **OK**.
- Restart Oracle SQL Developer.

#### Create a Database Connection to Hive

After the drivers are installed, you can create a connection to Hiveserver2.

If you are creating a Kerberos-secured connection, you will need a user ID, the Kerberos connection parameters, and the number of the port where Hiveserver2 is running on the Hadoop system (typically, port 10000). A keytab must exist for the user.

If you not using Kerberos, you will need a user ID (the oracle user or a user with equivalent privileges), the account password, and the Hiveserver2 port number.

See Create/Edit/Select Database Connection in the Oracle SQL Developer User's Guide for a explanation of the fields in the Oracle and Hive tabs in the **New/Select Database Connection** dialog.

## 2.1.4 Developing a CREATE TABLE Statement for ORACLE HIVE

Whichever method you use to create an Oracle external table over Hive (DBMS\_HADOOP, Oracle SQL Developer Data Modeler, Oracle Big Data Wizard in Oracle SQL Developer, or manual coding), you may need to set some access parameters to modify the default behavior of <code>ORACLE\_HIVE</code>.

## 2.1.4.1 Using the Default ORACLE\_HIVE Settings

The following statement creates an external table named <code>ORDER</code> to access Hive data:



```
CREATE TABLE order (cust_num VARCHAR2(10), order_num VARCHAR2(20), description VARCHAR2(100), order_total NUMBER (8,2))

ORGANIZATION EXTERNAL (TYPE oracle_hive);
```

Because no access parameters are set in the statement, the <code>ORACLE\_HIVE</code> access driver uses the default settings to do the following:

- Connects to the default Hadoop cluster.
- Uses a Hive table named order. An error results if the Hive table does not have fields named CUST\_NUM, ORDER\_NUM, DESCRIPTION, and ORDER\_TOTAL.
- Sets the value of a field to NULL if there is a conversion error, such as a CUST\_NUM value longer than 10 bytes.

## 2.1.4.2 Overriding the Default ORACLE\_HIVE Settings

You can set properties in the ACCESS PARAMETERS clause of the external table clause, which override the default behavior of the access driver. The following clause includes the com.oracle.bigdata.overflow access parameter. When this clause is used in the previous example, it truncates the data for the DESCRIPTION column that is longer than 100 characters, instead of throwing an error:

```
(TYPE oracle_hive
ACCESS PARAMETERS (
   com.oracle.bigdata.overflow={"action:"truncate", "col":"DESCRIPTION""} ))
```

The next example sets most of the available parameters for <code>oracle\_hive</code>:

```
CREATE TABLE order (cust_num VARCHAR2(10),
                    order_num VARCHAR2(20),
                    order_date DATE,
                    item_cnt NUMBER,
                    description VARCHAR2(100),
                    order total (NUMBER(8,2)) ORGANIZATION EXTERNAL
  (TYPE oracle hive
    ACCESS PARAMETERS (
       com.oracle.bigdata.tablename: order_db.order_summary
       com.oracle.bigdata.colmap: {"col":"ITEM_CNT", \
                                        "field": "order_line_item_count" }
        com.oracle.bigdata.overflow: {"action":"TRUNCATE", \
                                        "col": "DESCRIPTION" }
        com.oracle.bigdata.erroropt: [{"action":"replace", \
                                         "value":"INVALID_NUM" , \
                                         "col":["CUST_NUM","ORDER_NUM"]} ,\
                                        {"action": "reject", \
                                         "col": "ORDER_TOTAL}
))
```

The parameters make the following changes in the way that the <code>ORACLE\_HIVE</code> access driver locates the data and handles error conditions:

- com.oracle.bigdata.tablename: Handles differences in table names. ORACLE\_HIVE looks for a Hive table named ORDER\_SUMMARY in the ORDER.DB database.
- com.oracle.bigdata.colmap: Handles differences in column names. The Hive
   ORDER LINE ITEM COUNT field maps to the Oracle ITEM CNT column.



- com.oracle.bigdata.overflow: Truncates string data. Values longer than 100 characters for the DESCRIPTION column are truncated.
- com.oracle.bigdata.erroropt: Replaces bad data. Errors in the data for CUST\_NUM or ORDER\_NUM set the value to INVALID\_NUM.

## 2.2 Creating an Oracle External Table for Oracle NoSQL Database

You can use the <code>ORACLE\_HIVE</code> access driver to access data stored in Oracle NoSQL Database. However, you must first create a Hive external table that accesses the KVStore. Then you can create an external table in Oracle Database over it, similar to the process described in "Creating an Oracle External Table for Hive Data".

This section contains the following topics:

- Creating a Hive External Table for Oracle NoSQL Database
- Creating the Oracle Database Table for Oracle NoSQL Data
- About Column Data Type Mappings
- Example of Accessing Data in Oracle NoSQL Database

## 2.2.1 Creating a Hive External Table for Oracle NoSQL Database

To provide access to the data in Oracle NoSQL Database, you create a Hive external table over the Oracle NoSQL table. Oracle Big Data SQL provides a StorageHandler named oracle.kv.hadoop.hive.table.TableStorageHandler that enables Hive to read the Oracle NoSQL Database table format.

The following is the basic syntax of a Hive CREATE TABLE statement for a Hive external table over an Oracle NoSQL table:

```
CREATE EXTERNAL TABLE tablename colname coltype[, colname coltype,...]
STORED BY 'oracle.kv.hadoop.hive.table.TableStorageHandler'
TBLPROPERTIES (
        "oracle.kv.kvstore" = "database",
        "oracle.kv.hosts" = "nosql_node1:port[, nosql_node2:port...]",
        "oracle.kv.hadoop.hosts" = "hadoop_node1[,hadoop_node2...]",
        "oracle.kv.tableName" = "table_name");
```

#### **Hive CREATE TABLE Parameters**

#### tablename

The name of the Hive external table being created.

This table name will be used in SQL queries issued in Oracle Database, so choose a name that is appropriate for users. The name of the external table that you create in Oracle Database must be identical to the name of this Hive table.

Table, column, and field names are case insensitive in Oracle NoSQL Database, Apache Hive, and Oracle Database.

#### colname coltype

The names and data types of the columns in the Hive external table. See Table 2-1 for the data type mappings between Oracle NoSQL Database and Hive.



#### **Hive CREATE TABLE TBLPROPERTIES Clause**

#### oracle.kv.kvstore

The name of the KVStore. Only upper- and lowercase letters and digits are valid in the name.

#### oracle.kv.hosts

A comma-delimited list of host names and port numbers in the Oracle NoSQL Database cluster. Each string has the format *hostname:port*. Enter multiple names to provide redundancy in the event that a host fails.

#### oracle.kv.hadoop.hosts

A comma-delimited list of all host names in the Hadoop cluster with Oracle Big Data SQL enabled.

#### oracle.kv.tableName

The name of the table in Oracle NoSQL Database that stores the data for this Hive external table.



Apache Hive Language Manual DDL at

https://cwiki.apache.org/confluence/display/Hive/LanguageManual +DDL#LanguageManualDDL-Create/Drop/TruncateTable

## 2.2.2 Creating the Oracle Database Table for Oracle NoSQL Data

Use the following syntax to create an external table in Oracle Database that can access the Oracle NoSQL data through a Hive external table:

```
CREATE TABLE tablename(colname colType[, colname colType...])
ORGANIZATION EXTERNAL
(TYPE ORACLE_HIVE DEFAULT DIRECTORY directory
ACCESS PARAMETERS
(access parameters)
)
REJECT LIMIT UNLIMITED;
```

In this syntax, you identify the column names and data types. For more about this syntax, see "About the SQL CREATE TABLE Statement".

## 2.2.3 About Column Data Type Mappings

When Oracle Big Data SQL retrieves data from Oracle NoSQL Database, the data is converted twice to another data type:

- To a Hive data type when the data is read into the columns of the Hive external table.
- To an Oracle data type when the data is read into the columns of an Oracle Database external table.



. In order to execute a Big Data SQL query against data stored in an Oracle NoSQL Database table, a Hive *external table* must first be created with a schema mapped from the schema of the desired Oracle NoSQL Database table. Table 2-1 identifies the supported data types Oracle NoSQL Database table API and their mappings to Hive.

Table 2-1 Mapping Hive Data Types to the NoSQL Database Table API Data Model

Oracle NoSQL Database Table API	Hive
FieldDef.Type.STRING	STRING
FieldDef.Type.BOOLEAN	BOOLEAN
FieldDef.Type.BINARY	BINARY
FieldDef.Type.FIXED_BINARY	BINARY
FieldDef.Type.INTEGER	INT
FieldDef.Type.LONG	BIGINT
FieldDef.Type.FLOAT	FLOAT
FieldDef.Type.DOUBLE	DOUBLE
FieldDef.Type.ENUM	STRING
FieldDef.Type.ARRAY	ARRAY
FieldDef.Type.MAP	MAP <string, data_type=""></string,>
FieldDef.Type.RECORD	STRUCT <col_name :="" data_type,=""></col_name>

### Note:

To complete this mapping a corresponding Oracle Database *external table* must be created with a schema mapped from the schema of the Hive table.

Also note that the following Hive data types are not applicable to the mapping of Oracle NoSQL data types to Oracle Database data types: VARCHAR, CHAR, TINYINT, SMALLINT, DECIMAL, TIMESTAMP, DATE, UNION TYPE.

### See Also:

About Data Type Conversions provides details on Hive to Oracle Database data type mappings.

Predicate Pushdown in Oracle Big Data SQL requires that certain mappings between Hive Datatypes and Oracle Datatypes be present. See About Predicate Push Down.

# 2.2.4 Example of Accessing Data in Oracle NoSQL Database

This example uses the sample data provided with the Oracle NoSQL Database software:



- Creating the Oracle NoSQL Database Example Table
- Creating the Example Hive Table for vehicleTable
- Creating the Oracle Table for VEHICLES

### 2.2.4.1 Creating the Oracle NoSQL Database Example Table

Verify that the following files reside in the examples/hadoop/table directory:

```
create_vehicle_table.kvs
CountTableRows.java
LoadVehicleTable.java
```

This example runs on a Hadoop cluster node named some1node07 and uses a KVStore named SOME1KV.

#### To create and populate the sample table in Oracle NoSQL Database:

- 1. Open a connection to an Oracle NoSQL Database node on your Hadoop cluster.
- 2. Create a table named vehicleTable. The following example uses the load command to run the commands in create\_vehicle\_table.kvs:

```
$ cd NOSQL_HOME
$ java -jar lib/kvcli.jar -host somelnode07 -port 5000 \
   load -file examples/hadoop/table/create_vehicle_table.kvs
```

3. Compile LoadVehicleTable.java:

```
$ javac -cp examples:lib/kvclient.jar examples/hadoop/table/LoadVehicleTable.java
```

**4.** Execute the LoadVehicleTable class to populate the table:

The vehicleTable table contains the following fields:

Table 2-2 Fields in the vehicleTable Example

Field Name	Data Type
type	STRING
make	STRING
model	STRING
class	STRING
color	STRING
price	DOUBLE
count	INTEGER



### 2.2.4.2 Creating the Example Hive Table for vehicleTable

The following example creates a Hive table named VEHICLES that accesses vehicleTable in the SOME1KV KVStore. In this example, the system is configured with a Hadoop cluster in the first six servers (some1node01 to some1node06) and an Oracle NoSQL Database cluster in the next three servers (some1node07 to some1node09).

```
CREATE EXTERNAL TABLE IF NOT EXISTS vehicles
   (type STRING,
   make STRING,
    model STRING,
   class STRING,
    color STRING,
   price DOUBLE,
   count INT)
COMMENT 'Accesses data in vehicleTable in the SOME1KV KVStore'
STORED BY 'oracle.kv.hadoop.hive.table.TableStorageHandler'
TBLPROPERTIES
  ("oracle.kv.kvstore" = "SOME1KV",
   "oracle.kv.hosts" = "somelnode07.example.com:5000,somelnode08.example.com:5000",
   "oracle.kv.hadoop.hosts" =
"somelnode01.example.com,somelnode02.example.com,somelnode03.example.com,somelnode04.
example.com, somelnode05.example.com, somelnode06.example.com",
   "oracle.kv.tableName" = "vehicleTable");
```

The DESCRIBE command lists the columns in the VEHICLES table:

```
hive> DESCRIBE vehicles;
OK
                                             from deserializer
type
                       string
                                             from deserializer
make
                       string
model
                                             from deserializer
                       string
                                             from deserializer
class
                       string
                                             from deserializer
color
                       string
                       double
                                             from deserializer
price
                                              from deserializer
count
                       int.
```

A query against the Hive VEHICLES table returns data from the Oracle NoSQL vehicleTable table:

```
hive> SELECT make, model, class
     FROM vehicletable
     WHERE type='truck' AND color='red'
     ORDER BY make, model;
Total MapReduce jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 1
Chrysler
             Ram1500
                             RearWheelDrive
Chrysler
            Ram2500
                             FrontWheelDrive
Ford
            F150
                             FrontWheelDrive
Ford
            F250
                             RearWheelDrive
Ford
             F250
                            AllWheelDrive
             F350
                            RearWheelDrive
Ford
                             AllWheelDrive
              Sierra
GM
```



GM	Silverado1500	RearWheelDrive
GM	Silverado1500	AllWheelDrive

# 2.2.4.3 Creating the Oracle Table for VEHICLES

After you create the Hive table, the metadata is available in the Oracle Database static data dictionary views. The following SQL SELECT statement returns information about the Hive table created in the previous topic:

The next SQL CREATE TABLE statement generates an external table named VEHICLES over the Hive VEHICLES table, using the ORACLE\_HIVE access driver. The name of the table in Oracle Database must be identical to the name of the table in Hive. However, both Oracle NoSQL Database and Oracle Database are case insensitive.

```
CREATE TABLE vehicles
(type VARCHAR2(10), make VARCHAR2(12), model VARCHAR2(20),
class VARCHAR2(40), color VARCHAR2(20), price NUMBER(8,2),
count NUMBER)
ORGANIZATION EXTERNAL
(TYPE ORACLE_HIVE DEFAULT DIRECTORY DEFAULT_DIR
ACCESS PARAMETERS
(com.oracle.bigdata.debug=true com.oracle.bigdata.log.opt=normal))
REJECT LIMIT UNLIMITED;
```

This SQL SELECT statement retrieves all rows for red trucks from <code>vehicleTable</code> in Oracle NoSQL Database:

```
SQL> SELECT make, model, class
    FROM vehicles
    WHERE type='truck' AND color='red'
    ORDER BY make, model;
```

MODEL	CLASS
Ram1500	RearWheelDrive
Ram2500	FrontWheelDrive
F150	FrontWheelDrive
F250	AllWheelDrive
F250	RearWheelDrive
F350	RearWheelDrive
Sierra	AllWheelDrive
Silverado1500	RearWheelDrive
Silverado1500	4WheelDrive
Silverado1500	AllWheelDrive
	Ram1500 Ram2500 F150 F250 F250 F350 Sierra Silverado1500 Silverado1500



# 2.3 Creating an Oracle External Table for Apache HBase

You can also use the <code>ORACLE\_HIVE</code> access driver to access data stored in Apache HBase. However, you must first create a Hive external table that accesses the HBase table. Then you can create an external table in Oracle Database over it. The basic steps are the same as those described in "Creating an Oracle External Table for Oracle NoSQL Database".

# 2.3.1 Creating a Hive External Table for HBase

To provide access to the data in an HBase table, you create a Hive external table over it. Apache provides a storage handler and a SerDe that enable Hive to read the HBase table format.

The following is the basic syntax of a Hive CREATE TABLE statement for an external table over an HBase table:

```
CREATE EXTERNAL TABLE tablename colname coltype[, colname coltype,...]

ROW FORMAT

SERDE 'org.apache.hadoop.hive.hbase.HBaseSerDe'

STORED BY 'org.apache.hadoop.hive.hbase.HBaseStorageHandler'

WITH SERDEPROPERTIES (
   'serialization.format'='1',
   'hbase.columns.mapping'=':key,value:key,value:
```

### See Also:

Apache Hive Language Manual DDL at

https://cwiki.apache.org/confluence/display/Hive/LanguageManual +DDL#LanguageManualDDL-Create/Drop/TruncateTable

Hive HBase Integration at

https://cwiki.apache.org/confluence/display/Hive/ HBaseIntegration#HBaseIntegration-StorageHandlers

Class HBaseSerDe in the Hive API reference at

http://hive.apache.org/javadocs/r0.13.1/api/hbase-handler/index.html

# 2.3.2 Creating the Oracle Database Table for HBase

Use the following syntax to create an external table in Oracle Database that can access the HBase data through a Hive external table:

```
CREATE TABLE tablename(colname colType[, colname colType...])
ORGANIZATION EXTERNAL
(TYPE ORACLE_HIVE DEFAULT DIRECTORY DEFAULT_DIR
ACCESS PARAMETERS
(access parameters)
)
REJECT LIMIT UNLIMITED;
```



In this syntax, you identify the column names and data types. To specify the access parameters, see "About the SQL CREATE TABLE Statement".

# 2.4 Creating an Oracle External Table for HDFS Files

The ORACLE\_HDFS access driver enables you to access many types of data that are stored in HDFS, but which do not have Hive metadata. You can define the record format of text data, or you can specify a SerDe for a particular data format.

You must create the external table for HDFS files manually, and provide all the information the access driver needs to locate the data, and parse the records and fields. The following are some examples of CREATE TABLE ORGANIZATION EXTERNAL statements.

# 2.4.1 Using the Default Access Parameters with ORACLE\_HDFS

The following statement creates a table named <code>order</code> to access the data in all files stored in the <code>/usr/cust/summary</code> directory in HDFS:

Because no access parameters are set in the statement, the <code>ORACLE\_HDFS</code> access driver uses the default settings to do the following:

- Connects to the default Hadoop cluster.
- Reads the files as delimited text, and the fields as type STRING.
- Assumes that the number of fields in the HDFS files match the number of columns (three in this example).
- Assumes the fields are in the same order as the columns, so that CUST\_NUM data is
  in the first field, ORDER\_NUM data is in the second field, and ORDER\_TOTAL data is in the
  third field.
- Rejects any records in which the value causes a data conversion error: If the value for CUST\_NUM exceeds 10 characters, the value for ORDER\_NUM exceeds 20 characters, or the value of ORDER\_TOTAL cannot be converted to NUMBER.

# 2.4.2 Overriding the Default ORACLE\_HDFS Settings

You can use many of the same access parameters with <code>oracle\_hdfs</code> as <code>oracle\_hive</code>.

### 2.4.2.1 Accessing a Delimited Text File

The following example is equivalent to the one shown in "Overriding the Default ORACLE\_HIVE Settings". The external table access a delimited text file stored in HDFS.



```
CREATE TABLE order (cust_num VARCHAR2(10),
                    order_num VARCHAR2(20),
                    order_date DATE,
                    item cnt NUMBER,
                    description VARCHAR2(100),
                    order_total NUMBER(8,2))
   ORGANIZATION EXTERNAL
    TYPE oracle_hdfs
    DEFAULT DIRECTORY DEFAULT_DIR
       ACCESS PARAMETERS
        com.oracle.bigdata.colmap: {"col":"item_cnt",
"field": "order line item count" }
        com.oracle.biqdata.overflow: {"action":"TRUNCATE", "col":"DESCRIPTION"}
        com.oracle.bigdata.erroropt: [{"action":"replace", \
                                          "value":"INVALID NUM", \
                                          "col":["CUST_NUM","ORDER_NUM"]} , \
                                          {"action": "reject", "col": "ORDER_TOTAL}]
 LOCATION ('hdfs:/usr/cust/summary/*'));
```

The parameters make the following changes in the way that the <code>ORACLE\_HDFS</code> access driver locates the data and handles error conditions:

- com.oracle.bigdata.colmap: Handles differences in column names.

  ORDER\_LINE\_ITEM\_COUNT in the HDFS files matches the ITEM\_CNT column in the external table.
- com.oracle.bigdata.overflow: Truncates string data. Values longer than 100 characters for the DESCRIPTION column are truncated.
- com.oracle.bigdata.erroropt: Replaces bad data. Errors in the data for CUST\_NUM or ORDER\_NUM Set the value to INVALID\_NUM.

### 2.4.2.2 Accessing Avro Container Files

The next example uses a SerDe to access Avro container files.

```
CREATE TABLE order (cust_num VARCHAR2(10),
                    order_num VARCHAR2(20),
                    order_date DATE,
                    item_cnt NUMBER,
                    description VARCHAR2(100),
                    order_total NUMBER(8,2))
   ORGANIZATION EXTERNAL
     TYPE oracle_hdfs
     DEFAULT DIRECTORY DEFAULT_DIR
     ACCESS PARAMETERS (
        com.oracle.bigdata.rowformat: \
        SERDE 'org.apache.hadoop.hive.serde2.avro.AvroSerDe'
        com.oracle.bigdata.fileformat: \
        INPUTFORMAT 'org.apache.hadoop.hive.ql.io.avro.AvroContainerInputFormat'\
        OUTPUTFORMAT 'org.apache.hadoop.hive.ql.io.avro.AvroContainerOutputFormat'
        com.oracle.bigdata.colmap: { "col":"item_cnt", \
                 "field": "order_line_item_count" }
         com.oracle.bigdata.overflow: {"action":"TRUNCATE", \
                  "col":"DESCRIPTION"}
     LOCATION ('hdfs:/usr/cust/summary/*'));
```



The access parameters provide the following information to the <code>ORACLE\_HDFS</code> access driver:

- com.oracle.bigdata.rowformat: Identifies the SerDe that the access driver needs to use to parse the records and fields. The files are not in delimited text format.
- com.oracle.bigdata.fileformat: Identifies the Java classes that can extract records and output them in the desired format.
- com.oracle.bigdata.colmap: Handles differences in column names. ORACLE\_HDFS matches ORDER\_LINE\_ITEM\_COUNT in the HDFS files with the ITEM\_CNT column in the external table.
- com.oracle.bigdata.overflow: Truncates string data. Values longer than 100 characters for the DESCRIPTION column are truncated.

# 2.5 Creating an Oracle External Table for Kafka Topics

The Hive storage handler for Kafka enables Hive and Oracle Big Data SQL and Hive to query Kafka topics.

The ORACLE\_HIVE access driver can access Kafka data topics. You first create a Hive external table that accesses the Kafka topics and then create an Oracle Big Data SQL table over it.

# 2.5.1 Using Oracle's Hive Storage Handler for Kafka to Create a Hive External Table for Kafka Topics

The Hive storage handler for Kafka enables Hive and Oracle Big Data SQL to query Kafka topics.

To provide access to Kafka data, you create a Hive external table over the Kafka topics. The Oracle Big Data SQL storage handler that enables Hive to read the Kafka data format is <code>oracle.hadoop.kafka.hive.KafkaStorageHandler</code>.

You can use this storage handler to create external Hive tables backed by data residing in Kafka. Big Data SQL can then query the Kafka data through the external Hive tables.

The Hive DDL is demonstrated by the following example, where topic1 and topic2 are two topics in Kafka broker whose keys are serialized by Kafka's String serializer and whose values are serialized by kafka's Long serializer.

The example below shows the resulting Hive table. The Kafka key, value, offset, topic name, and partitionid are mapped to Hive columns. You can explicitly designate the offset for each topic/partition pair through a WHERE clause in you Hive guery.

```
hive> describe test_table;
OK
topic string from descrializer
```



```
partitionid
              int
                                  from deserializer
key
              string
                                  from deserializer
value
             bigInt
                                  from deserializer
offset
             bigint
                                   from deserializer
timestamptype smallInt
                              from deserializer
             timestamp
                               from deserializer
timestamp
Time taken: 0.084 seconds, Fetched: 7 row(s)
```

The content of the table is a snapshot of the Kafka topics when the Hive query is executed. When new data is inserted into the Kafka topics, you can use the offset column or the timestamp column to track the changes to the topic. The offsets are per topic/partition. For example, the following query will return new messages after the specified offsets in the where clause for each topic/partition:

```
hive> select * from test_table where (topic="topic1" and partitionid=0 and offset > 199) or (topic="topic1" and partitionid=1 and offset > 198) or (topic="topic2" and partitionid=0 and offset > 177) or (topic="topic2" and partitionid=1 and offset > 176);
```

You need to keep track of the offsets for all topic/partition. For example, you can use an Oracle table to store these offsets. A more convenient way to keep track of new data is using the timestamp column. You can query data after a specific time point using the following query:

```
hive> select * from test_table where timestamp > '2017-07-12 11:30:00';
```

See the Property Reference section below for descriptions of all table properties

#### **Property Reference**

Table 2-3 Table Properties of Hive Storage Handler for Kafka

Property Name	Requir ement	Description
oracle.kaf ka.table.to pics	•	A comma-separated list of Kafka topics. Each Kafka topic name must consists of only letters (uppercase and lowercase), numbers, .(dot), _(underscore), and -(minus). The maximum length for each topic name is 249. These topics must have the same serialization mechanisms. The resulting Hive table consists of records from all the topics listed here. A Hive column "topic" will be added and it will be set to the topic name for each record.
oracle.kaf ka.bootstr ap.servers	•	This property will be translated to the "bootstrap.servers" property for the underlying Kafka consumer. The consumer makes use of all servers, irrespective of which servers are specified here for bootstrapping. This list only impacts the initial hosts used to discover the full set of servers. This list should be in the form <code>host1:port1,host2:port2,</code> Since these servers are just used for the initial connection to discover the full cluster membership (which may change dynamically), this list need not contain the full set of servers. For availability reasons, you may want to list more than one server.



Table 2-3 (Cont.) Table Properties of Hive Storage Handler for Kafka

Property Name	Requir ement	Description
oracle.kaf ka.table.k ey.type	Optional	The key type for your record. If unset, then the key part of the Kafka record will be ignored in the Hive row. Only values of "string", "integer", "long", "double", "avro", "avro_confluent" are supported. "string", "integer", "double" and "long" correspond to the built-in primitive serialization types supported by Kafka. If this property is one of these primitive types, then the Kafka key for each record will be mapped to one single Hive Column. If this property is set to "avro" or "avro_confluent", then oracle.kafka.table.key.schema is required. The Kafka key for each record will be deserialized into an Avro Object. If the Avro schema is of record type then each first level field of the record will be mapped to a single Hive column. If the Avro schema is not of Record Type, then it will be mapped to a single Hive Column named "key".
		The difference between "avro" and "avro_confluent" is that the wire format for the serialization is slightly different. For "avro", the entire bytes array of the key consists of the bytes of avro serialization. For "avro_confluent", the bytes array consists of a magic byte, a version number, then the bytes of avro serialization of the key.
oracle.kaf ka.table.v alue.type	Optional	The value type of your record. If unset, then the value part of Kafka record will be ignored in the Hive row. Use of this property is similar to use of oracle.kafka.table.key.type. The difference between them is: when the Avro Schema for Kafka value is not of record type. The whole Avro object will be mapped to a single Hive Column named "value" instead of "key".
oracle.kaf ka.table.k ey.writer.s chema	Optional	An optional writer schema for the Kafka key's Avro serialization. It's required when the reader schema for the key is different from the schema in which the keys are written to Kafka brokers. It must be the exact schema in which Kafka keys are serialized.
oracle.kaf ka.table.k ey.schem a		The JSON string for the Kafka key's Avro reader schema. It doesn't need to be exactly the same as the Kafka key's writer Avro schema. As long as the reader schema is compatible with the Kafka key or the converted object from the converter, it is valid. This enables you to rename Hive columns and choose what fields to keep from the Kafka key in the Hive row. If the schema in this property is different from the schema in which the Kafka keys are serialized, then oracle.kafka.table.key.writer.schema is required.
oracle.kaf ka.table.v alue.writer .schema	Optional	An optional writer schema for the Kafka value's Avro serialization. Its use is similar to oracle.kafka.table.key.writer.schema.



Table 2-3 (Cont.) Table Properties of Hive Storage Handler for Kafka

Property Name	Requir ement	Description
oracle.kaf ka.table.v alue.sche ma	Require d when "oracle. kafka.ta ble.valu e.type" is "avro" or "avro_c onfluent	The JSON string for the Kafka value's Avro reader schema. Its use is similar to oracle.kafka.table.key.schema.
oracle.kaf ka.table.e xtra.colum ns	, default	A boolean flag to control whether to include extra Kafka columns: paritionid, offset, timestamptype.
oracle.kaf ka.chop.p artition		A Boolean flag to control whether to chop Kafka partitions into smaller chunks. This is useful when the number of Kafka partitions is small and the size of each Kafka partition is large.
oracle.kaf ka.partitio n.chunk.si ze	Optional	When oracle.kafka.chop.partition is true, this property controls the number of Kafka records in each partition chunk. It should be set a value estimated by (Ideal size of a split)/(Average size of a Kafka record). For example, if the ideal size of a split is 256 MB and the average size of s Kafka record is 256 Bytes, then this property should be set to 1000000.

### 2.5.2 Creating an Oracle Big Data SQL Table for Kafka Topics

Big Data SQL can use the ORACLE\_HIVE access driver to query data stored in Hive Tables.

After you create a Hive table over Kafka data by using the Hive storage handler for Kafka, there are no special procedures for generating a Big Data SQL table from the resulting Hive table. The default ORACLE\_HIVE settings can be overridden in the same way as with other Hive tables This is how to query the Hive external table that was created using the Hive storage handler for Kafka in the previous section..



) PARALLEL 2 REJECT LIMIT UNLIMITED

# 2.6 Using the Custom Parquet Reader for Oracle Big Data SQL

For reading parquet files, you have the option of using the custom Parquet reader for Oracle Big Data SQL.

The custom reader includes an optimization that reduces I/O – lazy retrieval and materialization of non-filter predicate columns. When evaluating a filter for a row-group (about 128 MB of data), the custom reader first retrieves only the columns for the filter predicate. The filter is applied and if there is no match then the reader moves to the next row group. If some row matches then the non-filter columns are read. Filter evaluation builds a data structure to efficiently de-serialize (materialize) only the values for the matching column indexes. This differs from the Hive implementation which reads AND de-serializes (materializes) all the columns in the select list, then evaluates the filter for the predicate columns, then finally assembles rows from the matching values. This optimization reduces I/O whenever there are no matching rows for the filter in the row group. In this case, the non-predicate columns are not read at all.

The custom reader also reduces CPU consumption in two ways:

- Skips decompression of individual Parquet pages (the minimal read unit) when those pages contain no values that satisfy the filter. This applies to primitive column types only, not to nested column types.
- Skips the de-serialization of individual column values if they do not belong to a matching row index.

#### Disabling or Re-Enabling the Custom Parguet Reader

The Parquet reader optimization is enabled by default. It can be disabled for an individual table by adding the following access parameter to the external table definition:

com.oracle.bigdata.useOracleParquet=false

You can add this setting to the cluster properties file to disable the optimization for all Parquet-based external tables. Remove the setting to return to the default.

### **Compatibility with Previously Created Parquet Format Data**

Use of the customer reader requires no changes to data format. However, for best performance, the format must provide min and max values for each column for each Parquet block. These values are used by the standard Hadoop Parquet InputFormat, as well as the custom Parquet reader, to optimize the query. The resulting optimization significantly improves query performance with both Hive and Oracle Big Data SQL.

Note that Parquet files created by Impala do not include min and max values for each column for each Parquet block.

To ensure that min and max values are available, it is recommended that you write Parquet files with Hive or other tools that generate output in the standard Hadoop Parquet InputFormat, such as PrestoDB and Spark.



To check if a file includes these values, you can use the parquet tools JAR to dump information about the file:

```
# hadoop jar parquet-tools-1.5.0-cdh5.12.0.jar meta <filename>.parq
```

On a CDH Hadoop distro, the parquet-tools command may also be configured in your path.

# 2.7 About the SQL CREATE TABLE Statement

The SQL CREATE TABLE statement has a clause specifically for creating external tables. The information that you provide in this clause enables the access driver to read data from an external source and prepare the data for the external table.

# 2.7.1 Basic Syntax

The following is the basic syntax of the CREATE TABLE statement for external tables:

```
CREATE TABLE table_name (column_name datatype, column_name datatype[,...])
ORGANIZATION EXTERNAL (external_table_clause);
```

You specify the column names and data types the same as for any other table. ORGANIZATION EXTERNAL identifies the table as an external table.

The *external\_table\_clause* identifies the access driver and provides the information that it needs to load the data. See "About the External Table Clause".

### 2.7.2 About the External Table Clause

CREATE TABLE ORGANIZATION EXTERNAL takes the *external\_table\_clause* as its argument. It has the following subclauses:

- TYPE Clause
- DEFAULT DIRECTORY Clause
- LOCATION Clause
- REJECT LIMIT Clause
- ORACLE HIVE Access Parameters



Oracle Database SQL Language Reference for the external\_table\_clause

### 2.7.2.1 TYPE Clause

The TYPE clause identifies the access driver. The type of access driver determines how the other parts of the external table definition are interpreted.

Specify one of the following values for Oracle Big Data SQL:

ORACLE\_HDFS: Accesses files in an HDFS directory.



ORACLE\_HIVE: Accesses a Hive table.



The <code>ORACLE\_DATAPUMP</code> and <code>ORACLE\_LOADER</code> access drivers are not associated with Oracle Big Data SQL.

### 2.7.2.2 DEFAULT DIRECTORY Clause

The DEFAULT DIRECTORY clause identifies an Oracle Database directory object. The directory object identifies an operating system directory with files that the external table reads and writes.

ORACLE\_HDFS and ORACLE\_HIVE use the default directory solely to write log files on the Oracle Database system.

### 2.7.2.3 LOCATION Clause

The LOCATION clause identifies the data source.

### 2.7.2.4 ORACLE HDFS LOCATION Clause

The LOCATION clause for ORACLE\_HDFS contains a comma-separated list of file locations. The files must reside in the HDFS file system on the default cluster.

A location can be any of the following:

- A fully qualified HDFS directory name, such as /user/hive/warehouse/hive\_seed/hive\_types. ORACLE\_HDFS uses all files in the directory.
- A fully qualified HDFS file name, such as /user/hive/warehouse/hive\_seed/ hive\_types/hive\_types.csv
- A URL for an HDFS file or a set of files, such as hdfs:/user/hive/warehouse/hive\_seed/hive\_types/\*. It is invalid to use the directory name alone.

The file names can contain any pattern-matching character described in Table 2-4.

Table 2-4 Pattern-Matching Characters

Character	Description
?	Matches any one character
*	Matches zero or more characters
[abc]	Matches one character in the set $\{a, b, c\}$
[a-b]	Matches one character in the range $\{ab\}$ . The character must be less than or equal to $b$ .
[^a]	Matches one character that is not in the character set or range {a}. The carat (^) must immediately follow the left bracket, with no spaces.
\c	Removes any special meaning of $c$ . The backslash is the escape character.



Table 2-4 (Cont.) Pattern-Matching Characters

Character	Description
{abcd}	Matches a string from the set $\{ab, cd\}$ . The escape character (\) removes the meaning of the comma as a path separator.
{abc{defh}	Matches a string from the set {ab, cde, cfh}. The escape character (\) removes the meaning of the comma as a path separator.

### 2.7.2.5 ORACLE\_HIVE LOCATION Clause

Do not specify the LOCATION clause for ORACLE\_HIVE; it raises an error. The data is stored in Hive, and the access parameters and the metadata store provide the necessary information.

### 2.7.2.6 REJECT LIMIT Clause

Limits the number of conversion errors permitted during a query of the external table before Oracle Database stops the query and returns an error.

Any processing error that causes a row to be rejected counts against the limit. The reject limit applies individually to each parallel query (PQ) process. It is not the total of all rejected rows for all PO processes.

### 2.7.2.7 ACCESS PARAMETERS Clause

The ACCESS PARAMETERS clause provides information that the access driver needs to load the data correctly into the external table. See "CREATE TABLE ACCESS PARAMETERS Clause".

# 2.8 About Data Type Conversions

When the access driver loads data into an external table, it verifies that the Hive data can be converted to the data type of the target column. If they are incompatible, then the access driver returns an error. Otherwise, it makes the appropriate data conversion.

Hive typically provides a table abstraction layer over data stored elsewhere, such as in HDFS files. Hive uses a serializer/deserializer (SerDe) to convert the data as needed from its stored format into a Hive data type. The access driver then converts the data from its Hive data type to an Oracle data type. For example, if a Hive table over a text file has a BIGINT column, then the SerDe converts the data from text to BIGINT. The access driver then converts the data from BIGINT (a Hive data type) to NUMBER (an Oracle data type).

Performance is better when one data type conversion is performed instead of two. The data types for the fields in the HDFS files should therefore indicate the data that is actually stored on disk. For example, JSON is a clear text format, therefore all data in a JSON file is text. If the Hive type for a field is DATE, then the SerDe converts the data from string (in the data file) to a Hive date. Then the access driver converts the data from a Hive date to an Oracle date. However, if the Hive type for the field is string, then the SerDe does not perform a conversion, and the access driver converts the



data from string to an oracle date. Queries against the external table are faster in the second example, because the access driver performs the only data conversion.

The table below identifies the data type conversions that  $\[mathbb{ORACLE\_HIVE}\]$  can make when loading data into an external table.

Table 2-5 Supported Hive to Oracle Data Type Conversions

Hive Data Type	VARCHAR 2, CHAR, NCHAR2, NCHAR, CLOB	NUMBER, FLOAT, BINARY_N UMBER, BINARY_F LOAT	BLOB	RAW	DATE, TIMESTAM P, TIMESTAM P WITH TZ, TIMESTAM P WITH LOCAL TZ	INTERVAL YEAR TO MONTH, INTERVAL DAY TO SECOND
INT SMALLINT TINYINT BIGINT	yes	yes	yes	yes	no	no
DOUBLE FLOAT	yes	yes	yes	yes	no	no
DECIMAL	yes	yes	no	no	no	no
BOOLEAN	yes <sup>1</sup>	yes <sup>2</sup>	yes <sup>2</sup>	yes	no	no
BINARY	yes	no	yes	yes	no	no
STRING	yes	yes	no	no	yes	yes
TIMESTAM P	yes	no	no	no	yes	no
STRUCT ARRAY UNIONTYP E MAP	yes	no	no	no	no	no

 $<sup>^{1}</sup>$  FALSE maps to the string FALSE, and TRUE maps to the string <code>TRUE</code>.

# 2.9 Querying External Tables

Users can query external tables using the SQL  $\mbox{select}$  statement, the same as they query any other table.

# 2.9.1 Granting User Access

Users who query the data on a Hadoop cluster must have READ access in Oracle Database to the external table and to the database directory object that points to the cluster directory. See "About the Cluster Directory".



<sup>&</sup>lt;sup>2</sup> FALSE maps to 0, and TRUE maps to 1.

# 2.9.2 About Error Handling

By default, a query returns no data if an error occurs while the value of a column is calculated. Processing continues after most errors, particularly those thrown while the column values are calculated.

Use the com.oracle.bigdata.erroropt and com.oracle.bigdata.overflow parameters to determine how errors are handled.

# 2.9.3 About the Log Files

You can use these access parameters to customize the log files:

- com.oracle.bigdata.log.exec
- com.oracle.bigdata.log.qc

# 2.10 About Oracle Big Data SQL on the Database Server (Oracle Exadata Machine or Other)

This section explains the changes that the Oracle Big Data SQL installation makes to the Oracle Database system (which may or may not be an Oracle Exadata Machine).

The section contains the following topics:

- About the bigdata\_config Directory
- Common Configuration Properties
- About the Cluster Directory

# 2.10.1 About the bigdata\_config Directory

The directory bigdata\_config contains configuration information that is common to all Hadoop clusters. This directory is located on the Oracle Database system under <code>\$ORACLE\_HOME/bigdatasql</code>. The oracle file system user (or whichever user owns the Oracle Database instance) owns <code>bigdata\_config</code>. The Oracle Database directory <code>ORACLE\_BIGDATA\_CONFIG</code> points to <code>bigdata\_config</code>.

# 2.10.2 Common Configuration Properties

The installation store these files in the bigdata\_config directory under <code>\$ORACLE\_HOME/bigdatasql</code>:

- bigdata.properties
- bigdata-log4j.properties

The Oracle DBA can edit these configuration files as necessary.



# 2.10.2.1 bigdata.properties

The bigdata.properties file in the common directory contains property-value pairs that define the Java class paths and native library paths required for accessing data in HDFS.

These properties must be set:

- bigdata.cluster.default
- java.classpath.hadoop
- java.classpath.hive
- java.classpath.oracle

The following list describes all properties permitted in bigdata.properties.

### bigdata.properties

Property	Description
bigdata.cluster.default	The name of the default Hadoop cluster. The access driver uses this name when the access parameters do not specify a cluster. Required.
	Changing the default cluster name might break external tables that were created previously without an explicit cluster name.
bigdata.cluster.list	A comma-separated list of Hadoop cluster names. Optional.
java.classpath.hadoop	The Hadoop class path. Required.
java.classpath.hive	The Hive class path. Required.
java.classpath.oracle	The path to the Oracle JXAD Java JAR file. Required.
java.classpath.user	The path to user JAR files. Optional.
java.libjvm.file	The full file path to the JVM shared library (such as libjvm.so). Required.
java.options	A comma-separated list of options to pass to the JVM. Optional.
	This example sets the maximum heap size to 2 GB, and verbose logging for Java Native Interface (JNI) calls:
	Xmx2048m,-verbose=jni



Property	Description
java.options2	A space-delimited list of options to pass to the JVM. Optional. The delimiter must be a space character, not a tab or other whitespace character.
	This example sets the maximum heap size to 2 GB, and verbose logging for Java Native Interface (JNI) calls:
	Xmx2048m -verbose=jni
	Notice that java.options is comma- delimited, while java.options2 is space delimited. These two properties can coexist in the same bigdata.proper ties file.
LD_LIBRARY_PATH	A colon separated (:) list of directory paths to search for the Hadoop native libraries. Recommended.
	If you set this option, then do not set <i>java.library</i> path in java.options.

Example 2-1 shows a sample bigdata.properties file.

### Example 2-1 Sample bigdata.properties File

```
# bigdata.properties
#
# Copyright (c) 2014, Oracle and/or its affiliates. All rights reserved.
#
# NAME
# bigdata.properties - Big Data Properties File
#
DESCRIPTION
# Properties file containing parameters for allowing access to Big Data
# Fixed value properties can be added here
#

java.libjvm.file=$ORACLE_HOME/jdk/jre/lib/amd64/server/libjvm.so
java.classpath.oracle=$ORACLE_HOME/hadoopcore/jlib/*:$ORACLE_HOME/hadoop/jlib/hver-2/
*:$ORACLE_HOME/dbjava/lib/*
java.classpath.hadoop=$HADOOP_HOME/*:$HADOOP_HOME/lib/*
java.classpath.hive=$HIVE_HOME/lib/*
LD_LIBRARY_PATH=$ORACLE_HOME/jdk/jre/lib
bigdata.cluster.default=hadoop_cl_1
```



### 2.10.2.2 bigdata-log4j.properties

The bigdata-log4j.properties file in the common directory defines the logging behavior of queries against external tables in the Java code. Any log4j properties are allowed in this file.

Example 2-2 shows a sample bigdata-log4j.properties file with the relevant log4j properties.

#### Example 2-2 Sample bigdata-log4j.properties File

```
# bigdata-log4j.properties
# Copyright (c) 2014, Oracle and/or its affiliates. All rights reserved.
      bigdata-log4j.properties - Big Data Logging Properties File
    DESCRIPTION
      Properties file containing logging parameters for Big Data
      Fixed value properties can be added here
bigsql.rootlogger=INFO,console
log4j.rootlogger=DEBUG, file
log4j.appender.console=org.apache.log4j.ConsoleAppender
log4j.appender.console.target=System.err
log4j.appender.console.layout=org.apache.log4j.PatternLayout
log4j.appender.console.layout.ConversionPattern=%d{yy/MM/dd HH:mm:ss} %p %c{2}: %m%n
log4j.appender.file=org.apache.log4j.RollingFileAppender
log4j.appender.file.layout=org.apache.log4j.PatternLayout
log4j.appender.file.layout.ConversionPattern=%d{yy/MM/dd HH:mm:ss} %p %c{2}: %m%n
log4j.logger.oracle.hadoop.sql=ALL, file
bigsql.log.dir=.
bigsgl.log.file=bigsgl.log
log4j.appender.file.File=$ORACLE_HOME/bigdatalogs/bigdata-log4j.log
```

### See Also:

Apache Logging Services documentation at

http://logging.apache.org/log4j/1.2/manual.html

## 2.10.3 About the Cluster Directory

The cluster directory contains configuration information for a Hadoop cluster. Each cluster that Oracle Database accesses using Oracle Big Data SQL has a cluster directory. This directory is located on the Oracle Database system under <code>\$ORACLE\_HOME/bigdatasql/clusters/</code>. For example, cluster <code>bdal\_cl\_1</code> would have a directory <code>\$ORACLE\_HOME/bigdatasql/clusters/bdal\_cl\_1</code> and under <code>\$ORACLE\_HOME/bigdatasql/clusters/bdal\_cl\_1/config would</code> be the following files for client configuration files for accessing the cluster:

bigdata.hosts (not editable by customers)



- core-site.xml
- hdfs-site.xml
- hive-site.xml
- mapred-site.xml (optional)
- log4j property files (such as hive-log4j.properties)

\$ORACLE\_HOME/bigdatasql/databases/<database name>/bigdata\_config/default\_cluster iS
a soft link to the directory of the default cluster.

A database directory object points to the cluster directory. Users who want to access the data in a cluster must have read access to the directory object.

### 2.10.4 About Permissions

On the Oracle database server, ensure that the <code>oracle</code> user (or whatever user owns the Oracle Database installation directory) has READ/WRITE access to the database directory that points to the log directory.

On the Hadoop side, when you run Database Acknowledge (# ./jaguar databaseack [config file]) this operation creates an account for the database owner and grants required permissions.



# Storing Oracle Data in Hadoop

Copy to Hadoop and Oracle Database Tablespaces in HDFS are two Oracle Big Data SQL resources for off-loading Oracle Database tables to the HDFS file system on a Hadoop cluster.

The table below compares these two tools.

Table 3-1 Comparison of Copy to Hadoop and Oracle Tablespaces in HDFS

Copy to Hadoop	Oracle Tablespaces in HDFS
Copies Oracle Database tables to Oracle Data Pump files stored in HDFS.	Oracle Database tables or partitions are stored within the tablespace in HDFS in their original Oracle-internal format.
Access is through a Hive external table and from the database with Oracle Big Data SQL.	Access is directly though the original Oracle Database tables. External tables are not needed.
Data is available (through Hive) to other processes in the Hadoop ecosystem and to Oracle Database (through Oracle Big Data SQL).	Data is directly available to Oracle Database only. Data is not accessible to other processes in Hadoop.

# 3.1 Using Copy to Hadoop

This section describes how to use Copy to Hadoop to copy Oracle Database tables to Hadoop.

- What Is Copy to Hadoop?
- Getting Started Using Copy to Hadoop
- Using Oracle Shell for Hadoop Loaders With Copy to Hadoop
- Using Copy to Hadoop to do Direct Copies
- Using Copy to Hadoop to do Staged Copies
- Querying the Data in Hive
- About Column Mappings and Data Type Conversions
- Working With Spark

### 3.1.1 What Is Copy to Hadoop?

Oracle Big Data SQL includes the Oracle Copy to Hadoop utility. This utility makes it simple to identify and copy Oracle data to the Hadoop Distributed File System. It can be accessed through the command-line interface Oracle Shell for Hadoop Loaders.

Data exported to the Hadoop cluster by Copy to Hadoop is stored in Oracle Data Pump format. The Oracle Data Pump files can be queried by Hive or Big Data SQL. The Oracle Data Pump format optimizes queries through Big Data SQL in the following ways:

- The data is stored as Oracle data types eliminating data type conversions.
- The data is queried directly without requiring the overhead associated with Java SerDes.

After Data Pump format files are in HDFS, you can use Apache Hive to query the data. Hive can process the data locally without accessing Oracle Database. When the Oracle table changes, you can refresh the copy in Hadoop. Copy to Hadoop is primarily useful for Oracle tables that are relatively static, and thus do not require frequent refreshes.

Copy to Hadoop is licensed under Oracle Big Data SQL. You must have an Oracle Big Data SQL license in order to use this utility.

# 3.1.2 Getting Started Using Copy to Hadoop

To install and start using Copy to Hadoop:

- 1. Follow the Copy to Hadoop and Oracle Shell for Hadoop Loaders installation procedures in the Oracle Big Data SQL Installation Guide.
  - As described in the installation guide, ensure that the prerequisite software is installed on both the Hadoop cluster (on Oracle Big Data Appliance or another Hadoop system) and on the Oracle Database server (Oracle Exadata Database Machine or other).
- Invoke Oracle Shell for Hadoop Loaders (OHSH) to do a direct, one-step copy or a staged, two-step copy of data in Oracle Database to Data Pump format files in HDFS, and create a Hive external table from the files.
  - OHSH will choose directcopy by default to do a direct, one-step copy. This is faster than a staged, two-step copy and does not require storage on the database server. However, there are situations where you should do a staged, two-step copy:
  - Copying columns from multiple Oracle Database source tables. (The direct, one-step copy copies data from one table only.)
  - Copying columns of type TIMESTAMPTZ or TIMESTAMPLTZ to Hive.
    - Since Hive does not have a data type that supports time zones or time offsets, you must cast these columns to TIMESTAMP when manually exporting these columns to Data Pump files
  - Copying data from a view. Views are not supported by the directcopy option.

The staged two-step copy using the manual steps is demonstrated in "Appendix A: Manual Steps for Using Copy to Hadoop for Staged Copies".

3. Query this Hive table the same as you would any other Hive table.



#### Tip:

For Hadoop power users with specialized requirements, the manual option for Direct Copy is recommended. See Manual Steps for Using Copy to Hadoop for Direct Copies in Appendix B.



### 3.1.2.1 Table Access Requirements for Copy to Hadoop

To copy table using Copy to Hadoop, an Oracle Database user must meet one of these requirements.

- The user is the owner of the table, or
- The user has one of the following privileges to access a table in another schema:
  - The select\_catalog\_role privilege (which provides SELECT privileges on data dictionary views).
  - The SELECT privilege on the table.

## 3.1.3 Using Oracle Shell for Hadoop Loaders With Copy to Hadoop

### 3.1.3.1 Introducing Oracle Shell for Hadoop Loaders

### What is Oracle Shell for Hadoop Loaders?

Oracle Shell for Hadoop Loaders (OHSH) is a helper shell that provides a simple to use command line interface to Oracle Loader for Hadoop, Oracle SQL Connector for HDFS, and Copy to Hadoop. It has basic shell features such as command line recall, history, inheriting environment variables from the parent process, setting new or existing environment variables, and performing environmental substitution in the command line.

The core functionality of Oracle Shell for Hadoop Loaders includes the following:

- Defining named external resources with which Oracle Shell for Hadoop Loaders interacts to perform loading tasks.
- Setting default values for load operations.
- Running load commands.
- Delegating simple pre and post load tasks to the Operating System, HDFS, Hive and Oracle. These tasks include viewing the data to be loaded, and viewing the data in the target table after loading.

### **Getting Started with Oracle Shell for Hadoop Loaders**

The examples directory in the OHSH kit contains many examples that define resources and load data using Oracle Shell for Hadoop Loaders. See <OHSH\_KIT>/examples/README.txt for a description of the examples and instructions for running them.

### **Getting Help**

The OHSH shell provides online help for all commands.

To get a list of all possible OHSH commands:

ohsh> help

To get help on a specific command, enter help, followed by the command:



ohsh> help show

The table below describes the help categories available.

Help Command	Description
help load	Describes load commands for Oracle and Hadoop tables.
help set	Shows help for setting defaults for load operations. It also describes what load methods are impacted by a particular setting.
help show	Shows help for inspecting default settings.
help shell	Shows shell-like commands.
help resource	Show commands for creating and dropping named resources.



To set up Oracle Shell for Hadoop Loaders, follow the instructions in the Oracle Big Data SQL Installation Guide.

# 3.1.4 Using Copy to Hadoop to do Direct Copies

This example shows how to use Oracle Shell for Hadoop Loaders (OHSH) to do a direct, one step copy from Oracle Database to Hadoop.

The example assumes OHSH and Oracle Big Data SQL have been installed and configured, and that the examples have been configured by following the instructions in the README.txt file found in the examples directory of the OHSH installation. The example scripts can be found in the examples directory of the OHSH installation.

#### Example 3-1 createreplace\_directcopy.ohsh

This script uses create or replace to create a Hive table called <code>cp2hadoop\_fivdti</code> from the Oracle table <code>FIVDTI</code> and then loads the Hive table with 10000 rows. It uses the <code>directcopy</code> command to run a map job on Hadoop and split the Oracle table into input splits. It then creates Data Pump format files in HDFS that include all the splits, and creates a Hive external table that maps to the Data Pump format files.

create or replace hive table hive0:cp2hadoop\_fivdti \
from oracle table olhp:fivdti using directcopy

The <code>load\_directcopy.ohsh</code> script shows how to load the Hive table with an additional 30 rows using the <code>directcopy</code> command.

load hive table hive0:cp2hadoop\_fivdti from oracle table olhp:fivdti  $\setminus$  using directcopy where "(i7 < 30)";

#### Optionally, the data can be converted to Parquet or ORC formats from OHSH:

%hive0 create table cp2hadoop\_fivdti\_parquet stored as parquet as select \* from cp2hadoop\_fivdti



# 3.1.5 Using Copy to Hadoop to do Staged Copies

This example shows how to use Oracle Shell for Hadoop Loaders (OHSH) to do a staged, two-step copy from Oracle Database to Hadoop.

This example assumes OHSH and Big Data SQL have been installed and configured, and that the examples have been configured according to the instructions in README.txt in the examples directory of the OHSH installation. The scripts below and many others are also available in the examples directory.

### Example 3-2 createreplace\_stage.ohsh

This script uses create or replace to create a Hive table called <code>cp2hadoop\_fivdti</code> from the Oracle table <code>FIVDTI</code>. It uses the <code>stage</code> command, which automatically does the following:

- Exports the contents of the source table in Oracle to Data Pump format files on local disk
- 2. Moves the Data Pump format files to HDFS.
- 3. Creates the Hive external table that maps to the Data Pump format files in HDFS.

```
create or replace hive table hive0:cp2hadoop_fivdti \
from oracle table olhp:fivdti using stage
```

#### Example 3-3 load\_stage.ohsh

The <code>load\_stage.ohsh</code> script shows how to load the Hive table with an additional 30 rows using the <code>stage</code> command.

```
load hive table hive0:cp2hadoop_fivdti from oracle table olhp:fivdti \ using stage where "(i7 < 30)";
```

#### **Manual Option**

The two-step method demonstrated in the <code>createreplace\_stage.ohsh</code> and <code>load\_stage.ohsh</code> example scripts automates some of the tasks required to do staged copies. However, there may be reasons to perform the steps manually, such as:

- You want to load columns from multiple Oracle Database source tables.
- You want to load columns of type TIMESTAMPZ or TIMESTAMPLTZ.

See Appendix A: Manual Steps for Using Copy to Hadoop for Staged Copies.

# 3.1.6 Querying the Data in Hive

The following OHSH command shows the number of rows in the Hive table after copying from the Oracle table.

```
%hiveO select count(*) from cp2hadoop_fivdti;
```

### 3.1.7 About Column Mappings and Data Type Conversions



### 3.1.7.1 About Column Mappings

The Hive table columns automatically have the same names as the Oracle columns, which are provided by the metadata stored in the Data Pump files. Any user-specified column definitions in the Hive table are ignored.

# 3.1.7.2 About Data Type Conversions

Copy to Hadoop automatically converts the data in an Oracle table to an appropriate Hive data type. Table 3-2 shows the default mappings between Oracle and Hive data types.

Table 3-2 Oracle to Hive Data Type Conversions

less than 10 BIGINT when the scale is 0 and the precision is less than 19 DECIMAL when the scale is greater than 0 or the precision is greater than 19  CLOB NCLOB INTERVALYM INTERVALYM INTERVALDS BINARY_DOUBLE BINARY_FLOAT BLOB BINARY_FLOAT BLOB BINARY ROWID UROWID RAW BINARY CHAR NCHAR VARCHAR2 VARCHAR2 NVARCHAR2 DATE TIMESTAMP TIMESTAMP TIMESTAMP TIMESTAMPTZ  DIRECTION IN SCALE IS 0 and the precision is greater than 0 or the precision is greater than 19 DECIMAL when the scale is 0 and the precision is greater than 19 DECIMAL When the scale is 0 and the precision is greater than 19 DECIMAL When the scale is 0 and the precision is greater than 0 or the precision is greater than 19  LOB  LOB  LOB  LOB  LOB  LOB  LOB  LO		
less than 10 BIGINT when the scale is 0 and the precision is less than 19 DECIMAL when the scale is greater than 0 or the precision is greater than 19  CLOB NCLOB INTERVALYM INTERVALYM INTERVALDS BINARY_DOUBLE BINARY_FLOAT BLOB BINARY_FLOAT BLOB BINARY ROWID UROWID RAW BINARY CHAR NCHAR VARCHAR2 VARCHAR2 NVARCHAR2 DATE TIMESTAMP TIMESTAMP TIMESTAMP TIMESTAMPTZ  DIRECTION IN SCALE IS 0 and the precision is greater than 0 or the precision is greater than 19 DECIMAL when the scale is 0 and the precision is greater than 19 DECIMAL When the scale is 0 and the precision is greater than 19 DECIMAL When the scale is 0 and the precision is greater than 0 or the precision is greater than 19  LOB  LOB  LOB  LOB  LOB  LOB  LOB  LO	Oracle Data Type	Hive Data Type
precision is less than 19 DECIMAL when the scale is greater than 0 or the precision is greater than 19  CLOB NCLOB INTERVALYM INTERVALDS BINARY_DOUBLE BINARY_FLOAT BLOB BINARY ROWID UROWID RAW BINARY CHAR NCHAR VARCHAR2 NVARCHAR2 DATE TIMESTAMP TIMESTAMP TIMESTAMPTZ  CLOB STRING STRING DOUBLE BOUBLE BINARY BINARY BINARY CHAR CHAR VARCHAR2 TIMESTAMP TIMESTAMP TIMESTAMPTZ Unsupported	NUMBER	INT when the scale is 0 and the precision is less than 10
CLOB NCLOB INTERVALYM INTERVALDS BINARY_DOUBLE BINARY_FLOAT BLOB ROWID UROWID RAW CHAR NCHAR VARCHAR2 DATE TIMESTAMP TIMESTAMP INTERVALDS  STRING STR		
INTERVALYM STRING INTERVALDS  BINARY_DOUBLE DOUBLE  BINARY_FLOAT FLOAT  BLOB BINARY  ROWID BINARY  UROWID  RAW BINARY  CHAR CHAR  NCHAR  VARCHAR2 VARCHAR2  DATE TIMESTAMP  TIMESTAMPTZ Unsupported		
INTERVALYM INTERVALDS  BINARY_DOUBLE  BINARY_FLOAT  BLOB  BINARY  ROWID  UROWID  RAW  BINARY  CHAR  NCHAR  VARCHAR2  DATE  TIMESTAMP  TIMESTAMP  BINARY  STRING  STRING  STRING  STRING  STRING  STRING  STRING  STRING  STRING  ANDUBLE  BINARY  BINARY  CHAR  CHAR  CHAR  VARCHAR2  VARCHAR2  Unsupported	CLOB	STRING
INTERVALDS  BINARY_DOUBLE  BINARY_FLOAT  BLOB  BINARY  ROWID  UROWID  RAW  BINARY  CHAR  NCHAR  VARCHAR2  NVARCHAR2  DATE  TIMESTAMP  TIMESTAMP  TIMESTAMP  DOUBLE  DOUBLE  BOUBLE  DOUBLE  DOUBLE  BINARY  CHAR  BINARY  CHAR  CHAR  CHAR  VARCHAR2  VARCHAR2  VARCHAR2  Unsupported	NCLOB	
BINARY_DOUBLE  BINARY_FLOAT  BLOB  BINARY  ROWID  UROWID  RAW  BINARY  CHAR  NCHAR  VARCHAR2  NVARCHAR2  DATE  TIMESTAMP  TIMESTAMP  TIMESTAMPTZ  DOUBLE  DOUBLE  DOUBLE  DOUBLE  DOUBLE  DATE  FLOAT  BINARY  CHAR  CHAR  VARCHAR  VARCHAR  TIMESTAMP	INTERVALYM	STRING
BINARY_FLOAT  BLOB  BINARY  ROWID  BINARY  UROWID  RAW  BINARY  CHAR  CHAR  NCHAR  VARCHAR2  NVARCHAR2  DATE  TIMESTAMP  TIMESTAMP  TIMESTAMPTZ  FLOAT  BINARY  CHAR  VARCHAR2  VARCHAR  VARCHAR2  UNsupported	INTERVALDS	
BLOB BINARY  ROWID BINARY  UROWID  RAW BINARY  CHAR CHAR  NCHAR  VARCHAR2  NVARCHAR2  DATE TIMESTAMP  TIMESTAMP  TIMESTAMPTZ  Unsupported	BINARY_DOUBLE	DOUBLE
ROWID UROWID  RAW BINARY  CHAR CHAR NCHAR  VARCHAR2 NVARCHAR2 DATE TIMESTAMP TIMESTAMP TIMESTAMPTZ  UROWID  BINARY  CHAR CHAR  CHAR  TIMESTAMP  TIMESTAMP  Unsupported	BINARY_FLOAT	FLOAT
UROWID  RAW BINARY  CHAR CHAR  NCHAR  VARCHAR2  NVARCHAR2  DATE TIMESTAMP  TIMESTAMP  TIMESTAMPTZ  Unsupported	BLOB	BINARY
RAW BINARY  CHAR CHAR  NCHAR  VARCHAR2  NVARCHAR2  DATE TIMESTAMP  TIMESTAMP  TIMESTAMP  TIMESTAMP  Unsupported	ROWID	BINARY
CHAR NCHAR  VARCHAR2  VARCHAR2  NVARCHAR2  DATE  TIMESTAMP  TIMESTAMP  TIMESTAMP  TIMESTAMP  Unsupported	UROWID	
NCHAR  VARCHAR2  NVARCHAR2  DATE  TIMESTAMP  TIMESTAMP  TIMESTAMP  Unsupported	RAW	BINARY
VARCHAR2  NVARCHAR2  DATE  TIMESTAMP  TIMESTAMP  TIMESTAMPTZ  Unsupported	CHAR	CHAR
NVARCHAR2  DATE TIMESTAMP  TIMESTAMP TIMESTAMP  TIMESTAMPTZ Unsupported	NCHAR	
DATE TIMESTAMP  TIMESTAMP  TIMESTAMP  Unsupported	VARCHAR2	VARCHAR
TIMESTAMP TIMESTAMP TIMESTAMPTZ Unsupported	NVARCHAR2	
TIMESTAMPTZ Unsupported	DATE	TIMESTAMP
2 22 11 12 22 2	TIMESTAMP	TIMESTAMP
TIMESTAMPI T71	TIMESTAMPTZ	Unsupported
TIMESTAMI ETZ	TIMESTAMPLTZ <sup>1</sup>	

To copy TIMESTAMPTZ and TIMESTAMPLTZ data to Hive, follow the instructions in Appendix A: Manual Steps for Using Copy to Hadoop to do Staged Copies. Cast the columns to TIMESTAMP when exporting them to the Data Pump files.

# 3.1.8 Working With Spark

The Oracle Data Pump files exported by Copy to Hadoop can be used in Spark.



The Spark installation must be configured to work with Hive. Launch a Spark shell by specifying the Copy to Hadoop jars.

```
prompt> spark-shell --jars orahivedp.jar,ojdbc7.jar,oraloader.jar,orail8n.jar,ora-
hadoop-common.jar
```

Verify the type of sqlContext in spark-shell:

```
scala> sqlContext
```

Your output will look like the following:

```
res0:org.apache.spark.sql.SQLContext = org.apache.spark.sql.hive.HiveContext@66ad7167
```

If the default sqlContext is not HiveContext, create it:

```
scala> val hiveContext = new org.apache.spark.sql.hive.HiveContext(sc)
```

You can now create a Data Frame df that points to a Hive external table over Oracle Data Pump files:

```
scala> val df = sqlContext.table("<hive external table>") <hive external table>:
org.apache.spark.sql.DataFrame = [ <column names> ]
```

Now you can access data via the data frame.

```
scala> df.count
scala> df.head
```

If a Hive external table had not been created and you only had the Oracle Data Pump files created by Copy to Hadoop, you can create the Hive external table from within Spark.

```
scala> sqlContext.sql("CREATE EXTERNAL TABLE <hive external table> ROW FORMAT SERDE
'oracle.hadoop.hive.datapump.DPSerDe' STORED AS INPUTFORMAT
'oracle.hadoop.hive.datapump.DPInputFormat' OUTPUTFORMAT
'org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat' LOCATION
'/user/oracle/oracle_warehouse/<hive database name>'")
```

### 3.1.9 Using Oracle SQL Developer with Copy to Hadoop

Oracle SQL Developer is a free, GUI-based development environment that provides easy to use tools for working Oracle Big Data Connectors, including Copy to Hadoop.

Using Oracle SQL Developer, you can copy data and create a new Hive table, or append data to an existing Hive external table that was created by Copy to Hadoop. In the GUI, you can initiate Copy to Hadoop in Oracle SQL Developer by right-clicking the Tables icon under any Hive schema. You can then append to an existing Hive external table by right-clicking the icon for that Hive table.

See Installing Oracle SQL Developer in this manual for instructions on where to obtain Oracle SQL Developer and how do the basic installation.



Oracle SQL Developer User's Guide for complete details on Oracle SQL Developer.



# 3.2 Storing Oracle Tablespaces in HDFS

You can store Oracle read-only tablespaces on HDFS and use Big Data SQL Smart Scan to off-load query processing of data stored in that tablespace to the Hadoop cluster. Big Data SQL Smart Scan performs data local processing - filtering query results on the Hadoop cluster prior to the return of the data to Oracle Database. In most circumstances, this can be a significant performance optimization. In addition to Smart Scan, querying tablespaces in HDFS also leverages native Oracle Database access structures and performance features. This includes features such as indexes, Hybrid Columnar Compression, Partition Pruning, and Oracle Database In-Memory.

Tables, partitions, and data in tablespaces in HDFS retain their original Oracle Database internal format. This is not a data dump. Unlike other means of accessing data in Hadoop (or other noSQL systems), you do not need to create Oracle External table. After copying the corresponding Oracle tablespaces to HDFS, you refer to the original Oracle table to access the data.

Permanent online, read only, and offline tablespaces (including ASM tablespaces) are eligible for the move to HDFS.



Since tablespaces allocated to HDFS are may not be altered, offline tablespaces must remain as offline. For offline tablespaces, then, what this feature provides is a hard backup into HDFS.

# 3.2.1 Advantages and Limitations of Tablespaces in HDFS

The following are some reasons to store Oracle Database tablespaces in HDFS.

- Because the data remains in Oracle Database internal format, I/O requires no resource-intensive datatype conversions.
- All Oracle Database performance optimizations such as indexing, Hybrid Columnar Compression, Partition Pruning, and Oracle Database In-Memory can be applied.
- Oracle user-based security is maintained. Other Oracle Database security features such as Oracle Data Redaction and ASO transparent encryption remain in force if enabled. In HDFS, tablespaces can be stored in zones under HDFS Transparent HDFS encryption.
- Query processing can be off-loaded. Oracle Big Data SQL Smart Scan is applied to Oracle Database tablespaces in HDFS. Typically, Smart Scan can provide a significant performance boost for queries. With Smart Scan, much of the query processing workload is off-loaded to the Oracle Big Data SQL server cells on the Hadoop cluster where the tablespaces reside. Smart Scan then performs predicate filtering in-place on the Hadoop nodes to eliminate irrelevant data so that only data that meets the query conditions is returned to the database tier for processing. Data movement and network traffic are reduced to the degree that smart scan predicate filtering can distill the dataset before returning it to the database.



- For each table in the tablespace, there is only a single object to manage the
  Oracle-internal table itself. To be accessible to Oracle Database, data stored in
  other file formats typically used in HDFS requires an overlay of an external table
  and a view.
- As is always the case with Oracle internal partitioning, partitioned tables and indexes can have partitions in different tablespaces some of which may be in Exadata, ZFSSA, and other storage devices. This feature adds HDFS as another storage option.

There are some constraints on using Oracle tablespaces in HDFS. As is the case with all data stored in HDFS, Oracle Database tables, partitions, and data stored in HDFS are immutable. Updates are done by deleting and replacing the data. This form of storage is best suited to off-loading tables and partitions for archival purposes. Also, with the exception of OD4H, data in Oracle tablespaces in HDFS is not accessible to other tools in the Hadoop environment, such as Spark, Oracle Big Data Discovery, and Oracle Big Data Spatial and Graph.

# 3.2.2 About Tablespaces in HDFS and Data Encryption

Oracle Database Tablespaces in HDFS can work with ASO ( Oracle Advanced Security) transparent table encryption as well as HDFS Transparent Encryption in HDFS.

#### **Tablespaces With Oracle Database ASO Encryption**

In Oracle Database, ASO transparent encryption may be enabled for a tablespace or objects within the tablespace. This encryption is retained if the tablespace is subsequently moved to HDFS. For queries against this data, the CELL\_OFFLOAD\_DECRYPTION setting determines whether Oracle Big Data SQL or Oracle Database decrypts the data.

- If CELL\_OFFLOAD\_DECRYPTION = TRUE, then the encryption keys are sent to the Oracle Big Data server cells in Hadoop and data is decrypted at the cells.
- If CELL\_OFFLOAD\_DECRYPTION = FALSE, encryption keys are not sent to the cells and therefore the cells cannot perform TDE decryption. The data is returned to Oracle Database for decryption.

The default value is TRUE.



In cases where <code>cell\_Offload\_decryption</code> is set to <code>False</code>, Smart Scan cannot read the encrypted data and is unable to provide the performance boost that results from the Hadoop-side filtering of the query result set. TDE Column Encryption prevents Smart Scan processing of the encrypted columns only. TDE Tablespace Encryption prevents Smart Scan processing of the entire tablespace.

#### **Tablespaces in HDFS Transparent Encryption Zones**

You can move Oracle Database tablespaces into zones under HDFS Transparent Encryption with no impact on query access or on the ability of Smart Scan to filter data.



# 3.2.3 Moving Tablespaces to HDFS

Oracle Big Data SQL provides two options for moving tablespaces from Oracle Database to the HDFS file system in Hadoop.

Using bds-copy-tbs-to-hdfs

The script bds-copy-tbs-to-hdfs.sh lets you select a preexisting tablespace in Oracle Database. The script automates the move of the selected tablespace to HDFS and performs necessary SQL ALTER operations and datafile permission changes for you. The DataNode where the tablespace is relocated is predetermined by the script. The script uses FUSE-DFS to move the datafiles from Oracle Database to the HDFS file system in the Hadoop cluster .

You can find bds-copy-tbs-to-hdfs.sh in the cluster installation directory - \$ORACLE\_HOME/BDSJaguar-3.2.0/<string identifer for the cluster>.

Manually Moving Tablespaces to HDFS

As an alternative to bds-copy-tbs-to-hdfs.sh, you can manually perform the steps to move the tablespaces to HDFS. You can either move an existing tablespace, or, create a new tablespace and selectively add tables and partitions that you want to off-load. In this case, you can set up either FUSE-DFS or an HDFS NFS gateway service to move the datafiles to HDFS.

The scripted method is more convenient. The manual method is somewhat more flexible. Both are supported.

### Before You Start:

As cited in the *Prerequisites* section of the installation guide, both methods require that the following RPMs are pre-installed:

- fuse
- fuse-libs

# yum -y install fuse fuse-libs

These RPMs are available in the Oracle public yum repository.

# 3.2.3.1 Using bds-copy-tbs-to-hdfs

On the Oracle Database server, you can use the script bds-copy-tbs-to-hdfs.sh to select and move Oracle tablespaces to HDFS. This script is in the bds-database-install directory that you extracted from the database installation bundle when you installed Oracle Big Data SQL.

### **Syntax**

bds-copy-tbs-to-hdfs.sh syntax is as follows:

```
bds-copy-tbs-to-hdfs.sh
bds-copy-tbs-to-hdfs.sh --install
bds-copy-tbs-to-hdfs.sh --uninstall
bds-copy-tbs-to-hdfs.sh --force-uninstall-script
```



```
bds-copy-tbs-to-hdfs.sh --tablespace=<tablespace name> [-pdb=<pluggable database
name>]
bds-copy-tbs-to-hdfs.sh --list=<tablespace name> [--pdb=<pluggable database name>]
bds-copy-tbs-to-hdfs.sh --show=<tablespace name> [--pdb=<pluggable database name>]
```

Additional command line parameters are described in the table below.

Table 3-3 bds-copy-tbs-to-hdfs.sh Parameter Options

Parameter List	Description
No parameters	Returns the FUSE-DFS status.
install	Installs the FUSE-DFS service. No action is taken if the service is already installed.
uninstall	Uninstalls the FUSE-DFS service and removes the mountpoint.
grid-home	Specifies the Oracle Grid home directory.
base-mountpoint	By default, the mountpoint is under /mnt. However, on some systems access to this directory is restricted. This parameter lets you specify an alternate location.
aux-run-mode	Because Oracle Big Data SQL is installed on the database side as a regular user (not a superuser), tasks that must be done as root and/or the Grid user require the installer to spawn shells to run other scripts under those accounts while bds-copy-tbs-to-hdfs.sh is paused. Theaux-run-mode parameter specifies a mode for running these auxiliary scripts.
	aux-run-mode= <mode></mode>
	Mode options are:
	<ul> <li>session – through a spawned session.</li> </ul>
	<ul> <li>su — as a substitute user.</li> </ul>
	<ul> <li>sudo — through sudo.</li> </ul>
	<ul> <li>ssh — through secure shell.</li> </ul>



Table 3-3 (Cont.) bds-copy-tbs-to-hdfs.sh Parameter Options

Parameter List	Description
force-uninstall-script	This option creates a secondary script that runs as root and forces the FUSE-DFS uninstall.

### **Caution:**

Limit use of this option to system recovery, an attempt to end a system hang, or other situations that may require removal of the **FUSE-DFS** service. Forcing the uninstall could potentially leave the database in an unstable state. The customer assumes responsibility for this choice. Warning message are displayed to remind you of the risk if you use this option.

tablespace= <tablespace name=""> [pdb=<pluggable database="" name="">]</pluggable></tablespace>	Moves the named tablespace in the named PDB to storage in HDFS on the Hadoop cluster. If there are no PDBs, then thepdb argument is discarded.
list= <tablespace name=""> [ pdb=<pluggable database="" name=""></pluggable></tablespace>	Lists tablespaces whose name equals or includes the name provided. Thepdb parameter is an optional scopelist=* returns all tablespacespdb=* returns matches for the tablespace name within all PDBs.
show= <tablespace name=""> [ pdb=<pluggable database="" name=""></pluggable></tablespace>	Shows tablespaces whose name equals or includes the name provided and are already moved to HDFS. Thepdb parameter is an optional scopeshow=* returns all tablespacespdb=* returns matches for the tablespace name within all PDBs.

### Usage

Use bds-copy-tbs-to-hdfs.sh to move a tablespace to HDFS as follows.



- 1. Log on as the <code>oracle</code> Linux user and cd to the <code>bds-database-install</code> directory where the database bundle was extracted. Find <code>bds-copy-tbs-to-hdfs.sh</code> in this directory.
- 2. Check that FUSE-DFS is installed.
  - \$ ./bds-copy-tbs-to-hdfs.sh
- Install the FUSE-DFS service (if it was not found in the previous check). This command will also start the FUSE-DFS the service.
  - \$ ./bds-copy-tbs-to-hdfs.sh --install

If this script does not find the mount point, it launches a secondary script. Run this script as root when prompted. It will set up the HDFS mount. You can run the secondary script in a separate session and then return to this session if you prefer.



### For RAC Databases: Install FUSE\_DFS on All Nodes:

On a RAC database, the script will prompt you that you must install FUSE-DFS on the other nodes of the database.

- 4. List the eligible tablespaces in a selected PDB or all PDBs. You can skip this step if you already know the tablespace name and location.
  - \$ ./bds-copy-tbs-to-hdfs.sh --list=mytablesapce --pdb=pdb1
- 5. Select a tablespace from the list and then, as oracle, run bds-copy-tbs-to-hdfs.sh again, but this time pass in the --tablespace parameter (and the --pdb parameter if specified). The script moves the tablespace to the HDFS file system.
  - \$ ./bds-copy-tbs-to-hdfs.sh --tablespace=mytablespace --pdb=pdb1

This command automatically makes the tablespace eligible for Smart Scan in HDFS. It does this in SQL by adding the "hdfs:" prefix to the datafile name in the tablespace definition. The rename changes the pointer in the database control file. It does not change the physical file name.



#### Tip:

If the datafiles are stored in ASM, the extraction will be made using RMAN. At this time, RMAN does not support a direct copy from ASM into HDFS. This will result in an error.

As workaround, you can use the <code>--staging-dir</code> parameter, which that enables you to do a two-stage copy – first to a file system directory and then into HDFS. The file system directory specified by <code>--staging-dir</code> must have sufficient space for the ASM datafile.

\$ ./bds-copy-tbs-to-hdfs.sh --tablespace=mytablespace --pdb=pdb1 --stagingdir=/home/user

For non-ASM datafiles, --staging-dir is ignored.



The tablespace should be back online and ready for access when you have completed this procedure.

### 3.2.3.2 Manually Moving Tablespaces to HDFS

As an alternative to bds-copy-tbs-to-hdfs.sh, you can use the following manual steps to move Oracle tablespaces to HDFS.



In the case of an ASM tablespace, you must first use RMAN or ASMCMD to copy the tablespace to the filesystem.

Oracle Big Data SQL includes FUSE-DFS and these instructions use it to connect to the HDFS file system. You could use an HDFS NFS gateway service instead. The documentation for your Hadoop distribution should provide the instructions for that method.

Perform all of these steps on the Oracle Database server. Run all Linux shell commands as root. For SQL commands, log on to the Oracle Database as the oracle user.

1. If FUSE-DFS is not installed or is not started, run bds-copy-tbs-to-hdfs.sh -- install. This script will install FUSE-DFS (if it's not already installed) and then start it.

The script will automatically create the mount point /mnt/fuse-<clustername>-hdfs.



The script bds-copy-tbs-to-hdfs.sh is compatible with FUSE-DFS 2.8 only.

In SQL, use CREATE TABLESPACE to create the tablespace. Store it in a local .dbf file. After this file is populated, you will move it to the Hadoop cluster. A single, bigfile tablespace is recommended.

### For example:

SQL> CREATE TABLESPACE movie\_cold\_hdfs DATAFILE '/u01/app/oracle/oradata/cdb/orcl/movie\_cold\_hdfs1.dbf' SIZE 100M reuse AUTOEXTEND ON nologging;

3. Use ALTER TABLE with the MOVE clause to move objects in the tablespace.

#### For example:

SQL> ALTER TABLE movie\_fact MOVE PARTITION 2010\_JAN TABLESPACE movie\_cold\_hdfs ONLINE UPDATE INDEXES;

You should check the current status of the objects to confirm the change. In this case, check which tablespace the partition belongs to.

SQL> SELECT table\_name, partition\_name, tablespace\_name FROM user\_tab\_partitions WHERE table\_name='MOVIE\_FACT';

4. Make the tablespace read only and take it offline.



```
SQL> ALTER TABLESPACE movie_cold_hdfs READ ONLY;
SQL> ALTER TABLESPACE movie_cold hdfs OFFLINE;
```

5. Copy the datafile to HDFS and then change the file permissions to read only.

```
hadoop fs -put /u01/app/oracle/oradata/cdb/orcl/movie_cold_hdfs1.dbf /user/oracle/tablespaces/
hadoop fs -chmod 440 /user/oracle/tablespaces/movie_cold_hdfs1.dbf
```

As a general security practice for Oracle Big Data SQL, apply appropriate HDFS file permissions to prevent unauthorized read/write access.

You may need to source <code>\$ORACLE\_HOME/bigdatasql/</code> hadoop\_<clustername>.env before running hadoop fs commands.

As an alternative, you could use the LINUX cp command to copy the files to FUSE.

Rename the datafiles, using ALTER TABLESPACE with the RENAME DATAFILE clause.

### Important:

Note the "hdfs:" prefix to the file path in the SQL example below. This is the keyword that tells Smart Scan that it should scan the file. Smart Scan also requires that the file is read only. The cluster name is optional. Also, before running the SQL statement below, the directory \$ORACLE\_HOME/dbs/hdfs:<clustername>/user/oracle/tablespaces should include the soft link movie\_cold\_hdfs1.dbf, pointing to /mnt/fuse-<clustername>-hdfs/user/oracle/tablespaces/movie\_cold\_hdfs1.dbf.

SQL> ALTER TABLESPACE movie\_cold\_hdfs RENAME DATAFILE '/u01/app/oracle/ oradata/cdb/orcl/movie\_cold\_hdfs1.dbf' TO 'hdfs:<clustername>/user/oracle/ tablespaces/movie\_cold\_hdfs1.dbf';

When you rename the datafile, only the pointer in the database control file changes. This procedure does not physically rename the datafile.

The tablespace must exist on a single cluster. If there are multiple datafiles, these must point to the same cluster.

7. Bring the tablespace back online and test it.

```
SQL> ALTER TABLESPACE movie_cold_hdfs ONLINE;
SQL> SELECT avg(rating) FROM movie_fact;
```

Below is the complete code example. In this case we move three partitions from local Oracle Database storage to the tablespace in HDFS.

```
mount hdfs
select * from dba_tablespaces;

CREATE TABLESPACE movie_cold_hdfs DATAFILE '/u01/app/oracle/oradata/cdb/orcl/
movie_cold_hdfs1.dbf' SIZE 100M reuse AUTOEXTEND ON nologging;

ALTER TABLE movie_fact
MOVE PARTITION 2010_JAN TABLESPACE movie_cold_hdfs ONLINE UPDATE INDEXES;
ALTER TABLE movie_fact
MOVE PARTITION 2010 FEB TABLESPACE movie cold hdfs ONLINE UPDATE INDEXES;
```



```
ALTER TABLE movie_fact

MOVE PARTITION 2010_MAR TABLESPACE movie_cold_hdfs ONLINE UPDATE INDEXES;

-- Check for the changes

SELECT table_name, partition_name, tablespace_name FROM user_tab_partitions WHERE table_name='MOVIE_FACT';

ALTER TABLESPACE movie_cold_hdfs READ ONLY;

ALTER TABLESPACE movie_cold_hdfs OFFLINE;

hadoop fs -put /u01/app/oracle/oradata/cdb/orcl/movie_cold_hdfs1.dbf /user/oracle/ tablespaces/
hadoop fs -chmod 444 /user/oracle/tablespaces/ movie_cold_hdfs1.dbf

ALTER TABLESPACE movie_cold_hdfs RENAME DATAFILE '/u01/app/oracle/oradata/cdb/orcl/ movie_cold_hdfs1.dbf' TO 'hdfs:hadoop_cl_1/user/oracle/tablespaces/ movie_cold_hdfs1.dbf';

ALTER TABLESPACE movie_cold_hdfs ONLINE;

-- Test
select avg(rating) from movie_fact;
```

# 3.2.4 Smart Scan for TableSpaces in HDFS

Smart Scan is an Oracle performance optimization that moves processing to the location where the data resides. In Big Data SQL, Smart Scan searches for datafiles whose path includes the "hdfs:" prefix. This prefix is the key that indicates the datafile is eligible for scanning.

After you have moved your tablespace data to HDFS and the tablespace and have prefixed the datafile path with the "hdfs:" tag, then queries that access the data in these files will leverage Big Data SQL Smart Scan by default. All of the Big Data SQL Smart Scan performance optimizations will apply. This greatly reduces the amount of data that moves from the storage tier to the database tier. These performance optimizations include:

- The massively parallel processing power of the Hadoop cluster is employed to filter data at its source.
- Storage Indexes can be leveraged to reduce the amount of data that is scanned.
- Data mining scoring can be off-loaded.
- Encrypted data scans can be off-loaded.

#### **Disabling or Enabling Smart Scan**

The initialization parameter \_CELL\_OFFLOAD\_HYBRID\_PROCESSING determines whether Smart Scan for HDFS is enabled or disabled. It is enabled by default.

To disable Smart Scan for tablespaces in HDFS do the following.

1. Set the parameter to FALSE in init or in a parameter file:

```
_CELL_OFFLOAD_HYBRID_PROCESSING=FALSE
```

The underscore prefix is required in this parameter name.

2. Restart the Oracle Database instance.

You can also make this change dynamically using the  ${\tt ALTER}$   ${\tt SYSTEM}$  directive in SQL. This does not require a restart.

SQL> alter system set \_cell\_offload\_hybrid\_processing=false;

One reason to turn off Smart Scan is if you need to move the Oracle tablespace datafiles out of HDFS and back to their original locations.

You can re-enable Smart Scan by resetting <code>\_CELL\_OFFLOAD\_HYBRID\_PROCESSING</code> to <code>TRUE</code>.



When  $_{\tt CELL\_OFFLOAD\_HYBRID\_PROCESSING}$  is set to False, Smart Scan is disabled for Oracle tablespaces residing in HDFS.



# Oracle Big Data SQL Reference

This chapter contains reference information for Oracle Big Data SQL:

- DBMS\_HADOOP PL/SQL Package
- DBMS\_BDSQL PL/SQL Package
- CREATE TABLE ACCESS PARAMETERS Clause
- Static Data Dictionary Views for Hive

# 4.1.1 DBMS HADOOP PL/SQL Package

The DBMS\_HADOOP package contains a function to generate the CREATE EXTERNAL TABLE DDL for a Hive table:

CREATE\_EXTDDL\_FOR\_HIVE

## 4.1.1.1 CREATE\_EXTDDL\_FOR\_HIVE

This function returns a SQL CREATE TABLE ORGANIZATION EXTERNAL statement for a Hive table. It uses the ORACLE\_HIVE access driver.

## **Syntax**

#### **Parameters**

Table 4-1 CREATE\_EXTDDL\_FOR\_HIVE Function Parameters

Parameter	Description
cluster_id	Hadoop cluster where the Hive metastore is located
db_name	Name of the Hive database
hive_table_name	Name of the Hive table
hive_partition	Whether the table is partitioned (TRUE) or not (FALSE)
table_name	Name of the Oracle external table to be created. It cannot already exist.



Table 4-1 (Cont.) CREATE\_EXTDDL\_FOR\_HIVE Function Parameters

Parameter	Description	
perform_ddl	Whether to execute the generated CREATE TABLE statement (TRUE) or just return the text of the command (FALSE).	
	Do not execute the command automatically if you want to review or modify it.	
text_of_ddl	The generated CREATE TABLE ORGANIZATION EXTERNAL statement.	

## **Usage Notes**

The Oracle Database system must be configured for Oracle Big Data SQL. See "About Oracle Big Data SQL on the Database Server (Oracle Exadata Machine or Other)".

The data type conversions are based on the default mappings between Hive data types and Oracle data types. See "About Data Type Conversions".

## 4.1.1.1.1 Example

The following query returns the CREATE EXTERNAL TABLE DDL for my\_hive\_table from the default Hive database. The connection to Hive is established using the configuration files in the ORACLE\_BIGDATA\_CONFIG directory, which identify the location of the HADOOP1 cluster.

```
DECLARE
   DDLtxt VARCHAR2(4000);
BEGIN
   dbms_hadoop.create_extddl_for_hive(
        CLUSTER_ID=>'hadoop1',
        DB_NAME=>'default',
        HIVE_TABLE_NAME=>'my_hive_table',
        HIVE_PARTITION=>FALSE,
        TABLE_NAME=>'my_xt_oracle',
        PERFORM_DDL=>FALSE,
        TEXT_OF_DDL=>DDLtxt
   );
   dbms_output.put_line(DDLtxt);
END;
//
```

The query returns the text of the following SQL command:



```
PARALLEL 2 REJECT LIMIT UNLIMITED
```

# 4.1.2 DBMS\_BDSQL PL/SQL Package

The DBMS\_BDSQL PL/SQL package contains procedures to add and remove a user map.

- ADD\_USER\_MAP
- REMOVE\_USER\_MAP

In previous releases of Oracle Big Data SQL, all queries against Hadoop and Hive data are executed as the <code>oracle</code> user and there is no option to change users. Although <code>oracle</code> is still the underlying user in all cases, you can now use Multi-User Authorization (based on Hadoop Secure Impersonation) to direct the oracle account to execute tasks on behalf of other designated users. This enables HDFS data access based on the user that is currently executing the query, rather than the singular <code>oracle</code> user.

The DBMS\_BDSQL package enables you to provide rules for identifying the currently connected user and to map the connected user to the user that is impersonated. Because there are numerous ways in which users can connect to Oracle Database, this user may be a database user, a user sourced from LDAP, from Kerberos, and so forth. Authorization rules on the files apply for that user and audits will reflect that user as well.



Grant the new BDSQL\_ADMIN role to designated administrators in order to allow them to invoke these procedures.

# 4.1.2.1 ADD\_USER\_MAP

Use the ADD\_USER\_MAP procedure to specify the rules for identifying the actual user who is running the query.

At query execution time, the database performs a lookup on the BDSQL\_USER\_MAP table to determine the current database user (current\_database\_user). It then uses the syscontext\_namespace and syscontext\_parm\_hadoop\_user parameters to identify the actual user.

#### **Syntax**



Table 4-2 ADD\_USER\_MAP Parameters

Parameter	Description
cluster_name	The name of the Hadoop cluster where the map will be applied. [DEFAULT] as cluster name designates the default cluster.
current_database_user	The current effective database user. This is what Oracle uses to check for authority. A value of '*' indicates that this row to be used if no other rows fit the criteria. There is no default and the value may not be NULL.
syscontext_namespace	Note that for the Oracle USERENV namespace, the only allowed values are GLOBAL_UID, CLIENT_IDENTIFIER, and AUTHENTICATED_IDENTITY.  If your Oracle Database installation uses Kerberos credentials, SSL, Active Directory, or LDAP for authentication, it is likely that your Hadoop system uses the same authentication framework. In that case, AUTHENTICATED_IDENTITY must be specified. This identifier only includes the username. The domain segment of the credential is truncated, as is the cluster name if included. For example, the username dirkrb will be used for authorization on the Hadoop cluster as the authenticated identity of dirkrb@HQ.TEST1.DBSEC2008.COM.
syscontext_parm_hadoop_user	The Hadoop user that will impersonate the current database user.

## Note:

The values for current\_database\_user and syscontext\_parm\_hadoop\_user can be the single asterisk character (\*) or any string that meets the requirements of Oracle simple\_sql\_name assertion:

- The name must begin with an alphabetic character. It may contain alphanumeric characters as well as the characters \_, \$, and # in the second and subsequent character positions.
- Quoted SQL names are also allowed.
- Quoted names must be enclosed in double quotes.
- Quoted names allow any characters between the quotes.
- Quotes inside the name are represented by two quote characters in a row, for example, "a name with "" inside" is a valid quoted name.
- The input parameter may have any number of leading and/or trailing white space characters.





SYS CONTEXT in the Database SQL Language Reference.

# 4.1.2.2 REMOVE\_USER\_MAP

Use REMOVE\_USER\_MAP to remove a row from BDSQL\_USER\_MAP table. This disassociates a specific Oracle Database user from specific Hadoop user.

## **Syntax**

```
procedure REMOVE_USER_MAP (
   cluster_name IN VARCHAR2 DEFAULT '[DEFAULT]',
   current_database_user IN VARCHAR2 NOT NULL
);
```



The reference page for ADD\_USER\_MAP describes the cluster\_name and current\_database\_user parameters.

# 4.1.2.3 Multi-User Authorization Security Table

SYS.BDSQL\_USER\_MAP is the multi-user authorization security table. Use the procedures ADD\_USER\_MAP and REMOVE\_USER\_MAP to update this table.

The primary key is (cluster\_name, current\_database\_user).

Table 4-3 SYS.BDSQL\_USER\_MAP

Column	Datatype	Description
cluster_name	varchar2	Name of the Hadoop cluster. The default is [DEFAULT].
current_databas e_user	varchar2	The current effective database user (no default, not NULL). Oracle uses this column to check for the authorization rule that corresponds to the given Oracle Database user. A value of '*' in a row is a directive to use this row if no other rows fit the criteria.
syscontext_name space	varchar2	This is the optional specification for the Oracle SYS_CONTEXT namespace. if customer security is set up. Note that for the Oracle USERENV namespace, the only allowed values are: 'GLOBAL_UID','CLIENT_IDENTIFIER', 'AUTHENTICATED_IDENTITY'.



Table 4-3 (Cont.) SYS.BDSQL USER MAP

Column	Datatype	Description
syscontext_parm _hadoop_user	varchar2	This column value has alternate interpretations.  If syscontext_namespace has a value, thensyscontext_parm_hadoop_user refers to the parameter that is specific to syscontext_namespace.
		However, when the value is '*', this is a directive to use the value of current_database_user for impersonation. The syscontext_namespace column must be NULL in this case.  • If syscontext_namespace is NULL, then syscontext_parm_hadoop_user contains the Hadoop user who is impersonated prior to HDFS files access.

Here a customer is using Active Directory, Kerberos, SSL, or LDAP for logon authentication against the database. AUTHENTICATED\_IDENTITY is specified in this case because customer uses the same Active Directory framework for Hadoop user management.

The example below is similar to running the following SQL query for the currently connected user:

```
select sys_context('USERENV', 'AUTHENTICATED_IDENTITY') from dual;
```

In this example, only the username (without the "@<domain>" segment) is used for authorization on the Hadoop cluster. There may also be cases where the format of AUTHENTICATED\_IDENTITY iS <username>/<cluster>@<domain\_name>.

cluster_name	current_database_u	syscontext_namesp	syscontext_parm_ha
	ser	ace	doop_user
[DEFAULT]	*	USERENV	AUTHENTICATED_ID ENTITY

In this example, "HRAPP" is an HR Application that always connects to the database using the HRAPP database user and then programmatically sets the application user through the <code>DBMS\_SESSION.SET\_IDENTIFIER</code> procedure. There are number of "lightweight" users who are designated with <code>CLIENT\_IDENTIFIER</code> (as in <code>sys\_context('USERENV', 'CLIENT\_IDENTIFIER')</code> [DEFAULT] \* USERENV GLOBAL\_UID, which is similar to running select <code>sys\_context('USERENV', 'CLIENT\_IDENTIFIER')</code> from <code>dual; )</code>.

The current database has other effective users who are enterprise users with logons managed by Oracle Internet Directory for Enterprise User Security. In these cases, the <code>GLOBAL\_UID</code> is used for Hadoop impersonation.

cluster_name	current_database_u ser	syscontext_namesp ace	syscontext_parm_ha doop_user
[DEFAULT]	HRAPP	USERENV	CLIENT_IDENTIFIER
[DEFAULT]	*	USERENV	GLOBAL_UID

In this example, BIAPP is a business intelligence application whose own context is its username. For customers using the application, their designated ID is used for Hadoop access. In other words, when the effective current user is 'BIAPP', we use



 $sys\_context("BIVPD","USERID")$  for the impersonation. For the rest of the users, we simply designate [DEFAULT] \* \* in order use their current database username for the impersonation.

cluster_name	current_database_u ser	syscontext_namesp ace	syscontext_parm_ha doop_user
[DEFAULT]	BIAPP	BIVPD	USERID
[DEFAULT]	*		*

In this example, the Oracle username SCOTT is impersonated by the hdpusr1 Hadoop user for HDFS access. The user ADAM is impersonated by hdpusr2 for HDFS access.

All other users have more limited access, so we use a syscontext\_namespace value of 'lowprivuser' to designate these users.

cluster_name	current_database_u ser	syscontext_namesp ace	syscontext_parm_ha doop_user
hadoop_cl_1	SCOTT		hdpusr1
hadoop_cl_1	ADAM	lowprivuser	hdpusr2
hadoop_cl_1	*		

## 4.1.3 CREATE TABLE ACCESS PARAMETERS Clause

This section describes the properties that you use when creating an external table that uses the <code>ORACLE\_HDFS</code> or <code>ORACLE\_HIVE</code> access drivers. In a <code>CREATE TABLE ORGANIZATION</code> <code>EXTERNAL</code> statement, specify the parameters in the <code>Opaque\_format\_spec</code> clause of <code>ACCESS PARAMETERS</code>.

This section contains the following topics:

- Syntax Rules for Specifying Properties
- ORACLE\_HDFS Access Parameters
- ORACLE HIVE Access Parameters
- Alphabetical list of properties

# 4.1.3.1 Syntax Rules for Specifying Properties

The properties are set using keyword-value pairs in the SQL CREATE TABLE ACCESS PARAMETERS clause and in the configuration files. The syntax must obey these rules:

The format of each keyword-value pair is a *keyword*, a colon or equal sign, and a *value*. The following are valid keyword-value pairs:

keyword=value keyword:value

The value is everything from the first non-whitespace character after the separator to the end of the line. Whitespace between the separator and the value is ignored. Trailing whitespace for the value is retained.

A property definition can be on one line or multiple lines.



- A line terminator is a line feed, a carriage return, or a carriage return followed by line feeds.
- When a property definition spans multiple lines, then precede the line terminators with a backslash (escape character), except on the last line. In this example, the value of the Keyword1 property is Value part 1 Value part 2 Value part 3.

- You can create a logical line by stripping each physical line of leading whitespace and concatenating the lines. The parser extracts the property names and values from the logical line.
- You can embed special characters in a property name or property value by preceding a character with a backslash (escape character), indicating the substitution. Table 4-4 describes the special characters.

**Table 4-4 Special Characters in Properties** 

Escape Sequence	Character
\b	Backspace (\u0008)
\t	Horizontal tab (\u0009)
\n	Line feed (\u000a)
\f	Form feed (\u000c)
\r	Carriage return (\u000d)
\"	Double quote (\u0022)
\'	Single quote (\u0027)
\\	Backslash (\u005c)
	When multiple backslashes are at the end of the line, the parser continues the value to the next line only for an odd number of backslashes.
\uxxxx	2-byte, big-endian, Unicode code point.
	When a character requires two code points (4 bytes), the parser expects $\setminus u$ for the second code point.

## 4.1.3.2 ORACLE\_HDFS Access Parameters

The access parameters for the <code>ORACLE\_HDFS</code> access driver provide the metadata needed to locate the data in HDFS and generate a Hive table over it.

## 4.1.3.2.1 Default Parameter Settings for ORACLE\_HDFS

If you omit all access parameters from the CREATE TABLE statement, then ORACLE\_HDFS uses the following default values:

```
com.oracle.bigdata.rowformat=DELIMITED
com.oracle.bigdata.fileformat=TEXTFILE
com.oracle.bigdata.overflow={"action":"error"}
com.oracle.bigdata.erroropt={"action":"setnull"}
```



## 4.1.3.2.2 Optional Parameter Settings for ORACLE HDFS

 ${\tt ORACLE\_HDFS} \ \ \textbf{supports the following optional com.oracle.bigdata} \ \ \textbf{parameters, which you} \\ \textbf{can specify in the } \ \ \textbf{opaque\_format\_spec clause} :$ 

- com.oracle.bigdata.colmap
- com.oracle.bigdata.erroropt
- com.oracle.bigdata.fields
- · com.oracle.bigdata.fileformat
- com.oracle.bigdata.log.exec
- com.oracle.bigdata.log.qc
- · com.oracle.bigdata.overflow
- com.oracle.bigdata.rowformat

Example 4-1 shows a CREATE TABLE statement in which multiple access parameters are set.

## Example 4-1 Setting Multiple Access Parameters for ORACLE\_HDFS

```
CREATE TABLE ORDER (CUST_NUM VARCHAR2(10),
                    ORDER_NUM VARCHAR2(20),
                    ORDER_DATE DATE,
                    ITEM_CNT NUMBER,
                    DESCRIPTION VARCHAR2(100),
                    ORDER_TOTAL (NUMBER8,2)) ORGANIZATION EXTERNAL
             (TYPE ORACLE HDFS
ACCESS PARAMETERS (
 com.oracle.bigdata.fields: (CUST_NUM,
                              ORDER NUM,
                              ORDER_DATE,
                              ORDER_LINE_ITEM_COUNT, \
                              DESCRIPTION,
                              ORDER_TOTAL)
  com.oracle.bigdata.colMap:
                                 {"col":"item_cnt", \
                                  "field": "order_line_item_count" }
  com.oracle.bigdata.overflow:
                                 {"action":"TRUNCATE", \
                                  "col": "DESCRIPTION" }
 com.oracle.bigdata.errorOpt: [{"action":"replace", \
                                  "value": "INVALID NUM", \
                                  "col":["CUST_NUM","ORDER_NUM"]} , \
                                 {"action": "reject", \
                                   "col": "ORDER_TOTAL } ]
LOCATION ("hdfs:/usr/cust/summary/*"));
```

## 4.1.3.3 ORACLE\_HIVE Access Parameters

ORACLE\_HIVE retrieves metadata about external data sources from the Hive catalog. The default mapping of Hive data to columns in the external table are usually appropriate. However, some circumstances require special parameter settings, or you might want to override the default values for reasons of your own.



## 4.1.3.3.1 Default Parameter Settings for ORACLE HIVE

If you omit all access parameters from the CREATE TABLE statement, then ORACLE\_HIVE uses the following default values:

```
com.oracle.bigdata.tablename=name of external table
com.oracle.bigdata.overflow={"action":"error"}
com.oracle.bigdata.erroropt={"action":"setnull"}
```

## 4.1.3.3.2 Optional Parameter Values for ORACLE\_HIVE

ORACLE\_HIVE supports the following optional com.oracle.bigdata parameters, which you can specify in the opaque\_format\_spec clause:

- com.oracle.bigdata.colmap
- · com.oracle.bigdata.erroropt
- com.oracle.bigdata.log.exec
- com.oracle.bigdata.log.gc
- com.oracle.bigdata.overflow
- com.oracle.bigdata.tablename

Example 4-2 shows a CREATE TABLE Statement in which multiple access parameters are set.

## Example 4-2 Setting Multiple Access Parameters for ORACLE\_HIVE

```
CREATE TABLE ORDER (cust_num VARCHAR2(10),
                    order_num VARCHAR2(20),
                    order_date DATE,
                    item_cnt NUMBER,
                    description VARCHAR2(100),
                    order_total (NUMBER8,2)) ORGANIZATION EXTERNAL
(TYPE oracle_hive
ACCESS PARAMETERS (
 com.oracle.bigdata.tableName: order_db.order_summary
 com.oracle.bigdata.colMap: {"col":"ITEM_CNT", \
                                  "field": "order_line_item_count" }
  com.oracle.bigdata.overflow: {"action":"ERROR", \
                                  "col": "DESCRIPTION" }
 com.oracle.bigdata.errorOpt: [{"action":"replace", \
                                   "value":"INV_NUM" , \
                                   "col":["CUST_NUM","ORDER_NUM"]} ,\
                                  {"action": "reject", \
                                   "col":"ORDER_TOTAL}]
));
```

# 4.1.3.4 com.oracle.bigdata.buffersize

Sets the buffer size in kilobytes for large record reads. Set this value if you need to read records that are greater than the default buffer size.

#### **Default Value**

1000 KB



#### **Syntax**

com.oracle.bigdata.buffersize: n

## **Example**

The following example sets the buffer size to 100 MB:

com.oracle.bigdata.buffersize: 100000

# 4.1.3.5 com.oracle.bigdata.colmap

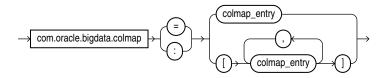
Maps a column in the source data to a column in the Oracle external table. You can define one or multiple pairs of column mappings. Use this property when the source field names exceed the maximum length of Oracle column names, or when you want to use different column names in the external table.

#### **Default Value**

A column in the external table with the same name as the Hive column

## **Syntax**

A JSON document with the keyword-value pairs is shown in the following diagram: colmap:



colmap entry:



#### **Semantics**

"col":name

"col": The keyword must be lowercase and enclosed in quotation marks.

name: The name of a column in the Oracle external table. It is case sensitive and must be enclosed in quotation marks.

"field":name

"field": The keyword must be lowercase and enclosed in quotation marks.

name: The name of a field in the data source. It is not case sensitive, but it must be enclosed in quotation marks. See "Syntax Rules for Specifying Properties".



#### **Examples**

This example maps a Hive column named <code>ORDER\_LINE\_ITEM\_COUNT</code> to an Oracle column named <code>ITEM\_CNT</code>:

The following example shows the mapping of multiple columns.

```
com.oracle.bigdata.colmap:[{"col":"KOL1", "field":"PROJECT_NAME"},
{ "col":"KOL2","field":"wsdl_name"},{"col":"KOL3", "field":"method"}]
```

## 4.1.3.6 com.oracle.bigdata.datamode

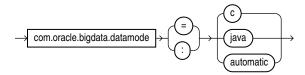
Specifies the method that SmartScan uses to scan a Hadoop data source. The method can make a significant difference in performance.

#### **Default Value**

automatic

#### **Syntax**

A JSON document with the keyword-value pairs shown in the following diagram: datamode:



#### **Semantics**

automatic

Automatically selects the appropriate mode, based on the metadata. It selects  ${\tt c}$  mode if possible, or  ${\tt java}$  mode if the data contains formats that are not supported by  ${\tt c}$  mode.

С

Uses Java to read the file buffers, but C code to process the data and convert it to Oracle format. Specify this mode for delimited data.

If the data contains formats that the C code does not support, then it returns an error.

java

Uses the Java SerDes and InputFormats to process the data and convert it to Oracle format. Specify this mode for Parquet, RCFile, and other data formats that require a SerDe.

# 4.1.3.7 com.oracle.bigdata.erroropt

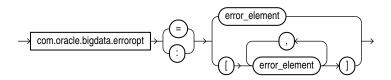
Describes how to handle errors that occur while the value of a column is calculated.

## **Default Value**

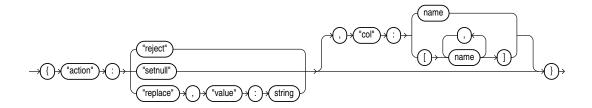
{ "action": "setnull" }

## **Syntax**

A JSON document with the keyword-value pairs is shown in the following diagram: erroropt:



error\_element:



#### **Semantics**

The "action", "reject", "setnull", "replace", "value", and "col" keywords must be lowercase and enclosed in quotation marks. See "Syntax Rules for Specifying Properties".

"action":value

value: One of these keywords:

- "reject": Does not load any rows.
- "setnull": Sets the column to NULL.
- "replace": Sets the column to the specified value.

"value":string

 $\it string$ : Replaces a bad value in the external table. It must be enclosed in quotation marks.

"col":name

name: Identifies a column in an external table. The column name is case sensitive, must be enclosed in quotation marks, and can be listed only once.

## **Example**

This example sets the value of the <code>cust\_num</code> or <code>order\_num</code> columns to <code>invalid</code> if the Hive value causes an error. For any other columns, an error just causes the Hive value to be rejected.



# 4.1.3.8 com.oracle.bigdata.fields

Lists the field names and data types of the data source.

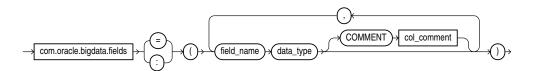
#### **Default Value**

Not defined

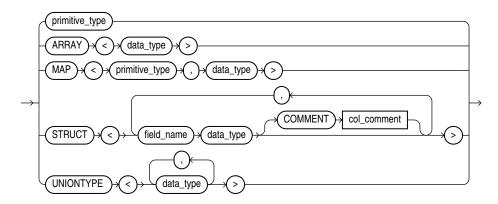
## **Syntax**

A JSON document with the keyword-value pairs is shown in the following diagram:

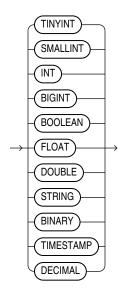
fields:



data\_type:



primitive\_type:



#### **Semantics**

The syntax is the same as a field list for a Hive table. If you split the field list across multiple lines, you must use a backslash to escape the new line characters.

field\_name

The name of the Hive field. Use only alphanumeric characters and underscores (\_). The maximum length is 128 characters. Field names are case-insensitive.

data\_type

The data type of the Hive field. Optional; the default is  ${\tt STRING}$ . The character set must be UTF8.

The data type can be complex or primitive:

**Hive Complex Data Types** 

ARRAY: Indexable list

MAP: Key-value tuples

STRUCT: List of elements

UNIONTYPE: Multiple data types

**Hive Primitive Data Types** 

INT: 4 byte integer

BIGINT: 8 byte integer

SMALLINT: 2 byte integer

TINYINT: 1 byte integer

BOOLEAN: TRUE **Of** FALSE

FLOAT: single precision

DOUBLE: double precision

STRING: character sequence

## See Also:

"Data Types" in the Apache Hive Language Manual at

https://cwiki.apache.org/confluence/display/Hive/LanguageManual+Types

## COMMENT col\_comment

A string literal enclosed in single quotation marks, which is stored as metadata for the Hive table (comment property of TBLPROPERTIES).

# 4.1.3.9 com.oracle.bigdata.fileformat

Describes the row format of the data source, based on the  ${\tt ROW}$  FORMAT clause for a Hive table generated by  ${\tt ORACLE\_HDFS}$ .

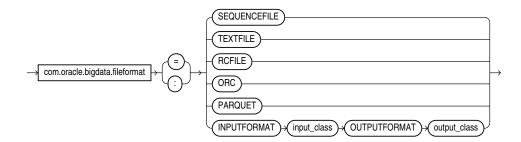
#### **Default Value**

TEXTFILE

#### **Syntax**

A JSON document with the keyword-value pairs is shown in the following diagram.

#### fileformat:



#### **Semantics**

**ORC** 

Optimized row columnar file format

**PARQUET** 

Column-oriented, binary file format

**RCFILE** 

Record columnar file format

SEQUENCEFILE

Compressed file format

**TEXTFILE** 

Plain text file format



#### **INPUTFORMAT**

Identifies a Java class that can extract records from the data file.

#### **OUTPUTFORMAT**

Identifies a Java class that can format the output records in the desired format

## 4.1.3.10 com.oracle.bigdata.log.exec

Specifies how the access driver generates log files generated by the C code for a query, when it is running as parallel processes on CDH.

The access driver does not create or write log files when executing on a Hadoop cluster node; the parallel query processes write them. The log files from the Java code are controlled by <code>log4j</code> properties, which are specified in the configuration file or the access parameters. See "bigdata-log4j.properties".

#### **Default Value**

Not defined (no logging)

#### **Syntax**

[directory\_object:]file\_name\_template

#### **Semantics**

directory\_object

The Oracle directory object for the HDFS path on the Hadoop cluster where the log file is created.

file name template

A string used to generate file names. Table 4-4 describes the optional variables that you can use in the template.

Table 4-5 Variables for com.oracle.bigdata.log.exec

Variable	Value
%p	Operating system process identifier (PID)
%a	A number that uniquely identifies the process.
88	A percent sign (%)

#### **Example**

The following example generates log file names that include the PID and a unique number, such as xtlogp\_hive14\_3413\_57:

com.oracle.bigdata.log.exec= xtlogp\_hive14\_%p\_%a

## 4.1.3.11 com.oracle.bigdata.log.qc

Specifies how the access driver generates log files for a query.



## **Default Value**

Not defined (no logging)

#### **Syntax**

[directory\_object:]file\_name\_template

#### **Semantics**

directory\_object

Name of an Oracle directory object that points to the path where the log files are written. If this value is omitted, then the logs are written to the default directory for the external table.

file\_name\_template

A string used to generate file names. Table 4-6 describes the optional variables that you can use in the string.

Table 4-6 Variables for com.oracle.bigdata.log.qc

Variable	Value
%p	Operating system process identifier (PID)
%%	A percent sign (%)

## **Example**

This example creates log file names that include the PID and a percent sign, such as xtlogp\_hive213459\_%:

com.oracle.bigdata.log.qc= xtlogp\_hive21%p\_%%

# 4.1.3.12 com.oracle.bigdata.overflow

Describes how to handle string data that is too long for the columns in the external table. The data source can be character or binary. For Hive, the data source can also be STRUCT, UNIONTYPES, MAP, OF ARRAY.

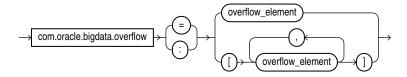
#### **Default Value**

{ "action": "error" }

#### **Syntax**

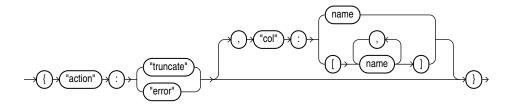
A JSON document with the keyword-value pairs is shown in the following diagram:

overflow ::=





## overflow\_element ::=



#### **Semantics**

The "action", "truncate", "error", and "col" tags must be lowercase and enclosed in quotation marks. See "Syntax Rules for Specifying Properties".

"action":value

The value of "action" can be one of the following keywords:

- truncate: Shortens the data to fit the column.
- error: Throws an error. The com.oracle.bigdata.erroropt property controls the result of the error.

"col":name

name: Identifies a column in the external table. The name is case sensitive and must be enclosed in quotation marks.

### **Example**

This example truncates the source data for the DESCRIPTION column, if it exceeds the column width:

```
com.oracle.bigdata.overflow={"action":"truncate", \
                              "col": "DESCRIPTION" }
```

# 4.1.3.13 com.oracle.bigdata.rowformat

Provides the information the access driver needs to extract fields from the records in a file.



## Important:

The com.oracle.bigdata.rowformat is unrelated to the access parameter syntax of traditional external tables that use "type ORACLE LOADER." There are keywords such as fields, terminated, and others that appear in both clauses, but the commonality in naming is coincidental and does not imply common functionality. The com.oracle.bigdata.rowformat access parameter is passed without change to the default Hive serde. The Hive serde to extract columns from rows is deliberately limited. Complex cases are handled by specialized serdes.

#### **Default Value**

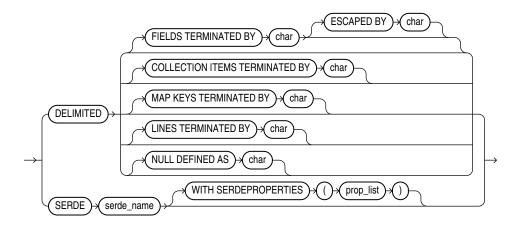
DELIMITED



#### **Syntax**

A JSON document with the keyword-value pairs is shown in the following diagram.

rowformat:



#### **Semantics**

#### **DELIMITED**

Describes the characters used to delimit the fields in a record:

- FIELDS TERMINATED BY: The character that delimits every field in the record. The
  optional ESCAPED BY character precedes the delimit character when it appears
  within a field value.
- COLLECTION ITEMS TERMINATED BY: The character that marks the end of an array element. Used when a column is a collection or a nested record. In this case the resulting value will be a JSON array.
- MAP KEYS TERMINATED BY: The character that marks the end of an entry in a MAP field. Used when a column is a collection or a nested record. The resulting value is a JSON object.
- LINES TERMINATED BY: The character that marks the end of a record.
- NULL DEFINED AS: The character that indicates a null value.

#### **SERDE**

Identifies a SerDe that can parse the data and any properties of the SerDe that the access driver might need.

#### **Example**

This example specifies a SerDe for an Avro container file:

```
com.oracle.bigdata.rowformat:
    SERDE'org.apache.hadoop.hive.serde2.avro.AvroSerDe'
```

The next example specifies a SerDe for a file containing regular expressions:

```
com.oracle.bigdata.rowformat=\
    SERDE 'org.apache.hadoop.hive.contrib.serde2.RegexSerDe' \
```



```
WITH SERDEPROPERTIES \
("input.regex" = "(\\\d{6}) (\\\d{5}) (.{29}) .*")
```

## 4.1.3.14 com.oracle.bigdata.tablename

Identifies the Hive table that contains the source data.

#### **Default Value**

DEFAULT.external\_table\_name

#### **Syntax**

[hive\_database\_name.]table\_name

#### **Semantics**

The maximum length of *hive\_database\_name* and *table\_name* is 128 UTF-8 characters (512 bytes).

hive\_database\_name

The Hive database where the source data resides. DEFAULT is the name of the initial Hive database.

table name

The Hive table with the data. If you omit *table\_name*, then <code>ORACLE\_HIVE</code> searches for a Hive table with the same name as the external table. Table names are case-insensitive.

## **Example**

This setting indicates that the source data is in a table named <code>ORDER\_SUMMARY</code> in the <code>Hive ORDER\_DB</code> database:

com.oracle.bigdata.tablename ORDER\_DB.ORDER\_SUMMARY

# 4.1.4 Static Data Dictionary Views for Hive

The Oracle Database catalog contains several static data dictionary views for Hive tables. You can query these data dictionary views to discover information about the Hive tables that you can access.

For you to access any Hive databases from Oracle Database, you must have read privileges on the  $\mbox{oracle\_BigData\_config}$  directory object.

- ALL HIVE DATABASES
- ALL\_HIVE\_TABLES
- ALL\_HIVE\_COLUMNS
- DBA\_HIVE\_DATABASES
- DBA HIVE TABLES
- DBA\_HIVE\_COLUMNS
- USER\_HIVE\_DATABASES



- USER\_HIVE\_TABLES
- USER HIVE COLUMNS

# 4.1.4.1 ALL\_HIVE\_DATABASES

 ${\tt ALL\_HIVE\_DATABASES}$  describes all databases in the Hive metastore accessible to the current user.

#### **Related Views**

- DBA\_HIVE\_DATABASES describes all the databases in the Hive metastore.
- USER\_HIVE\_DATABASES describes the databases in the Hive metastore owned by the current user.

Column	Datatype	NULL	Description
CLUSTER_ID	VARCHAR2(4000)	NOT NULL	Hadoop cluster where the Hive metastore is located
DATABASE_NAME	VARCHAR2(4000)	NOT NULL	Hive database name
DESCRIPTION	VARCHAR2(4000)		Hive database description
DB_LOCATION	VARCHAR2(4000)	NOT NULL	
HIVE_URI	VARCHAR2(4000)		Hive database URI

## See Also:

- "DBA\_HIVE\_DATABASES"
- "USER HIVE DATABASES"

# 4.1.4.2 ALL\_HIVE\_TABLES

 ${\tt ALL\_HIVE\_TABLES}$  describes all tables in the Hive metastore accessible to the current user.

The Oracle Big Data SQL configuration must identify the default Hive database for the current user. The current user must also have READ privileges on the ORA\_BIGSQL\_CONFIG database directory. See "About the bigdata\_config Directory".

#### **Related Views**

- DBA\_HIVE\_TABLES describes all tables in the Hive metastore.
- USER\_HIVE\_TABLES describes the tables in the database owned by the current user in the Hive metastore.

Column	Datatype	NULL	Description
CLUSTER_ID	VARCHAR2(4000)	NOT NULL	Hadoop cluster where the Hive metastore is located
DATABASE_NAME	VARCHAR2(4000)	NOT NULL	Name of the Hive database



Column	Datatype	NULL	Description
TABLE_NAME	VARCHAR2(4000)	NOT NULL	Name of the Hive table
LOCATION	VARCHAR2(4000)		
NO_OF_COLS	NUMBER		Number of columns in the Hive table
CREATION_TIME	DATE		Time when the table was created
LAST_ACCESSED_TIME	DATE		Time of most recent access
OWNER	VARCHAR2(4000)		Owner of the Hive table
TABLE_TYPE	VARCHAR2(4000)	NOT NULL	Type of Hive table, such as external or managed
PARTITIONED	VARCHAR2(4000)		Whether the table is partitioned (YES) or not ( $\mathbb{N}$ O)
NO_OF_PART_KEYS	NUMBER		Number of partitions
INPUT_FORMAT	VARCHAR2(4000)		Input format
OUTPUT_FORMAT	VARCHAR2(4000)		Output format
SERIALIZATION	VARCHAR2(4000)		SerDe serialization information
COMPRESSED	NUMBER		Whether the table is compressed (YES) or not ( $\mathbb{N}$ 0)
HIVE_URI	VARCHAR2(4000)		Hive database URI

## See Also:

- "DBA\_HIVE\_TABLES"
- "USER\_HIVE\_TABLES"

# 4.1.4.3 ALL\_HIVE\_COLUMNS

 ${\tt ALL\_HIVE\_COLUMNS}$  describes the columns of all Hive tables accessible to the current user.

The Oracle Big Data SQL configuration must identify the default Hive database for the current user. The current user must also have READ privileges on the ORA\_BIGSQL\_CONFIG database directory. See "About the bigdata\_config Directory".

## **Related Views**

- DBA\_HIVE\_COLUMNS describes the columns of all tables in the Hive metastore.
- USER\_HIVE\_COLUMNS describes the columns of the tables in the Hive database owned by the current user.

Column	Datatype	NULL	Description
CLUSTER_ID	VARCHAR2(4000)	NOT NULL	Hadoop cluster where the Hive metastore is located



Column	Datatype	NULL	Description
DATABASE_NAME	VARCHAR2(4000)	NOT NULL	Name of the Hive database; if blank, then the default database
TABLE_NAME	VARCHAR2(4000)	NOT NULL	Name of the Hive table
COLUMN_NAME	VARCHAR2(4000)	NOT NULL	Name of the Hive column
HIVE_COLUMN_TYPE	VARCHAR2(4000)	NOT NULL	Data type of the Hive column
ORACLE_COLUMN_TYPE	VARCHAR2(4000)	NOT NULL	Oracle data type equivalent to Hive data type
LOCATION	VARCHAR2(4000)		
OWNER	VARCHAR2(4000)		Owner of the Hive table
CREATION_TIME	DATE		Time when the table was created
HIVE_URI	VARCHAR2(4000)		Hive database URI

## See Also:

- "DBA\_HIVE\_COLUMNS"
- "USER\_HIVE\_COLUMNS"

# 4.1.4.4 DBA\_HIVE\_DATABASES

DBA\_HIVE\_DATABASES describes all the databases in the Hive metastore. Its columns are the same as those in ALL\_HIVE\_DATABASES.

See Also:
"ALL\_HIVE\_DATABASES"

# 4.1.4.5 DBA\_HIVE\_TABLES

 ${\tt DBA\_HIVE\_TABLES}$  describes all tables in the Hive metastore. Its columns are the same as those in  ${\tt ALL\_HIVE\_TABLES}$ .

The Oracle Big Data SQL configuration must identify the default Hive database for the current user. See "About the bigdata\_config Directory".

See Also:

"ALL\_HIVE\_TABLES"

## 4.1.4.6 DBA\_HIVE\_COLUMNS

DBA\_HIVE\_COLUMNS describes the columns of all tables in the Hive metastore. Its columns are the same as those in ALL\_HIVE\_COLUMNS.

```
See Also:

"ALL_HIVE_COLUMNS"
```

# 4.1.4.7 USER\_HIVE\_DATABASES

USER\_HIVE\_DATABASES describes the databases in the Hive metastore owned by the current user. Its columns (except for owner) are the same as those in ALL HIVE DATABASES.

```
See Also:

"ALL_HIVE_DATABASES"
```

## 4.1.4.8 USER HIVE TABLES

USER\_HIVE\_TABLES describes the tables in the database owned by the current user in the Hive metastore. Its columns (except for OWNER) are the same as those in ALL\_HIVE\_TABLES.

The Oracle Big Data SQL configuration must identify the default Hive database for the current user. The current user must also have READ privileges on the ORA\_BIGSQL\_CONFIG database directory. See "About the bigdata\_config Directory".

```
See Also:

"ALL_HIVE_TABLES"
```

# 4.1.4.9 USER\_HIVE\_COLUMNS

USER\_HIVE\_COLUMNS describes the columns of the tables in the Hive database owned by the current user. Its columns (except for OWNER) are the same as those in ALL\_HIVE\_COLUMNS.

The Oracle Big Data SQL configuration must identify the default Hive database for the current user. The current user must also have READ privileges on the ORA\_BIGSQL\_CONFIG database directory. See "About the bigdata\_config Directory".

**✓** See Also:

"ALL\_HIVE\_COLUMNS"



# **Appendices**

C

# Using mtactl to Manage the MTA extproc

The multithreaded agent control utility (mtactl) enables Oracle Big Data SQL users to start, stop, and configure the MTA (Multi-Threaded Agent) extproc in both Oracle Clusterware Ready Service (CRS) and non-CRS Oracle Database environments.



In non-CRS environments, customers must run  ${\tt mtactl}$  in order to start the MTA extproc.

## **Usage**

In this usage description,  $mta\_sid$  is the SID that a given multithreaded extproc agent services.

```
mtactl {start|restart|stop|status|delete|show|bdsql} <mta_sid>
mtactl unset <parameter> <mta_sid>
mtactl set <parameter> <parameter_value> <mta_sid>
mtactl -help
mtactl <command> -help
```



## **Multithreaded Agent Control Utility Commands**

Table C-1 mtactl Commands

Comm and	Full Syntax	Description
start	mtactl start <mta_sid></mta_sid>	Start an MTA extproc for this SID, with the existing init parameter values stored in the repository. Use the default values if the repository does not exist.



If you used Oracle Big Data SQL 3.1, be aware that the behavior of restart and start are now reversed from what they were in 3.1 start now uses init values from the repositor y if available .restart always uses the default values.

restar mtactl start <mta\_sid>
t
stop mtactl stop <mta\_sid>
status mtactl status <mta\_sid>

Clean up the repository and restart the MTA extproc agent for the SID, with the default values.

Stop the MTA extproc agent that services the given SID.

Display status for the MTA extproc that services the given SID.



Table C-1 (Cont.) mtactl Commands

Comm and	Full Syntax	Description
delete	<pre>mtactl delete <mta_sid></mta_sid></pre>	Clean up the repository for the given SID.
show	<pre>mtactl show <mta_sid></mta_sid></pre>	Display the init parameters for the MTA extproc that services the given SID.
bdsql	<pre>mtactl bdsql <mta_sid></mta_sid></pre>	Display additional operations. These are for setting up the MTA extproc for use with Oracle Big Data SQL.
set	<pre>mtactl set <init parameter=""> <value> <mta_sid></mta_sid></value></init></pre>	Set the init parameters for the MTA extproc that services the given SID. Supported parameters are:
		<pre>max_dispatchers tcp_dispatchers max_task_threads max_sessions listener_address</pre>
unset	<pre>mtactl unset <init parameter=""> <mta_sid></mta_sid></init></pre>	Unset the init parameters in the repository for the MTA extproc that services the given SID.

## **Examples**

 $\$  mtactl start BDSQL\_hadoop\_cl\_1 //note: using existing or default init parameter values

```
$ mtactl delete BDSQL_hadoop_cl_1
```



<sup>\$</sup> mtactl set max\_sessions 200 BDSQL\_hadoop\_cl\_1

<sup>\$</sup> mtactl set max\_dispatchers 5 BDSQL\_hadoop\_cl\_1

<sup>\$</sup> mtactl set max\_task\_threads 5 BDSQL\_hadoop\_cl\_1

<sup>\$</sup> mtactl set listener\_address "(ADDRESS=(PROTOCOL=ipc)(KEY=crs))" BDSQL\_hadoop\_cl\_1

<sup>\$</sup> mtactl start BDSQL\_hadoop\_cl\_1 (note: use customized parameter values)

<sup>\$</sup> mtactl restart BDSQL\_hadoop\_cl\_1 //note: using default init parameter values

F

# Change History for Previous Releases

The following are changes in previous versions of the product.

# F.1 Changes in Oracle Big Data SQL 3.1

The following are changes in Oracle Big Data SQL Release 3.1.

#### **Expanded Deployments**

Release 3.1 broadens support for Oracle Big Data SQL connectivity between Oracle Engineered Systems and commodity servers.

In earlier Oracle Big Data SQL releases, the following Oracle Database/Hadoop connections are possible:

- Oracle Exadata Database Machine to Oracle Big Data Appliance.
- Oracle Database on commodity servers to commodity Hadoop systems.

As of Release 3.1, Oracle Big Data SQL supports all of the following Oracle Database/ Hadoop system connections:

- Oracle Database on commodity servers to Oracle Big Data Appliance.
- Oracle Database on commodity servers to commodity Hadoop systems.
- Oracle Exadata Database Machine to Oracle Big Data Appliance.
- Oracle Exadata Database Machine to commodity Hadoop systems.

The phrase "Oracle Database on commodity servers" refers to non-Exadata Linux systems that are officially-supported as Oracle Database platforms. "Commodity Hadoop systems" refers to Hortonworks HDP systems or to Cloudera CDH-based systems other than Oracle Big Data Appliance. In all cases, Oracle Database servers and Hadoop systems must meet the prerequisites identified in the Oracle Big Data SQL Master Compatibility Matrix (Doc ID 2119369.1 in My Oracle Support).

## **Oracle SPARC SuperCluster Support**

Release 3.1 provides support for Oracle SPARC SuperCluster, with certain limitations:

- Ethernet connections between BDA (and commodity Hadoop systems) and SPARC SuperCluster are not supported.
- Oracle Database Tablepaces in HDFS (with Smart Scan technology) is not supported on this platform.

#### **Ethernet Option for Connections to the Exadata Database Machine**

The preferred method of connecting Oracle Big Data Appliance and Oracle Exadata Database for any purpose is through InfiniBand. Previous releases of Oracle Big Data SQL have required InfiniBand for these connections. In Release 3.1, Ethernet networking between the Exadata Database Machine and Oracle Big Data Appliance is



now supported. This enables you to use Oracle Big Data SQL with these two Engineered Systems in environments where InfiniBand is not feasible, such as when the two systems are geographically distant from each other.

Release 3.1 also enables Ethernet connections between commodity Hadoop systems and the Oracle Exadata Database Machine.

Oracle Big Data SQL connectivity between commodity Hadoop systems and commodity Oracle Database servers has been Ethernet-based throughout previous releases.

Note that Ethernet connections between Oracle Big Data Appliance (or commodity Hadoop systems) and Oracle SPARC SuperCluster are not supported at this time.

# Simplified Deployment on Oracle Database Servers – Oracle Grid Infrastructure is now Optional

In previous releases of Oracle Big Data SQL, Oracle Grid Infrastructure is a prerequisite of the installation for all Oracle Database servers, including standalone servers that are not part of an Oracle RAC system. In Release 3.1, you have the option to install Oracle Big Data SQL on servers where Oracle Grid Infrastructure is not present. Note that in these cases, the installer makes some configuration file changes that require a restart of Oracle Database.

## Unified Platform Support in the Oracle Big Data SQL Installer

Previous Oracle Big Data SQL releases included two separate installation procedures – one for Oracle Engineered Systems and another for commodity servers. In Release 3.1, you use the same installation process for both Oracle and non-Oracle platforms. This is also true for maintenance. For all of the supported Hadoop/Oracle Database combinations there is a uniform set of steps to update the Oracle Big Data SQL configuration when there are changes to the Hadoop cluster or Oracle Database server.

# New Features to Simplify ILM – Oracle Database Tablespaces in HDFS (With Smart Scan Technology)

Oracle Database ILM (Information Lifecycle Management) can now be extended to use Hadoop to store read-only Oracle Database tablespaces. When you move tablespaces from Oracle Database to HDFS, the tables, partitions, and data retain their original Oracle Database internal format, remain accessible to queries, and support the full range of Oracle Database performance optimizations and security features, including the following:

- Smart Scan for HDFS, which enables off-load of query processing to Oracle Big Data SQL on the Hadoop cluster. Smart Scan also provides filtering of query results in Hadoop prior to the return of the data to Oracle Database. In most circumstances, this can be a significant performance optimization. Indexing, Hybrid Columnar Compression, Partition Pruning, and Oracle Database In-Memory are also supported.
- Oracle Advanced Security Option (ASO) Transparent Encryption and Data Redaction.

Tablespaces stored in HDFS are read-only, therefore this storage is best suited to data archiving.

See Section 3.2 in the Oracle Big Data SQL User's Guide for details.



#### **Enhancements in Oracle Shell for Hadoop Loaders 1.2**

Oracle Shell for Hadoop Loaders (OHSH) is an intuitive command line tool for data migration. You can set up resources to connect to Hive, HDFS, or Oracle Database, and then access each of these data sources through the uniform OHSH interface. OHSH is one of the ways to use Copy to Hadoop. Copy to Hadoop users can download OHSH from OTN.

#### OHSH 1.2 includes the following changes:

- On-disk logging of load operations in the \$HOME/.ohsh shadow directory.
- The ability to minimize output when doing load commands. (See the help command for set outputlevel.)
- Loading Hive tables from Oracle tables not living in the oracle user's schema.
- Wallet and TNS usage by OHSH relies on the setting of environmental variables
   WALLET\_LOCATION and TNS\_ADMIN. The set threadmin and set walletlocation commands are no longer supported.

In addition, you no longer set <code>HIVEO\_URL</code> to the fully-qualified URL of remote HiveServer2 in order to create a <code>%hiveO</code> resource. In OHSH 1.2, set the environmental variable <code>HS2\_HOST\_PORT</code> in <code>bin/ohsh</code>, which is the <code><hostname>:<port> pair of HiveServer2</code>.

#### **Enhancements to Copy To Hadoop**

The new method directcopy is added to Copy to Hadoop.

This is a direct, single-step method of copying data from Oracle Database to HDFS. See *Using Copy to Hadoop to do Direct Copies* in the *Oracle Big Data SQL User's Guide* for more information.

#### Granting Access – Users Now Require the BDSQL\_USER Role

In prior Oracle Big Data SQL releases, all users are granted Big Data SQL access implicitly. Release 3.1 introduces the BDSQL\_USER role. Users requiring Oracle Big Data SQL access must be granted this role explicitly.

You must also now grant read privileges on the BigDataSQL configuration directory object.

For example, to grant access to user1:

```
SQL> grant BDSQL_USER to user1;
SQL> grant read on directory ORACLE_BIGDATA_CONFIG to user1;
```

#### Installation Instructions Moved to Oracle Big Data SQL Installation Guide

The *Oracle Big Data SQL Installation Guide* provides instructions how to install and uninstall the software. In releases prior to Oracle Big Data Appliance 3.1, installation instructions are in the user's guide.

# F.2 Changes in Oracle Big Data SQL 3.0.1

The following are changes in Oracle Big Data SQL Release 3.0.1.



In addition to continued support for Oracle Big Data Appliance and the Exadata Database Machine, Release 3.0.1 includes further improvements to Big Data SQL support for platforms other than Oracle Engineered Systems. These currently include Hortonworks Data Platform (HDP) and Cloudera's Distribution Including Apache Hadoop (CDH) running on non-Big Data Appliance commodity hardware where these Hadoop systems are used in conjunction with Oracle Database running on non-Exadata servers.

#### **Support for Newer CDH and HDP Platforms**

Oracle Big Data SQL 3.0.1 runs on Hortonworks Hadoop Data Platform (HDP) 2.3 – 2.4.x and Cloudera's Distribution Including Apache Hadoop (CDH) 5.4.7 – 5.7. CDH 5.4.7 support includes both commodity CDH systems and Oracle Big Data Appliance 4.3.

## **Improved Installation Automation**

The installation on HDP and CDH now automatically saves the YARN resource management configuration for later restoration.

The oracle user and oinstall group setup are now managed by the installation program.

#### Uniform Uninstall Procedure for HDP and CDH

The uninstall procedure is now the same for both HDP and CDH systems.

#### Other Changes for HDP and CDH

- The default JDK is upgraded to jdk-8u66-linux-x64
- bds-database-create-bundle.sh includes a new --jdk-tar-path parameter for specifying a JDK other than the default.
- bds-database-install.sh includes two new parameters:
  - --cell-config-dir--root-script-only
- The installation process now creates and populates cellinit.ora if the file does not exist or is empty.

#### **Oracle Shell for Hadoop Loaders**

Oracle Shell for Hadoop Loaders is a helper shell that provides a simple-to-use command line interface to Copy to Hadoop, Oracle Loader for Hadoop, and Oracle SQL Connector for HDFS.

#### **Support for Oracle Big Data Appliance 4.5**

Oracle Big Data SQL Release 3.0.1 is a mammoth installation option in Oracle Big Data Appliance 4.5. You can also enable Oracle Big Data SQL Release 3.0.1 at any time after the mammoth installation by using the bdacli utility.



## Note:

In Release 3.0.1, Oracle Big Data SQL does not yet support connections between Oracle Big Data Appliance and non-Exadata Oracle Database servers, or between non-Big Data Appliance Hadoop systems and the Exadata Database Machine.



A

# Manual Steps for Using Copy to Hadoop for Staged Copies

To manually copy data from Oracle Database to Hadoop using Copy to Hadoop, take the following steps:

 On the Oracle Database server, connect to Oracle Database and generate Data Pump format files containing the table data and metadata.

See "Generating the Data Pump Files".

2. Copy the files to HDFS on the Hadoop cluster.

See "Copying the Files to HDFS".

Connect to Apache Hive and create an external table from the files.

See "Creating a Hive Table".

4. Query this Hive table the same as you would any other Hive table.

# A.1 Generating the Data Pump Files

The SQL CREATE TABLE statement has a clause specifically for creating external tables, in which you specify the ORACLE\_DATAPUMP access driver. The information that you provide in this clause enables the access driver to generate a Data Pump format file that contains the data and metadata from the Oracle database table.

This section contains the following topics:

- About Data Pump Format Files
- Identifying the Target Directory
- About the CREATE TABLE Syntax
- Copying the Files to HDFS

# A.1.1 About Data Pump Format Files

Data Pump files are typically used to move data and metadata from one database to another. Copy to Hadoop uses this file format to copy data from an Oracle database to HDFS.

To generate Data Pump format files, you create an external table from an existing Oracle table. An **external table** in Oracle Database is an object that identifies and describes the location of data outside of a database. External tables use **access drivers** to parse and format the data. For Copy to Hadoop, you use the ORACLE\_DATAPUMP access driver. It copies the data and metadata from internal Oracle tables and populates the Data Pump format files of the external table.

# A.1.2 Identifying the Target Directory

You must have read and write access to a database directory in Oracle Database. Only Oracle Database users with the CREATE ANY DIRECTORY system privilege can create directories.

This example creates a database directory named EXPORTDIR that points to the / exportdir directory on the Oracle Database server (Oracle Exadata Database Machine or other):

```
SQL> CREATE DIRECTORY exportdir AS '/exportdir';
```

# A.1.3 About the CREATE TABLE Syntax

The following is the basic syntax of the CREATE TABLE statement for Data Pump format files:

```
CREATE TABLE table_name
ORGANIZATION EXTERNAL (
    TYPE oracle_datapump
DEFAULT DIRECTORY database_directory
LOCATION ('filename1.dmp','filename2.dmp'...)
) PARALLEL n
AS SELECT * FROM tablename;
```

#### **DEFAULT DIRECTORY**

Identifies the database directory that you created for this purpose. See "Identifying the Target Directory".

## **LOCATION**

Lists the names of the Data Pump files to be created. The number of names should match the degree of parallelism (DOP) specified by the PARALLEL clause. Otherwise, the DOP drops to the number of files.

The number of files and the degree of parallelism affect the performance of Oracle Database when generating the Data Pump format files. They do not affect querying performance in Hive.

## **PARALLEL**

Sets the degree of parallelism (DOP). Use the maximum number that your Oracle DBA permits you to use. By default the DOP is 1, which is serial processing. Larger numbers enable parallel processing.

## **AS SELECT**

Use the full SQL SELECT syntax for this clause. It is not restricted. The *tablename* identifies the Oracle table to be copied to HDFS.



For descriptions of these parameters:

- Oracle Database SQL Language Reference
- Oracle Database Utilities



# A.2 Copying the Files to HDFS

The Oracle Big Data SQL installation installs Hadoop client files on the Oracle Datatabase server (Oracle Exadata Database Machine or other). The Hadoop client installation enables you to use Hadoop commands to copy the Data Pump files to HDFS. You must have write privileges on the HDFS directory.

To copy the dmp files into HDFS, use the hadoop fs -put command. This example copies the files into the HDFS customers directory owned by the oracle user:

\$ hadoop fs -put customers\*.dmp /user/oracle/customers

# A.3 Creating a Hive Table

To provide access to the data in the Data Pump files, you create a Hive external table over the Data Pump files. Copy to Hadoop provides SerDes that enable Hive to read the files. These SerDes are read only, so you cannot use them to write to the files.



Apache Hive Language Manual DDL at

https://cwiki.apache.org/confluence/display/Hive/LanguageManual +DDL#LanguageManualDDL-Create/Drop/TruncateTable

## A.3.1 About Hive External Tables

For external tables, Hive loads the table metadata into its metastore. The data remains in its original location, which you identify in the LOCATION clause. If you drop an external table using a HiveQL DROP TABLE statement, then only the metadata is discarded, while the external data remains unchanged. In this respect, Hive handles external tables in fundamentally the same way as Oracle Database.

External tables support data sources that are shared by multiple programs. In this case, you use Oracle Database to update the data and then generate a new file. You can replace the old HDFS files with the updated files, while leaving the Hive metadata intact.

The following is the basic syntax of a Hive CREATE TABLE statement for creating a Hive external table for use with a Data Pump format file:

```
CREATE EXTERNAL TABLE tablename

ROW FORMAT

SERDE 'oracle.hadoop.hive.datapump.DPSerDe'

STORED AS

INPUTFORMAT 'oracle.hadoop.hive.datapump.DPInputFormat'

OUTPUTFORMAT 'org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat'

LOCATION 'hdfs_directory'
```



# A.4 Example Using the Sample Schemas

This example shows all steps in the process of creating a Hive table from an Oracle table using Copy to Hadoop.

# A.4.1 About the Sample Data

The Oracle tables are from the Sales History (SH) sample schema. The CUSTOMERS table provides extensive information about individual customers, including names, addresses, telephone numbers, birth dates, and credit limits. The COUNTRIES table provides a list of countries, and identifies regions and subregions.

This query shows a small selection of data in the CUSTOMERS table:

```
SELECT cust_first_name first_name,
   cust_last_name last_name,
   cust_gender gender,
   cust_year_of_birth birth
FROM customers
ORDER BY cust_city, last_name
FETCH FIRST 10 ROWS ONLY;
```

## The query returns the following rows:

FIRST_NAME	LAST_NAME	GENDER	BIRTH
Lise	Abbey	F	1963
Lotus	Alden	M	1958
Emmanuel	Aubrey	M	1933
Phil	Ball	M	1956
Valentina	Bardwell	F	1965
Lolita	Barkley	F	1966
Heloise	Barnes	M	1980
Royden	Barrett	M	1937
Gilbert	Braun	M	1984
Portia	Capp	F	1948

To reproduce this example, install the sample schemas in Oracle Database and connect as the  ${\tt SH}$  user.



*Oracle Database Sample Schemas* for descriptions of the tables and installation instructions for the schemas.

# A.4.2 Creating the EXPDIR Database Directory

These SQL statements create a local database directory named EXPDIR and grant access to the SH user:

```
SQL> CREATE DIRECTORY expdir AS '/expdir'; Directory created.
```



```
SQL> GRANT READ, WRITE ON DIRECTORY expdir TO SH; Grant succeeded.
```

# A.4.3 Creating Data Pump Format Files for Customer Data

The following examples show how to create the Data Pump files and check their contents.

Copy to Hadoop supports only the syntax shown in the examples. Data pump files created with the Export utility or Oracle Data Pump are not compatible.

## A.4.3.1 CREATE TABLE Example With a Simple SELECT Statement

This example shows a very simple SQL command for creating a Data Pump format file from the CUSTOMERS table. It selects the entire table and generates a single output file named customers.dmp in the local /expdir directory.

```
CREATE TABLE export_customers
ORGANIZATION EXTERNAL
(
TYPE oracle_datapump
DEFAULT DIRECTORY expdir
LOCATION('customers.dmp')
)
AS SELECT * FROM customers;
```

# A.4.3.2 CREATE TABLE Example With a More Complex SQL SELECT Statement

The next example shows more complexity in the syntax. It joins the CUSTOMERS and COUNTRIES tables on the COUNTRY\_ID columns to provide the country names. It also limits the rows to customers in the Americas. The command generates two output files in parallel, named americas1.dmp and americas2.dmp, in the local /expdir directory.

```
CREATE TABLE export_americas
  ORGANIZATION EXTERNAL
  TYPE oracle_datapump
  DEFAULT DIRECTORY expdir
  LOCATION('americas1.dmp', 'americas2.dmp')
  PARALLEL 2
AS SELECT a.cust_first_name first_name,
  a.cust_last_name last_name,
  a.cust_gender gender,
  a.cust_year_of_birth birth,
  a.cust_email email,
  a.cust_postal_code postal_code,
  b.country_name country
FROM customers a,
    countries b
WHERE a.country id=b.country id AND
     b.country_region='Americas'
ORDER BY a.country_id, a.cust_postal_code;
```



# A.4.4 Verifying the Contents of the Data Files

You can check the content of the output data files before copying them to Hadoop. The previous CREATE TABLE statement created an external table named EXPORT\_AMERICAS, which you can describe and query the same as any other table.

The DESCRIBE statement shows the selection of columns and the modified names:

# SQL> DESCRIBE export\_americas; Name Null? Type FIRST\_NAME NOT NULL VARCHAR2(20) LAST\_NAME NOT NULL VARCHAR2(40) GENDER NOT NULL CHAR(1) BIRTH NOT NULL NUMBER(4) EMAIL VARCHAR2(50) POSTAL\_CODE NOT NULL VARCHAR2(10) COUNTRY NOT NULL VARCHAR2(40)

A SELECT statement like the following shows a sample of the data:

```
SELECT first_name, last_name, gender, birth, country
FROM export_americas
WHERE birth > 1985
ORDER BY last_name
FETCH FIRST 5 ROWS ONLY;
```

FIRST_NAME	LAST_NAME	GENDER	BIRTH	COUNTRY
Opal	Aaron	M	1990	United States of America
KaKit	Abeles	M	1986	United States of America
Mitchel	Alambarati	M	1987	Canada
Jade	Anderson	M	1986	United States of America
Roderica	Austin	M	1986	United States of America

## A.4.5 Copying the Files into Hadoop

The following commands list the files in the local expdir directory, create a Hadoop subdirectory named customers, and copy the files to it. The user is connected to the Hadoop cluster (Oracle Big Data Appliance or other) as the oracle user.

```
$ cd /expdir
$ ls americas*.dmp
americas1.dmp americas2.dmp
$ hadoop fs -mkdir customers
$ hadoop fs -put *.dmp customers
$ hadoop fs -ls customers
Found 2 items
-rw-r--r-- 1 oracle oracle 798720 2014-10-13 17:04 customers/americas1.dmp
-rw-r--r-- 1 oracle oracle 954368 2014-10-13 17:04 customers/americas2.dmp
```

## A.4.6 Creating a Hive External Table

This HiveQL statement creates an external table using the Copy to Hadoop SerDes. The LOCATION clause identifies the full path to the Hadoop directory containing the Data Pump files:

```
CREATE EXTERNAL TABLE customers

ROW FORMAT SERDE 'oracle.hadoop.hive.datapump.DPSerDe'

STORED AS

INPUTFORMAT 'oracle.hadoop.hive.datapump.DPInputFormat'

OUTPUTFORMAT 'org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat'

LOCATION '/user/oracle/customers';
```

The describe command shows the columns of the customers external table.

## hive> DESCRIBE customers;

ΟK

011		
first_name	varchar(20)	from deserializer
last_name	varchar(40)	from deserializer
gender	char(1)	from deserializer
birth	int	from deserializer
email	varchar(50)	from deserializer
postal_code	varchar(10)	from deserializer
country	varchar(40)	from deserializer



B

# Using Copy to Hadoop With Direct Copy

Copy to Hadoop with the direct copy option copies data from an Oracle Database table directly to Oracle Datapump files stored in HDFS.

Copy to Hadoop simplifies the data copy because it does not require intermediate storage on the database server. The rest of steps needed to make the data accessible to the Hadoop ecosystem such as creating Hive external table to access the copied data and running Hive queries is common to both copy options. (stage and direct).

The intended audience for this section are power users of Hadoop with specialized requirements . All other users should use the Oracle Shell for Hadoop Loader (OHSH) CLI for Copy to Hadoop operations. See Using Oracle Shell for Hadoop Loaders With Copy to Hadoop.

# B.1 Manual Steps for Using Copy to Hadoop for Direct Copies

Follow these steps.

## **Getting Started**

- 1. First confirm that Copy to Hadoop is installed and configured.
- 2. Ensure that the user account has sufficient privileges to copy a database table. (See Table Access Requirements for Copy to Hadoop.)
- 3. Make sure that the table contains supported column types. (See About Column Mappings and Data Type Conversions.)
- Log in to either a node in the Hadoop cluster or a system set up as a Hadoop client for the cluster.
- 5. If you are connecting to a secure cluster then run kinit to authenticate the user, as described in the Oracle Big Data Appliance Software User's Guide.
- 6. Run the Copy To Hadoop job using direct copy. See "Running the Copy to Hadoop Job for Direct Copy" below.
- 7. After the job succeeds, check the jobReport.log file in the \_ctoh subdirectory of the output directory in HDFS. Check that the RowCount listed in the log file is equal to the number of rows in the database table.
- 8. Connect to Apache Hive and create an external table from the Data Pump files. (See Creating a Hive Table.)

## Running the Copy to Hadoop Job (Using Direct Copy)

Set the environment variables required by Copy to Hadoop.
 Locate the installation directory of Copy to Hadoop and set the CP2HADOOP\_HOME
Bash shell variable. For example:



# export CP2HADOOP\_HOME="/opt/oracle/orahivedp-3.1.0"

## Add the Copy to Hadoop JARs to HADOOP\_CLASSPATH. For example:

# export HADOOP\_CLASSPATH="\${CP2HADOOP\_HOME}/jlib/\*:\${HADOOP\_CLASSPATH}"



## Tip:

When using Copy to Hadoop, you should always list \$CP2HADOOP\_HOME/jlib/\* first in HADOOP\_CLASSPATH. Another way to avoid JAR conflicts is to define an appropriately ordered HADOOP\_CLASSPATH within a script that uses it.

#### 2. Run the Job.

## This is the command syntax:

# hadoop jar \${CP2HADOOP\_HOME}/jlib/orahivedp.jar oracle.hadoop.ctoh.CtohDriver \ [-D <configuration-property>=<configuration-value>]+

## **Example 1: Running the Job on a Secure Cluster Using Oracle Wallet**

```
# hadoop jar ${CP2HADOOP_HOME}/jlib/orahivedp.jar oracle.hadoop.ctoh.CtohDriver \
-D oracle.hadoop.ctoh.connection.tnsEntry=<my-oracle-tns-entry> \
-D oracle.hadoop.ctoh.connection.walletLoc=<local-oracle-wallet-dir> \
-D oracle.hadoop.ctoh.connection.tnsAdmin=<local-oracle-wallet-dir> \
-D oracle.hadoop.ctoh.connection.clusterWalletLoc=<oracle-wallet-dir-on-hadoop-
cluster> \
-D oracle.hadoop.ctoh.connection.clusterTnsAdmin=<oracle-wallet-dir-on-hadoop-
cluster> \
-D mapreduce.output.fileoutputformat.outputdir=<mytab-hdfs-output-dir> \
-D oracle.hadoop.ctoh.splitterType="BLOCK_SPLITTER" \
-D oracle.hadoop.ctoh.table=<dbSchema.dbTable> \
-D oracle.hadoop.ctoh.maxSplits=10
```

## Example 2: Running the Job on A Unsecured Hadoop Cluster (for Demo **Purposes Only**)

```
# hadoop jar ${CP2HADOOP_HOME}/jlib/orahivedp.jar oracle.hadoop.ctoh.CtohDriver \
-D oracle.hadoop.ctoh.jdbc.url="jdbc:oracle:thin:@myhost:1521/myservice" \
-D oracle.hadoop.ctoh.connection.username="myuser" \
-D oracle.hadoop.ctoh.connection.password="mypassword" \
-D mapreduce.output.fileoutputformat.outputdir="mytable_output_dir" \
-D oracle.hadoop.ctoh.splitterType="BLOCK_SPLITTER" \
-D oracle.hadoop.ctoh.table="otherUser.mytable" \
-D oracle.hadoop.ctoh.maxSplits=10
```

## **Performance Tuning Tips**

You can control the degree of parallelism of the Copy to Hadoop job by specifying the number of map processes using the oracle.hadoop.ctoh.maxSplits property. The higher the number of map processes, the higher the parallelism. Note that each process connects to the database, so this value also determines the number of simultaneous connections to the database. Typically, a number such as 64 works well.

## **Required Configuration Properties**

See Section "Configuration Properties Reference" for more information on these and other properties.



- oracle.hadoop.ctoh.table
- mapreduce.output.fileoutputformat.outputdir
- oracle.hadoop.ctoh.maxSplits
- oracle.hadoop.ctoh.splitterType

## **Connection Properties for a secure Hadoop cluster using Oracle Wallet:**

- oracle.hadoop.ctoh.connection.walletLoc
- oracle.hadoop.ctoh.connection.tnsAdmin
- oracle.hadoop.ctoh.connection.tnsEntry
- The following properties are also required if the Oracle Wallet directory on the Hadoop cluster is different from the directory on the Hadoop client:
  - oracle.hadoop.ctoh.connection.clusterWalletLoc
  - oracle.hadoop.ctoh.connection.clusterTnsAdmin

# Connection Properties for Unsecured Hadoop clusters (for Demo Purposes Only):

For demo purposes, use the following properties in place of the properties used with secured clusters.

- oracle.hadoop.ctoh.connection.username
- oracle.hadoop.ctoh.connection.password
- oracle.hadoop.ctoh.jdbc.url

## An Incremental Copy using Copy to Hadoop

To incrementally copy data from the same Oracle table to a pre-existing destination directory in HDFS, the following additional properties are required. (This configuration assumes that a Copy To Hadoop job was run initially to copy data from an Oracle Database table to datapump files in an HDFS directory.)

- oracle.hadoop.ctoh.whereClause
- oracle.hadoop.ctoh.datapump.output
- oracle.hadoop.ctoh.datapump.basename

oracle.hadoop.ctoh.datapump.output specifies a preexisting HDFS location that contains the datapump files from a previous run of Copy To Hadoop.

oracle.hadoop.ctoh.whereClause identifies the subset of rows to be copied from the Oracle table for the incremental load.

oracle.hadoop.ctoh.datapump.basename specifies a unique prefix for the datapump files. This property is used to generate unique datapump file names to prevent file name collisions during an incremental load.



# B.2 Copy to Hadoop Property Reference

This reference describes customer-accessible properties of Copy to Hadoop.

**Copy to Hadoop Configuration Property Reference (for Direct Copy)** 

Property	Description
oracle.hadoop.ctoh.home	Type: String
	<b>Default Value</b> : Value of the CP2HADOOP_HOME environment variable.
	<b>Description</b> : This configuration property is used to locate jars required for the Copy to Hadoop job.
oracle.hadoop.ctoh.table	Type: String
	Default Value: None.
	<b>Description</b> : The name of the database table whose content is copied to Hdoop as datapump files. It can also be schema qualified. For example, to specify the table EMPLOYEE in schema MANAGER, you can use MANAGER.EMPLOYEE
mapreduce.output.fileoutputformat.outputd	Type: String
ir	Default Value: None.
	<b>Description</b> : The name of the output directory where datapump files are created by the Hadoop job. The job output logs are also stored in the_ctoh subdirectory.
oracle.hadoop.ctoh.datapump.output	Type: String
	Default Value: None.
	<b>Description</b> : Specifies the destination directory for datapump files. If this property is not specified, the datapump files will live in the directory specified by the mapreduce.output.fileoutputformat.output dir property.
oracle.hadoop.ctoh.datapump.basename	Type: String.
	Default Value: dppart
	<b>Description</b> : The prefix or base-name of generated data pump files. For example if a user specifies this property as "dp_tbl", then the generated datapump file is dp_tbl-m-00000.dmp.
oracle.hadoop.ctoh.datapump.extension	Type:
	Default Value: dmp
	<b>Description</b> : The suffix of generated data pump files. For example if a user specifies this property as ".dp", then the generated datapump file is dppart—m-00000.dp.



Property	Description
oracle.hadoop.ctoh.maxSplits	Type: Integer.
	Default Value: None.
	<b>Description</b> : The number of datapump files that are created by the Hadoop job. This is also the number of mappers that will be produced by the Hadoop job.
oracle.hadoop.ctoh.splitterType	Туре:
	Default Value: None.
	<b>Description</b> : The type of splitters that will be used to split the table data.
	<ul> <li>BLOCK_SPLITTER: This splitter divides the table into block ranges.</li> </ul>
	<ul> <li>ROW_SPLITTER: The splitter divides the table into row ranges.</li> </ul>
	<ul> <li>PARTITION_SPLITTER: If a table is partitioned, the partition splitter can be used. When this splitter is specified, the number of datapump files created by the Hadoop job is at most equal to the number of partitions in the table.</li> </ul>
oracle.hadoop.ctoh.columns	Type:
	Default Value: None.
	Description: Specifies the subset of columns to be copied. For example if a user specifies "NAME,MANAGER" then the data for columns NAME and MANAGER are copied. If this property is not specified then all columns are copied (unless filtered by a WHERE clause).
oracle.hadoop.ctoh.whereClause	Type: String.
	Default Value: None.
	<b>Description</b> : This property is used to copy a subset of rows. For example, to copy employees whose ids are less than 1000 and greater than 500, then specify the following WHERE clause: EMPLOYEE_ID < 1000 AND EMPLOYEE_ID > 500.
oracle.hadoop.ctoh.jdbc.url	Type: String.
	Default Value: None.
	<b>Description</b> : The JDBC url to connect to the database. This property can used for demo purposes and non-secure clusters. Use Oracle Wallet with Secure Hadoop clusters in production environments.
oracle.hadoop.ctoh.connection.username	Type: String.
	Default Value: None.
	<b>Description</b> : The name of the Oracle Database user. This property can used for demo purposes and non secure clusters. Use Oracle Wallet with Secure Hadoop clusters in production environments.



Property	Description
oracle.hadoop.ctoh.connection.password	Type: String.
	Default Value: None.
	Description: The password of the Oracle Database user. This property can used for demo purposes and non secure clusters. Use Oracle Wallet with Secure Hadoop clusters in production environments.
oracle.hadoop.ctoh.connection.walletLoc	Type: String.
	Default Value: None.
	<b>Description</b> : Location of the Oracle Wallet directory on the Hadoop client.
	When using an Oracle Wallet, you must also set the following properties:
	• oracle.hadoop.ctoh.connection.tnsAdm in
	• oracle.hadoop.ctoh.connection.tnsEnt ry
oracle.hadoop.ctoh.connection.clusterWall	Type:
etLoc	Default Value: Value of
	oracle.hadoop.ctoh.connection.walletLoc.
	Description: Location of the Oracle wallet directory on the Hadoop cluster. NOTE: This directory must be visible to all nodes in the Hadoop cluster. When using this property, you must also set the following properties:
	• oracle.hadoop.ctoh.connection.cluste rTnsAdmin
	• oracle.hadoop.ctoh.connection.tnsEnt ry
oracle.hadoop.ctoh.connection.tnsAdmin	Type: String.
	Default Value: Not Defined.
	Description: Location of the directory on the Hadoop client, containing the SQL*Net configuration files such as sqlnet.ora and tnsnames.ora. Set this property so that you can use TNS entry names in the database connection strings. When using Oracle Wallet, this property value can be the same as the value of the oracle.hadoop.ctoh.connection.walletLoc property.
	You must set this property when using an Oracle Wallet as an external password store. See oracle.hadoop.ctoh.connection.walletLoc.



Property	Description
oracle.hadoop.ctoh.connection.clusterTns	Type: String.
Admin	<b>Default Value</b> : Value of the property oracle.hadoop.ctoh.connection.tnsAdmin.
	Description: Location of the directory on the Hadoop cluster containing SQL*Net configuration files such as sqlnet.ora and tnsnames.ora. NOTE: This directory must be visible to all nodes in the Hadoop cluster.
	Set this property so that you can use TNS entry names in database connection strings. When using Oracle Wallet, this property value can be the same as the value of oracle.hadoop.ctoh.connection.clusterWalletLoc.
	You must set this property when using an Oracle Wallet as an external password store (as Oracle recommends). See oracle.hadoop.ctoh.connection.clusterWalletLoc.
oracle.hadoop.ctoh.connection.tnsEntry	Type: String.
	Default Value: None.
	Description: The TNS entry name defined in the tnsnames.ora file. Use this property with oracle.hadoop.ctoh.connection.tnsAdmin. When using Oracle Wallet, make sure that the tnsEntry name matches your wallet credential.
oracle.hadoop.ctoh.cachePath	Type: String.
	Default Value: \$
	$\label{lem:mappeduce.output.fileoutputformat.outputdir} \{ \texttt{mapreduce.output.fileoutputformat.outputdir} \} //ctohCache$
	Description: Identifies the full path to an HDFS directory where cCopy to Hadoop can create files that are loaded into the MapReduce distributed cache. The distributed cache is a facility for caching large, application-specific files and distributing them efficiently across the nodes in a cluster.



D

# Diagnostic Tips and Details

The following is a collection of notes that can be useful for troubleshooting and for general understanding of Oracle Big Data SQL.

# D.1 Running Diagnostics with bdschecksw

On the Oracle Database server, you can use the Big Data SQL Diagnostics Collection tool, bdschecksw, to do a simple sanity test of Oracle Big Data SQL. This script gathers and analyzes diagnostic information about the Oracle Big Data SQL installation from both the Oracle Database and the Hadoop cluster sides of the installation. The script is in \$ORACLE HOME/bin on the database server.

## **Syntax**

bdschecksw Required\_Params [Options]

The table below describes the requireed and optional parameters used with bdschecksw.

Table D-1 bdschecksw Parameters

Parameter	Description	Required or Optional
-h,help	Display command help and exit.	Optional
-d,dbhome ORACLE_HOME	The path to the Oracle installation on the Oracle Database server.	Required only if the ORACLE_HOME environment variable is not defined on the Oracle Database server where bdschecksw is executed.
-s,sid=[ORACLE_SID]	The SID of the Oracle Database server.	Required only if the ORACLE_SID environment variable is not defined on the Oracle Database server where bdschecksw is executed.

Table D-1 (Cont.) bdschecksw Parameters

Parameter	Description	Required or Optional
-g,gihome=Grid_Infrastructure_home Oracle_Database_node_IP_address	The Grid Infrastructure path on the Oracle Database server.	Required only if the GI_HOME environment variable is not defined on the Oracle Database server where bdschecksw is executed.
<pre>-y,giuser Oracle_Database_node_IP_address Grid_Infrastructure_home</pre>	GI_HOME administrator name or owner (OS user name) of GI_HOME.	
-q,sqlplus Oracle_Database_node_IP_address username	SQLPlus username on the Oracle Database server. The user is prompted for the password.	Required.
-c,cell DNS short name [n]	The Hadoop cluster cell nodes. Use DNS short names (FQDN minus domain name) with a space as a delimiter. IP addresses are not recommended, because a node may exist on multiple networks.	Required for Hortonworks HDP only.
-u,uname Hadoop_cluster_node_username	Credentials to run remote commands on the Hadoop cluster from the Oracle Database server. This is usually the oracle user.	The username and password are always required.
-p,pdb= <i>PDB_CONTAINER</i>	The Pluggable Database container name on the Oracle Database server.	Required only if the Pluggable Database (PDB) container is configured on the Oracle Database server where bdschecksw is executed.



Table D-1 (Cont.) bdschecksw Parameters

Parameter	Description	Required or Optional
-k,key SSH_identity_key_file	Path to an SSH (Secure Shell) key file.	The optional SSH identity (or key) file is used on top of -u and -p to selects a file from which the identity (private key) for public key authentication is read.
-r,cluster Hadoop_cluster_name	Hadoop cluster name.	Optional.
-t,trace	Turn on extproc and log4j tracing during test execution.	Optional.
-f,file=file_path	Redirect output to the file.	Optional.
-i,interactive	Enter command line arguments interactively (rather than all at once).	Optional.
-x	Extensive mode.	Optional. Requires root privilege.
-v,verbose	Verbose reporting mode.	Optional. (Recommend ed for full details in the report.)

## **Exit Status**

The bdschecksw script returns one of the following status codes.

Table D-2 Status Codes for bdschecksw

Status	Description
0	Success
1	Minor problem (for example, no response in interactive mode).
2	Significant problem (for example, an invalid command line argument).

## **Example**

\$ ./bdschecksw -d /u03/app/oracle/product/\$ORACLE\_HOME/dbhome\_1 -s orcl -p pdborcl
-g /u03/app/oracle/product/\$ORACLE\_HOME/grid -q sys -u oracle -v



# D.2 How to do a Quick Test

Here is an all-around series of basic checks to ensure that Oracle Big Data SQL is working.

- On the Oracle Database server, source the environment using the hadoop\_<hcluster>.env file in \$ORACLE\_HOME/bigdatasql.
- 2. If Kerberos is enabled, kinit as the oracle Linux user on the Oracle Database server. If possible, also kinit on each of the Big Data SQL datanodes as the oracle user.



You can run this test without running kinit on the datanodes, but then offloading in the test will not work. You will eventually need to kinit on the datanodes in order to verify that offloading is working.

- 3. Create a text file and add several of lines of random text.
- 4. Copy the text file into hdfs as /user/oracle/test.txt.

```
$ hadoop fs -put test.txt /user/oracle/test.txt
```

5. Define an Oracle external table of type <code>ORACLE\_HDFS</code>:

- 6. Define a Hive table:
  - a. Connect to Hive via Hue, the Hive/Beeline command line, or using Oracle SQL Developer with a Hive JDBC driver.
  - b. CREATE TABLE bds\_test\_hive (line string);
  - C. LOAD DATA INPATH '/user/oracle/test.txt' OVERWRITE INTO TABLE bds\_test\_hive;
- 7. Define an external ORACLE\_HIVE table:



# D.3 Oracle Big Data SQL Database Objects

Table D-3 Database Objects in Big Data SQL

Туре	Object
Directories	• DEFAULT_DIR - points to \$ORACLE_HOME/bigdatasql/databases/ <database name="">/default_dir.</database>
	• ORACLE_BIGDATA_CONFIG - points to \$ORACLE_HOME/bigdatasql/databases/ <database name="">/bigdata_config.</database>
	ORA_BIGDATA_CL_ <hcluster> — expected to have a null value for its path.</hcluster>
	This is a way of limiting access. There always must be a directory object associated with an external table. Because the directory object is used for privilege checking, this is a requirement even for Hive/HDFS, where the files do not reside under the directory.
Database Links (public) These allow Big Data SQL to reach the MTA	• BDSQL\$_DEFAULT_CLUSTER — the connect string's SID should equal bds_ <dbname>_<hcluster> . And the <hcluster> should be the default cluster (as defined by bigdata.cluster.default) in \$ORACLE_HOME/ bigdatasql/databases/<database name="">/bigdata_config/ bigdata.properties.</database></hcluster></hcluster></dbname>
(multi-threaded agent)	• BDSQL\$_ <hcluster> - the connect string's SID should equal bds_<dbname>_<hcluster>.</hcluster></dbname></hcluster>
Data Dictionary Views	• User_hive_tables, all_hive_tables, dba_hive_tables – queries all Hive tables for all Hive databases for all Hadoop clusters.
	• User_hive_databases, all_hive_databases, dba_hive_databases – queries all Hive databases for all Hadoop clusters.
	<ul> <li>User_hive_columns, all_hive_columns, dba_hive_columns - queries all hHive tables for all Hive databases for all Hadoop clusters.</li> </ul>
	<ul> <li>V\$cell – the Oracle Big Data SQL server processes running on datanodes will appear here (if properly detected by the diskmon).</li> </ul>
Functions and	• DBMS_HADOOP
Procedures for	<pre>- Create_extddl_for_hive()</pre>
Hive Data Dictionary See	• GetHiveTable — pipeline function that returns data back from the extproc external procedure. Used by the *_hive_[tables/databases/columns] views and DBMS HADOOP.
cathive.sql, dbmshadp.sql.	<ul> <li>HiveMetadata – ODCI framework defining the external procedureGetHiveTable.</li> </ul>
	<ul> <li>SYS.DBMSHADOOPLIB (libkubsagt12.so) - C library for the external procedure.</li> </ul>
	<ul> <li>HiveMetadata.jar – java library called by libkubsagt12.so.</li> </ul>
Tables	SYS.HIVE_URI\$ - security table for non-DBA users.



Table D-3 (Cont.) Database Objects in Big Data SQL

Туре	Object
Statistics	<ul> <li>All statistics have %XT% in the name:         <ul> <li>cell XT granules requested for predicate offload</li> <li>cell XT granule bytes requested for predicate offload</li> <li>cell interconnect bytes returned by XT smart scan</li> <li>cell XT granule predicate offload retries</li> <li>cell XT granule IO bytes saved by storage index</li> </ul> </li> <li>Use this query:         <ul> <li>select n.name, s.value /* , s.inst_id, s.sid */ from v\$statname n, v\$mystat s</li> <li>where n.name like '%XT%' and s.statistic# = n.statistic#;</li> </ul> </li> <li>If needed, also use: grant select any catalog to <dbuser>;</dbuser></li> </ul>

# D.4 Other Database-Side Artifacts

Table D-4 \$ORACLE\_HOME/bigdatasql Directory

Subdirectory or Filename	Description of Contents
clusters directory	Contains settings related to all clusters installed on this ORACLE_HOME. It includes a subdirectory for each cluster, which contains:  config directory – configuration files downloaded for Cloudera Manager or Ambari.  fuse directory – settings for the FUSE-DFS service for the this cluster.  hadoop, hbase, and hive soft links to the actual client directories. (For example:
	hadoop-2.6.0-cdh5.12.0,hbase-1.2.0-cdh5.12.0, hive-1.1.0-cdh5.12.0, although the versions installed may different for your system.)  *.conf - IPsec configuration file.  *.keytab - Kerberos keytab file for the database owner.



Table D-4 (Cont.) \$ORACLE\_HOME/bigdatasql Directory

## Subdirectory or Filename

## **Description of Contents**

databases directory

Contains settings related to all databases running on this ORACLE\_HOME. It includes a subdirectory for each database running on this ORACLE\_HOME. Each database subdirectory contains a bigdata\_config directory, which includes:

- bigdata.properties Defines the location of JAR files. If your Hive tables use non-standard Java libraries, you may need to copy those libraries to the database and update the classpath entries in this file. Also defines the Hadoop cluster list and default cluster. Restart the extprocbds\_<dbname>\_<hcluster> after changing.
- bigdata-log4j.properties This controls the logging of the Java pieces, such as the metadata discovery phase of querying external tables or the query fetch phase if the cells are unavailable. Change log4j.logger.oracle.hadoop.s ql to INFO to log more. Restart the extprocbds\_<dbname>\_<hcluster> after changing.
- <hcluster> directory a soft link
  to \$ORACLE\_HOME/bigdatasql/clusters/
  <cluster name>/config, which contains
  the client configuration files (like hivesite.xml) copied from the Hadoop
  cluster.

Each database subdirectory also includes:

- default\_cluster soft link to the \$ORACLE\_HOME/bigdatasql/ clusters/ subdirectory that is the default cluster for this database.
- default\_dir directory for external tables associated with this database.

Soft link to the JDK installation. The version installed by Oracle Big Data SQL is  $jdk1.8.0\_141$ , although a different version may be present.

The Oracle Big Data SQL Java JAR directory

jdk

jlib directory



Table D-4 (Cont.) \$ORACLE\_HOME/bigdatasql Directory

Subdirectory or Filename	<b>Description of Contents</b>
bigdata_config directory	<ul> <li>bigdata.properties – Defines the location of JAR files. If your Hive tables use non-standard Java libraries, you may need to copy those libraries to the database and update the classpath entries in this file. Also defines the Hadoop cluster list and default cluster. Restart the extprocbds_<dbname>_<hcluster> after changing.</hcluster></dbname></li> <li>bigdata-log4j.properties - This controls the logging of the Java pieces, such as the metadata discovery phase of querying external tables or the query fetch phase if the cells are unavailable. Change log4j.logger.oracle.hadoop.s ql to INFO to log more. Restart the extprocbds_<dbname>_<hcluster> after changing.</hcluster></dbname></li> <li><hcluster> directory - a soft link to \$ORACLE_HOME/bigdatasql/clusters/<cluster name="">/config, which contains the client configuration files (like hive-</cluster></hcluster></li> </ul>
	site.xml) copied from the Hadoop cluster.
default_dir directory	This directory is usually empty.
log directory	Contains Java log files from the two type of extprocs. Use the PID that is part of the filename to identify which extproc you are looking at (only one log file will have the PID of the currently running extprocbds_ <pre><dbname>_<hcluster> process).</hcluster></dbname></pre>
hadoop_ <cluster name="">.env</cluster>	Sets the Hadoop client environment. There is one of these .env files for each cluster installation. You can source this file and then run hadoop fs commands to quickly test Hadoop connectivity
orahivedp	A soft link to the installed cp2hadoop (Copy to Hadoop) toolkit. The version of the toolkit installed by Oracle Big Data SQL 3.2 is orahivedp-3.2.0.



Table D-5 External Procedure Agents

Agent	Description
extproc	Runs the external procedure code used by the *_hive_tables, *_hive_databases views and the DBMS_HADOOP procedure.
	\$ORACLE_HOME/hs/admin/ extproc.ora configures the extproc. You can add TRACE_LEVEL=ON to get more trace (but may need to first comment outset EXTPROC_DLLS= to fix the "Error in assignment statement" message). The C portion's log files are in \$ORACLE_HOME/hs/log/orcl_agt*, but these are usually not interesting for diagnostic purposes. The JVM portion's log files are written to \$ORACLE_HOME/bigdatasql/log (but you need to set up bigdatalog4j.properties first).
extprocbds_ <dbname>_<hcluster></hcluster></dbname>	The BDS Multi-threaded Agent that is used when querying external tables. This is started/ stopped by Oracle Clusterware which in turn runs the mtactl utility. This is registered to Oracle Clusterware when bds_database_install.sh runs on the last database server node.
	If you don't have this extprocbds_ <dbname>_<hcluster> running, then you probably didn't run bds_database_install.sh on every database server in your RAC cluster. The C portion's log files are in \$ORACLE_HOME/hs/log (but you need to edit \$ORACLE_HOME/hs/admin/initbds - add TRACE_LEVEL=ON and then restart to see logging). The JVM portion's log files go into \$ORACLE_HOME/bigdatasql/log (but you need to setup bigdata-log4j.properties and restart). This is the recommended way to restart (although the quicker way is to run kill -9 on the process):</hcluster></dbname>
	<pre>\$ crsctl stop resource bds_<dbname>_<hcluster> \$ crsctl start resource bds_<dbname>_<hcluster></hcluster></dbname></hcluster></dbname></pre>





Both of External Procedures in the table above make callbacks to the database, which can be blocked by the use of the Secure External Password Store feature. If you use Secure External Password Store (SQLNET.WALLET\_OVERRIDE=TRUE), see Document 2126903.1 in My Oracle Support.

Table D-6 Log Files and Trace Files

Directory or Filename	Description	
\$ORACLE_HOME/ bigdatasql/log	Contains log files from Java code run via the extprocbds_ <dbname>_<hcluster> - one shared file with PID equal to the extprocbds_<dbname>_<hcluster> PID and extproc (one per session if the session uses *_hive_tables or DBMS_HADOOP). Tip: This is good diagnostic information.</hcluster></dbname></hcluster></dbname>	
<pre>\$ORACLE_HOME/hs /log</pre>	Contains log files from the C code of the extproc processes (one per session) and the multi-threadedextbds_ <dbname>_<hcluster> process. The exproc is usually not interesting for diagnostic purposes. The extprocbds_* has a bit more interesting information (but you need to set TRACE_LEVEL=ON in initbds_*.ora).</hcluster></dbname>	
Database diag directory	Contains log files from the database session. These can yield good information.	
	To identify the exact database session log file location:	
	<pre>select value from v\$diag_info WHERE name = 'Default Trace File';</pre>	
	To turn on external table logging:	
	<pre>alter session set "_xt_trace"="low","compilation","execution"; • To turn on additional logging:</pre>	
	<pre>alter session set events 'trace[KCFIS] disk high, memory high';</pre>	
/u01/oracle/ diag/crs/ <hostname>/crs/ trace/diskmon.trc</hostname>	Contains diskmon logs and errors. In a commodity Oracle Database server to commodity Hadoop environment (support in Oracle Big Data SQL 3.0 and greater), check this trace file for communication errors or fencing (ENTITY_FENCED). Restart diskmon if needed (use crsctl). In a commodity-to-commodity environment, you can simply kill the diskmon process, but do not do that in an Oracle Exadata Database Machine environment.	
	If you want to get additional diskmon tracing, you can set environment parameters before you invoke the crsctl command to start the cluster. Since the cluster is likely already running, you'll first have to shut the cluster down. Then, set the environment and then start it back up. Here is how you do it in the Oracle Big Data SQL 3.x commodity database server scenario using Oracle Restart. (Note that the crsctl commands will be different if using RAC, ASM, and/or Exadata):	
	<pre>crsctl stop has export CELLCLIENT_TRACE_LEVEL="all,4" export CELLCLIENT_AUTOFLUSH_LEVEL="all,4" crsctl start has</pre>	



Table D-6 (Cont.) Log Files and Trace Files

Directory or Filename	Description
/etc/oracle/ cell/network- config/ cellinit.ora	Record the IP address and subnet range for the database server. For Oracle Big Data SQL on commodity servers, this file also includes parameters which switch the protocol away from InfiniBand RDS to TCP/UDP (_skgxp_dynamic_protocol=2). On commodity servers, the database server's diskmon (running out of the Oracle Grid home)
<pre>/etc/oracle/ cell/network- config/ celliniteth.ora</pre>	communicates with the BDS processes on the data nodes listening on TCP port 5042.
<pre>Kerberos files: kinit, klist, etc/ krb5.conf, krb5- workstation*.rpm</pre>	If your Hadoop cluster uses Kerberos, you'll need to setup Kerberos on the database and have a mechanism (such as crontab) to keep a valid Kerberos ticket at all times for the oracle Linux user. You will need a similar ticket renewal mechanism on the BDS datanodes as well.
	The Oracle Big Data SQL installer now provides a directive in the Jaguar configuration file that will automatically set up a cron job for this on both the Hadoop cluster and the Oracle Database system. See the description of the configuration file in the installation guide.

# D.5 Hadoop Datanode Artifacts

The table below identifies objects on the Hadoop server that can provide helpful information for troubleshooting Big Data SQL.

Table D-7 Hadoop-side Datanode Artifacts That are Useful for Troubleshooting

Datanode Artifact	Description
bdscli command	List quarantine detail
	<ul> <li>Drop quarantine all</li> </ul>
	<ul> <li>List alerthistory</li> </ul>
	<ul> <li>Drop alerthistory</li> </ul>
	<ul> <li>List bdsql detail</li> </ul>



Table D-7 (Cont.) Hadoop-side Datanode Artifacts That are Useful for Troubleshooting

Datanode Artifact	Description
Log files	<ul> <li>/var/log/bigdatasql/DM – installer log files</li> <li>/var/log/bigdatasql/cloudera or /var/log/bigdatasql/ambari – Ambari or CM service log files.</li> <li>/opt/oracle/bigdatasql/bdcell-<cell version="">/bigdata.properties</cell></li> <li>/opt/oracle/bigdatasql/bdcell-<cell version="">/bigdata-log4j.properties</cell></li> <li>This defaults to logging off. Change tolog4j.logger.oracle.had oop.sql=INFO and restart.</li> <li>/opt/oracle/bigdatasql/bdcell-<cell< li=""> </cell<></li></ul>
	version>/log directory  - bigdata-log4j.log - logs entries from the JVM pieces of Big Data SQL (logging defaults to off, so edit bigdata- log4j.properties first and restart). This can be particularly useful information.
	<ul> <li>/var/log/oracle/diag/bdsql/cell/ <hostname>/trace/ - general cell trace files for the Management Server, Restart Server, and Monitor Server. The alert.log file will have details about quarantine and de-quarantine events.</hostname></li> </ul>
	<ul> <li>/var/log/oracle/diag/bdsql/ cell/SYS_*/trace/ - Oracle Big Data SQL offload server trace files for the C portion. These are not useful for troubleshooting in most cases.</li> </ul>
Other datanode artifacts	<pre>/opt/oracle/cell/cellsrv/ deploy/config/cellinit.ora - records the cell's IP address.</pre>

# D.6 Step-by-Step Process for Querying an External Table

- 1. User issues a SELECT query involving an Oracle Big Data SQL external table.
- Database sees that one of the objects in the SQL is an External table of type ORACLE\_HIVE
- 3. Database identifies the cluster name from the com.oracle.bigdata.cluster parameter on the External table definition else uses the default cluster.
- 4. Database identifies the Hive table name from the com.oracle.bigdata.tablename parameter, else assumes the Hive table name is the same as the Oracle table name.



**5.** Database knows that the ORACLE\_HIVE External table implementation uses an external procedure which is invoked through the <code>extprocbds\_<dbname>\_<hcluster></code> multi-threaded agent.

## Note:

The first phase of the query requires getting the Hive metadata. If you get an error during this first phase, you'll likely see an error that begins as follows. Notice the "OPEN" in ODCIEXTTABLEOPEN)

ORA-29913: error in executing ODCIEXTTABLEOPEN callout

- 6. Database uses the public database link BDSQL\$\_DEFAULT\_CLUSTER or BDSQL\$\_<hcluster> to find the connect string to ask the listener to connect the database session to a thread of the extprocbds\_dbname>\_hcluster> multi-threaded agent
  - **a.** extprocbds\_<dbname>\_<hcluster> was previously started by Oracle Clusterware and is using configuration information from the \$ORACLE\_HOME/bigdatasql/databases/<database name>/bigdata\_config directory.
  - b. extprocbds\_<br/>dbname>\_<hcluster> has spawned a JVM running Hadoop client libraries using the above configuration information. The Hadoop client libraries were copied from the Oracle Big Data Appliance to the Oracle Database server when you ran the bds-exa-install.sh script.
- 7. extprocbds\_<dbname>\_<hcluster> uses its JVM and the Hive metastore client library to call the Hive metastore (using a URL such as thrift://hostname>:9083) to get metadata (columns, inputformat, serde, other table properties) for the Hive table.
  - a. At this point, if the Hive metastore is protected by Kerberos authentication, the Hive client libraries running in the extprocbds JVM on the Oracle Database server will try to send the local Kerberos ticket to the Hive server. This will be the ticket owned by the oracle Linux user account who is running the database
- **8.** extprocbds\_<dbname>\_<hcluster> calls the Hive metastore to get a list of input paths that hold the data behind the Hive table.
- 9. extprocbds\_<abname>\_<hcluster> converts the list of input paths into a list of splits/ blocks using Hadoop MapReduce libraries and logic. Then it asks the HDFS namenode for the location (including replicas) of all of the splits /blocks.
  - **a.** Again, if HDFS is protected by Kerberos, the Kerberos ticket from the oracle Linux user account on the database will be need to be used.
  - b. If compression is used, at this point the JVM might have to load specific compression Java or native libraries. If these are non-standard libraries, you will need to install them on both the Oracle Database server and the Hadoop side. For instance, LZO compression requires an additional install and configuration performed on both the database-side on the Hadoop-side.

At this point, the "description" phase is done and the database knows the structure of the Hive table as well as the location of all of the blocks of data (including replicas). This information is also known as the metadata payload. We now begin the "fetch" phase.

10. The database intelligent storage layer, KCFIS (Kernel Cache File Intelligent Storage), which is also used on Oracle Exadata systems, compares the



hostnames of where the blocks of data are stored to a list of active BDSQL server hosts being maintained by the Grid's diskmon process. (You can see diskmon's list of BDSQL server hosts in V\$CELL).



The second phase of the query requires fetching the data. If you get an error during this second phase, you'll likely see an error that begins as follows. Notice the "FETCH" in ODCIEXTTABLEFETCH):

ORA-29913: error in executing ODCIEXTTABLEFETCH callout

11. Assuming that the list of datanode hostnames matches the list of BDSQL hostnames, the database sends a list of local blocks (also called Granules) to each of the BDSQL servers. The database also sends the BDSQL servers metadata about the table, columns, and structure it is accessing. It does this in parallel and asynchronously for performance



The database statistics "cell XT granules requested for predicate offload" and "cell XT granule bytes requested for predicate offload" are updated at this point

- 12. The BDSQL process running on the data nodes checks the SQL\_ID against its local list of quarantined SQL\_IDS. If the SQL\_ID matches the quarantine list, then the BDSQL process on the datanode will return an error. However, the user should not see this error. Instead, the database will first try another cell, then try to do the work itself. (See Steps 15 and 16).
- 13. Assuming the SQL\_ID is not quarantined by the BDSQL process on the datanode, the BDSQL process will do its SmartScan work against the list of blocks/granules sent to it.



## Tip:

See the blog entry Big Data SQL Quick Start. Storage Indexes - Part10 in The Data Warehouse Insider for details about Storage Indexes and other aspects of SmartScan processing.

- a. The BDSQL offload process has previously read its configuration information from /opt/oracle/bigdatasql/bdcell-<cell version>/bigdata.properties.
- **b.** The BDSQL process has previously loaded a JVM based on the properties defined in the above configuration.
- c. If the Hive table has special InputFormat or Serde classes, the JVM will load those classes assuming it can find them based on the classpath defined in the above configuration. For some common InputFormats (such as delimited text), Oracle has written C code that can handle those formats faster than regular Java code.



- d. If Kerberos authentication is used, then the BDSQL's JVM will send its local Kerberos ticket to the HDFS datanode process. This is the Kerberos ticket associated with the oracle Linux user on the datanode where BDSQL is running.
- e. If Sentry authorization is used, the <code>oracle</code> Linux user's Kerberos ticket's identity needs to have been granted access to the Hive table and underlying HDFS data.
- f. The BDSQL server will update statistics like "cell XT granule IO bytes saved by StorageIndex" as it runs.
- 14. The database kcfis layer will collect results as they are returned from the BDSQL processes on the datanodes and send the next batch of blocks/granules to the BDSQL processes.
  - a. The database will update the "cell interconnect bytes returned by XT smart scan" statistic as bytes are returned
- **15.** If there are issues with a BDSQL process for a given block, the database will try to send the work to a different BDSQL process (it will pick a location that has a replica of the block that failed).
  - a. The database will update the "cell XT granule predicate offload retries" statistic.
- **16.** If the database is unable to get the BDSQL processes to successfully offload a block even after retrying, then the database will "fallback" and have the JVM in the extprocbds\_<db>\_<cluster> do the work.
  - This will be slower as the raw data will need to be moved to the database server for processing.
  - b. If the Hive table has any special InputFormats or Serdes, the extprocbds\_<db>\_<cluster>'s JVM will need to load them based on the classpath configuration defined on the database's bigdata.properties file.
- **17.** The results from the external table data source continue to be collected until all input paths/blocks/granules are handled.

# D.7 Step-by-Step for a Hive Data Dictionary Query

- 1. User queries one of the Oracle Big Data SQL data dictionary views, such as all\_hive\_tables.
  - In Oracle Big Data SQL 2.0 and earlier, if this was the <code>user\_hive\_\*</code> view and the user was not a DBA, then the user needed to be listed in the <code>sys.hive\_uri\$</code> table. Oracle Big Data SQL 3.0 removed the <code>hive\_uri\$</code> check.
- 2. The view accesses the GetHiveTable pl/sql pipelined table function.
- 3. The GetHiveTable function is implemented by the HiveMetadata type which is implemented as an external procedure using the SYS.DBMSHADOOPLIB library.
- 4. The Oracle Database spawns a new instance of the "extproc" for this database session. The extproc reads the \$ORACLE\_HOME/hs/admin/extproc.ora file for settings.
  - You can set  ${\tt TRACE\_LEVEL=ON}$  for tracing of the C code. Log file will be written to  ${\tt SORACLE\_HOME/hs/log}$ .



By default, there may be an error in the <code>extproc.ora</code>, causing an "Error in assignment statement" message in the log . The statement "SET EXTPROC\_DLLS=" (with no value after the equal sign) is not valid. Comment this line out if you want to use <code>TRACE\_LEVEL=ON</code> .

- 5. The extproc attaches the libkubsagt.so library (as in SYS.DBMSHADOOPLIB).
- 6. Libkubsagt12.so initiates a JVM and loads the HiveMetadata.jar.
  - a. The JVM uses the configuration information in \$ORACLE\_HOME/ bigdatasql/bigdata\_config/ to identify the list of clusters and their Hive metastore connection information.
  - b. Logging for the JVM is based on \$ORACLE\_HOME/bigdatasql/ bigdata\_config/bigdata-log4j.properties. Log files will be written to \$ORACLE\_HOME/bigdatasql/log. There will be a new log file for each database session.
- 7. The Java code in HiveMetadata. jar uses the Hive metastore client libraries to connect to the Hive metastore to retrieve data about all of the databases and all of the tables.
  - **a.** If the Hive metastore is protected by Kerberos, the JVM will try to send the Kerberos ticket of the oracle Linux user who is running the database
- 8. The Java code returns the request data back to the database.

# D.8 Key Adminstration Tasks for Oracle Big Data SQL

Restarting the extprocbds\_<db>\_<hcluster>:

```
$ crsctl stop res bds_<dbname>_<hcluster>
```

Quick way, but not the best way: kill the  ${\tt extprocbds\_*}$  process and wait for it to come back

Restarting the extproc.

This begins a new database session.

- Restarting the Oracle Big Data SQL software on the datanodes:
  - Use Cloudera Manager or the Ambari Web UI.
  - Quick way, but not the best way: kill the bdsqloflsrv process and wait for it to come back.
  - Command line method on an Oracle Big Data Appliance (logged on as root on node1):

```
$ bdacli stop big_data_sql_cluster
$ bdacli start big_data_sql_cluster
```

Checking for Oracle Big Data SQL quarantines on a single datanode:

```
$ bdscli -e list quarantine detail
```

To check for quarantines on all datanodes:

```
$ dcli -g cells.lst bdscli -e list quarantine detail
```

Clearing Oracle Big Data SQL quarantines on a single datanode:

```
$ bdscli -e drop quarantine all
```



## To clear quarantines on all datanodes:

\$ dcli -g cells.lst bdscli -e drop quarantine all

- Checking statistics for proper offloading:
  - Use the Sql Monitor hint: /\*+ MONITOR\*/.
  - Query XT statistics. Ensure that "retries" is zero and "bytes returned" is greater than zero.
- Looking for log files on the datanodes:
  - 1. Clear quarantines on all datanodes
  - 2. Set Log property in /opt/oracle/bigdatasql/bdcell-12.1/bigdata-log4j.properties on datanodes.
  - 3. Restart bdsgloflsrv on datanodes.
  - **4.** Cd to the log file directory: /opt/oracle/bigdatasql/bdcell-12.1/ log.
  - 5. tail -f bigdata-log4j.log
  - 6. Ensure that your query has data on the node you are looking at (i.e. your query should need to access files with many blocks. If you only touch a small number of blocks, the result may be that your datanode is not be asked to do any work)
  - 7. Make a new database session (to reset XT statistics) and Run query.



#### Tip:

Use the  $\slash \text{MONITOR*/}$  hint if you want to be sure to see it in SQL Monitor.

You should see new entries in the datanode's bigdata-log4j.log.

- **8.** On the Oracle Database server, query XT statistics and check that retries=0 and bytes returned>0.
- Looking for log files on the database:
  - Clear quarantines on all data nodes
  - 2. Make a new database session (to reset XT statistics)
  - 3. Find out what instance your session is connected to (in case you got load-balanced to a different database server than the one you logged on to):

select host\_name from v\$instance;

- 4. Log in to that instance's database server at the Linux level.
- **5. Set log properties in** \$ORACLE\_HOME/bigdatasql/bigdata-log4j.properties.
- **6.** Restart extprocbds\_<db>\_<hcluster> on that instance to pick up the log property changes
- 7. Turn on XT tracing:

This command turns on external table logging:



```
alter session set "_xt_trace"="low", "compilation", "execution";
```

This command adds additional tracing:

alter session set events 'trace[KCFIS] disk high, memory high';

8. Run the query.



## Tip:

Use the /\*+ MONITOR \*/ hint if you want to be sure to see it in SQL Monitor.

9. Query XT statistics and see if retries=0 and bytes returned>0.

```
select n.name, s.value /* , s.inst_id, s.sid */ from v$statname n, gv$mystat
s where n.name like '%XT%' and s.statistic# = n.statistic#;
```

- 10. Look at JVM log file: \$ORACLE\_HOME/bigdatasql. (Look for the one with the same PID as the extprocbds\_\* process.)
- 11. Look at database trace\_file:

```
select value from v$diag_info WHERE name = 'Default Trace File';
```

# D.9 Additional Java Diagnostics

You can add JVM properties to the bigdata.properties file as shown below. This
can be good for hard-to-spot low-level Kerberos issues.

```
java.options=-Dsun.security.krb5.debug=true
```

The extproc and extprocbds\_<dbname>\_<hcluster> processes run the JVMs on the
database and the bdsqloflsrv process runs the JVM on the datanode. You can
see this by running the "ips" command:

```
$ORACLE_HOME/bigdatasql/jdk*/bin/jps
```

 If you are very comfortable with your Java skills, you can also use Oracle JVisualVM or Oracle JConsole to connect to the JVMs.

# D.10 Checking for Correct Oracle Big Data SQL Patches

Patch and Datapatch errors can have a number of different causes and effects. One thing you can do is check to ensure that the expected patches are loaded.

If you see "wrong number or types of arguments in a call to 'FETCH\_OPEN' in the error stack

Here is an example of an error stack that may warrant a query of DBA\_REGISTRY\_SQLPATCH to determine if the correct patches are loaded:

```
ORA-29913: error in executing ODCIEXTTABLEFETCH callout
ORA-29400: data cartridge error
ORA-06550: line 1, column 25:
PLS-00306: wrong number or types of arguments in call to 'FETCH_OPEN'
ORA-06550: line 1, column 14:PL/SQL: Statement ignored
```



This may indicate one of several problems.

- A Bundle Patch was not applied correctly
- Datapatch did not run and therefore a patch required by the installed version of Oracle Big Data SQL is not installed

In this case, run the following query to determine if patches identified as requirements for this installation in the *Oracle Big Data SQL Master Compatibility Matrix* (Doc ID 2119369.1 in My Oracle Support) have been applied.

 $\verb|select PATCH_ID, PATCH_UID, VERSION, STATUS, DESCRIPTION from DBA_REGISTRY_SQLPATCH order by BUNDLE_SERIES|\\$ 

For example, for Oracle Big Data SQL 3.0.1 with BP 12.1.0.2.160419 (22806133), the query should return these results.

PATCH_ID DESCRIPTION	PATCH_UID	VERSION	STATUS		
				_	
22806133	19983161	12.1.0.2	SUCCESS	DATABASE	BUNDLE
PATCH: 12.1.	0.2.160419	(22806133)			

If the query fails or the correct patch for the installed bundle is not found, see 1609718.1 in My Oracle Support for more information about issues with Datapatch.



# **Licensing Information**

# E.1 Oracle Big Data SQL Licensing

The licensing for Oracle Big Data SQL is separate from the licensing for other Oracle products.

When you purchase licensing for Oracle Big Data SQL, note the following:

- A separate license must be procured for each Hadoop cluster.
- All data nodes within the Hadoop cluster must be licensed. Partial licensing within a node is not available.
- All disks accessed by Oracle Big Data SQL (either within or outside of the Hadoop cluster) must be licensed. This includes disks used by Kafka or NoSQL stores when Oracle Big Data SQL accesses those sources.
- There is no additional license required for the Oracle Database server side.
- Oracle Copy to Hadoop licensing is included.

## Third Party Licensing for Oracle Shell for Hadoop Loaders

Oracle Shell for Hadoop Loaders is included in the Oracle Big Data SQL installation bundle. The following are third-party projects contained in Oracle Shell for Hadoop Loaders.

- ANTLR 4.7
- Apache Commons Exec 1.3

Unless otherwise specifically noted, or as required under the terms of the third party license (e.g., LGPL), the licenses and statements herein, including all statements regarding Apache-licensed code, are intended as notices only.

## E.1.1 ANTLR 4.7

Copyright (c) 2015 Terence Parr, Sam Harwell

All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

- Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.
- Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

**3.** The name of the author may not be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE AUTHOR "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE AUTHOR BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

## E.1.2 JLine 2.13

Copyright (c) 2002-2015 Marc Prud'hommeaux <mwp1@cornell.edu>

All rights reserved.

Redistribution and use in source and binary forms, with or without modification, are permitted provided that the following conditions are met:

Redistributions of source code must retain the above copyright notice, this list of conditions and the following disclaimer.

Redistributions in binary form must reproduce the above copyright notice, this list of conditions and the following disclaimer in the documentation and/or other materials provided with the distribution.

Neither the name of JLine nor the names of its contributors may be used to endorse or promote products derived from this software without specific prior written permission.

THIS SOFTWARE IS PROVIDED BY THE COPYRIGHT HOLDERS AND CONTRIBUTORS "AS IS" AND ANY EXPRESS OR IMPLIED WARRANTIES, INCLUDING, BUT NOT LIMITED TO, THE IMPLIED WARRANTIES OF MERCHANTABILITY AND FITNESS FOR A PARTICULAR PURPOSE ARE DISCLAIMED. IN NO EVENT SHALL THE COPYRIGHT OWNER OR CONTRIBUTORS BE LIABLE FOR ANY DIRECT, INDIRECT, INCIDENTAL, SPECIAL, EXEMPLARY, OR CONSEQUENTIAL DAMAGES (INCLUDING, BUT NOT LIMITED TO, PROCUREMENT OF SUBSTITUTE GOODS OR SERVICES; LOSS OF USE, DATA, OR PROFITS; OR BUSINESS INTERRUPTION) HOWEVER CAUSED AND ON ANY THEORY OF LIABILITY, WHETHER IN CONTRACT, STRICT LIABILITY, OR TORT (INCLUDING NEGLIGENCE OR OTHERWISE) ARISING IN ANY WAY OUT OF THE USE OF THIS SOFTWARE, EVEN IF ADVISED OF THE POSSIBILITY OF SUCH DAMAGE.

## E.1.3 Apache Commons Exec 1.3

Include the following License ONLY ONCE in the documentation even if there are multiple products licensed under the license.

The following applies to all products licensed under the Apache 2.0 License:

You may not use the identified files except in compliance with the Apache License, Version 2.0 (the "License.").

You may obtain a copy of the License at http://www.apache.org/licenses/LICENSE-2.0. A copy of the license is also reproduced below.

## E.1.4 Apache Licensed Code

The following is included as a notice in compliance with the terms of the Apache 2.0 License, and applies to all programs licensed under the Apache 2.0 license:

You may not use the identified files except in compliance with the Apache License, Version 2.0 (the "License.")

You may obtain a copy of the License at

http://www.apache.org/licenses/LICENSE-2.0

A copy of the license is also reproduced below.

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied.

See the License for the specific language governing permissions and limitations under the License.

## E.1.5 Apache License

Version 2.0, January 2004

http://www.apache.org/licenses/

TERMS AND CONDITIONS FOR USE, REPRODUCTION, AND DISTRIBUTION

## 1. Definitions

"License" shall mean the terms and conditions for use, reproduction, and distribution as defined by Sections 1 through 9 of this document.

"Licensor" shall mean the copyright owner or entity authorized by the copyright owner that is granting the License.

"Legal Entity" shall mean the union of the acting entity and all other entities that control, are controlled by, or are under common control with that entity. For the purposes of this definition, "control" means (i) the power, direct or indirect, to cause the direction or management of such entity, whether by contract or otherwise, or (ii) ownership of fifty percent (50%) or more of the outstanding shares, or (iii) beneficial ownership of such entity.

"You" (or "Your") shall mean an individual or Legal Entity exercising permissions granted by this License.

"Source" form shall mean the preferred form for making modifications, including but not limited to software source code, documentation source, and configuration files.

"Object" form shall mean any form resulting from mechanical transformation or translation of a Source form, including but not limited to compiled object code, generated documentation, and conversions to other media types.



"Work" shall mean the work of authorship, whether in Source or Object form, made available under the License, as indicated by a copyright notice that is included in or attached to the work (an example is provided in the Appendix below).

"Derivative Works" shall mean any work, whether in Source or Object form, that is based on (or derived from) the Work and for which the editorial revisions, annotations, elaborations, or other modifications represent, as a whole, an original work of authorship. For the purposes of this License, Derivative Works shall not include works that remain separable from, or merely link (or bind by name) to the interfaces of, the Work and Derivative Works thereof.

"Contribution" shall mean any work of authorship, including the original version of the Work and any modifications or additions to that Work or Derivative Works thereof, that is intentionally submitted to Licensor for inclusion in the Work by the copyright owner or by an individual or Legal Entity authorized to submit on behalf of the copyright owner. For the purposes of this definition, "submitted" means any form of electronic, verbal, or written communication sent to the Licensor or its representatives, including but not limited to communication on electronic mailing lists, source code control systems, and issue tracking systems that are managed by, or on behalf of, the Licensor for the purpose of discussing and improving the Work, but excluding communication that is conspicuously marked or otherwise designated in writing by the copyright owner as "Not a Contribution."

"Contributor" shall mean Licensor and any individual or Legal Entity on behalf of whom a Contribution has been received by Licensor and subsequently incorporated within the Work.

- 2. Grant of Copyright License. Subject to the terms and conditions of this License, each Contributor hereby grants to You a perpetual, worldwide, non-exclusive, no-charge, royalty-free, irrevocable copyright license to reproduce, prepare Derivative Works of, publicly display, publicly perform, sublicense, and distribute the Work and such Derivative Works in Source or Object form.
- 3. Grant of Patent License. Subject to the terms and conditions of this License, each Contributor hereby grants to You a perpetual, worldwide, non-exclusive, nocharge, royalty-free, irrevocable (except as stated in this section) patent license to make, have made, use, offer to sell, sell, import, and otherwise transfer the Work, where such license applies only to those patent claims licensable by such Contributor that are necessarily infringed by their Contribution(s) alone or by combination of their Contribution(s) with the Work to which such Contribution(s) was submitted. If You institute patent litigation against any entity (including a cross-claim or counterclaim in a lawsuit) alleging that the Work or a Contribution incorporated within the Work constitutes direct or contributory patent infringement, then any patent licenses granted to You under this License for that Work shall terminate as of the date such litigation is filed.
- **4. Redistribution**. You may reproduce and distribute copies of the Work or Derivative Works thereof in any medium, with or without modifications, and in Source or Object form, provided that You meet the following conditions:
  - You must give any other recipients of the Work or Derivative Works a copy of this License; and
  - **b.** You must cause any modified files to carry prominent notices stating that You changed the files; and
  - c. You must retain, in the Source form of any Derivative Works that You distribute, all copyright, patent, trademark, and attribution notices from the Source form of the Work, excluding those notices that do not pertain to any part of the Derivative Works; and



d. If the Work includes a "NOTICE" text file as part of its distribution, then any Derivative Works that You distribute must include a readable copy of the attribution notices contained within such NOTICE file, excluding those notices that do not pertain to any part of the Derivative Works, in at least one of the following places: within a NOTICE text file distributed as part of the Derivative Works; within the Source form or documentation, if provided along with the Derivative Works; or, within a display generated by the Derivative Works, if and wherever such third-party notices normally appear. The contents of the NOTICE file are for informational purposes only and do not modify the License. You may add Your own attribution notices within Derivative Works that You distribute, alongside or as an addendum to the NOTICE text from the Work, provided that such additional attribution notices cannot be construed as modifying the License.

You may add Your own copyright statement to Your modifications and may provide additional or different license terms and conditions for use, reproduction, or distribution of Your modifications, or for any such Derivative Works as a whole, provided Your use, reproduction, and distribution of the Work otherwise complies with the conditions stated in this License.

- 5. Submission of Contributions. Unless You explicitly state otherwise, any Contribution intentionally submitted for inclusion in the Work by You to the Licensor shall be under the terms and conditions of this License, without any additional terms or conditions. Notwithstanding the above, nothing herein shall supersede or modify the terms of any separate license agreement you may have executed with Licensor regarding such Contributions.
- 6. Trademarks. This License does not grant permission to use the trade names, trademarks, service marks, or product names of the Licensor, except as required for reasonable and customary use in describing the origin of the Work and reproducing the content of the NOTICE file.
- 7. Disclaimer of Warranty. Unless required by applicable law or agreed to in writing, Licensor provides the Work (and each Contributor provides its Contributions) on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied, including, without limitation, any warranties or conditions of TITLE, NON-INFRINGEMENT, MERCHANTABILITY, or FITNESS FOR A PARTICULAR PURPOSE. You are solely responsible for determining the appropriateness of using or redistributing the Work and assume any risks associated with Your exercise of permissions under this License.
- 8. Limitation of Liability. In no event and under no legal theory, whether in tort (including negligence), contract, or otherwise, unless required by applicable law (such as deliberate and grossly negligent acts) or agreed to in writing, shall any Contributor be liable to You for damages, including any direct, indirect, special, incidental, or consequential damages of any character arising as a result of this License or out of the use or inability to use the Work (including but not limited to damages for loss of goodwill, work stoppage, computer failure or malfunction, or any and all other commercial damages or losses), even if such Contributor has been advised of the possibility of such damages.
- 9. Accepting Warranty or Additional Liability. While redistributing the Work or Derivative Works thereof, You may choose to offer, and charge a fee for, acceptance of support, warranty, indemnity, or other liability obligations and/or rights consistent with this License. However, in accepting such obligations, You may act only on Your own behalf and on Your sole responsibility, not on behalf of any other Contributor, and only if You agree to indemnify, defend, and hold each



Contributor harmless for any liability incurred by, or claims asserted against, such Contributor by reason of your accepting any such warranty or additional liability.

## **END OF TERMS AND CONDITIONS**

## **APPENDIX:** How to apply the Apache License to your work

To apply the Apache License to your work, attach the following boilerplate notice, with the fields enclosed by brackets "[]" replaced with your own identifying information. (Don't include the brackets!) The text should be enclosed in the appropriate comment syntax for the file format. We also recommend that a file or class name and description of purpose be included on the same "printed page" as the copyright notice for easier identification within third-party archives.

Copyright [yyyy] [name of copyright owner]

Licensed under the Apache License, Version 2.0 (the "License"); you may not use this file except in compliance with the License. You may obtain a copy of the License at

http://www.apache.org/licenses/LICENSE-2.0

Unless required by applicable law or agreed to in writing, software distributed under the License is distributed on an "AS IS" BASIS, WITHOUT WARRANTIES OR CONDITIONS OF ANY KIND, either express or implied. See the License for the specific language governing permissions and limitations under the License.

This product includes software developed by The Apache Software Foundation (http://www.apache.org/) (listed below):



# Index

A	D
access drivers, 1-2, A-1 ACCESS PARAMETERS clause, 2-24 special characters, 4-8 syntax rules, 4-7 ACCESS PARAMETERS Clause syntax, 4-7 ALL_HIVE_COLUMNS view, 4-23 ALL_HIVE_DATABASES view, 4-22 ALL_HIVE_TABLES view, 2-2, 4-22 array overflows, 4-18	data dictionary views, 4-21 data mode, 4-12 data source name, 4-21 data type conversion (Big Data SQL), 2-24 data types (HDFS), 4-14 DBA_HIVE_COLUMNS view, 4-25 DBA_HIVE_DATABASES view, 4-24 DBA_HIVE_TABLES view, 4-24 DBMS_HADOOP package, 4-1 DBMS_OUTPUT package, 2-3 DEFAULT DIRECTORY clause, 2-23 delimited text files, 4-19
BDSQL_USER, <i>F-3</i>	E
bigdata_config directory, 2-26 BigDataSQL configuration directory object, F-3 binary overflows, 4-18	error handling, 4-12 error handling (Big Data SQL), 2-26 external tables about, 1-2, A-1
<u>C</u>	, ,
catalog views, <i>4-21</i> Cell XT, <i>1-10</i>	F
character overflows, 4-18 column mapping, 4-11 com.oracle.bigdata.buffersize, 4-10	field extraction, 4-19 field names, 4-14
com.oracle.bigdata.colmap, 4-11 com.oracle.bigdata.datamode, 4-12	Н
com.oracle.bigdata.erroropt, 4-12 com.oracle.bigdata.fields, 4-14 com.oracle.bigdata.fileformat, 4-16 com.oracle.bigdata.log.exec, 4-17 com.oracle.bigdata.log.qc, 4-17 com.oracle.bigdata.overflow, 4-18 com.oracle.bigdata.rowformat, 4-19 com.oracle.bigdata.tablename, 4-21 CREATE TABLE ORGANIZATION EXTERNAL syntax, 2-22, A-2 CREATE TABLE statement	Hadoop log files, 4-10, 4-17 Hive columns, 4-23 Hive data     access from Oracle Database, 2-1 Hive databases, 4-22 Hive table sources, 4-21 Hive tables, 4-22 Hive views, 4-21 L
generating automatically for Hive, 4-1 CREATE_EXTDDL_FOR_HIVE function syntax, 4-1	licenses, third-party, <i>E-1</i> LOCATION clause, <i>2-23</i> log files, <i>4-17</i>



0	REJECT LIMIT clause, 2-24 row format description, 4-16 row formats, 4-19		
Oracle Big Data SQL			
access drivers, 1-2			
data type conversion, 2-24 general description, 1-1	S		
installation changes on the Oracle Database	sequence files, 4-16		
server ., 2-26 Oracle Database	SerDe parsing, 4-19		
access to Hive data, 2-1	Smart Scan, 1-2		
Data Modeler, 2-1	SmartScan mode, 4-12		
DBMS HADOOP, 2-1	source name, 4-21 static data dictionary views, 4-21		
SQL Developer, 2-1	Statistics, 1-10		
Use in Oracle Big Data SQL, 2-1	struct overflows, 4-18		
Oracle Exadata Machine			
Big Data SQL installation changes, 2-26	Т		
ORACLE_HDFS access driver, 2-15	<u> </u>		
ORACLE_HIVE	text files, 4-16		
access parameters, 4-9	text overflows, 4-18		
ORACLE_HIVE examples, 2-6 ORC files, 4-16	third-party licenses, <i>E-1</i>		
overflow handling, 4-18	TYPE clause, 2-22		
overnow nanding, 4 10			
Р	U		
Danwart Flora 4.40	union overflows, 4-18		
Parquet files, 4-16	user access from Oracle Database, 2-25		
parsing HDFS files, 4-19 PL/SQL packages, 4-1	USER_HIVE_COLUMNS view, 4-25		
PUT_LINE function, 2-3	USER_HIVE_DATABASES view, 4-25		
TOT_ENVERGINGTION, 20	USER_HIVE_TABLES view, 4-25		
D	users		
R	BDSQL_USER, F-3		
RC files, 4-16			

